



The Future of Object Persistence

St Louis JUG

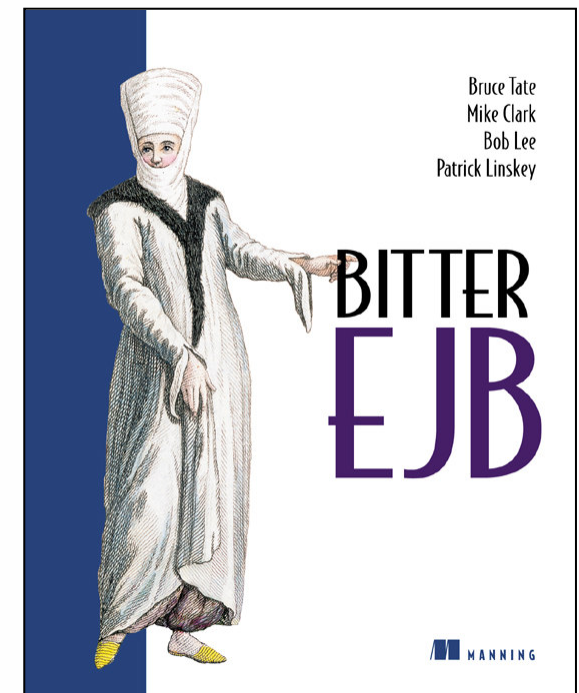
13 January 2005

About Me



Patrick Linskey

- CTO at SolarMetric
- Involved in object/relational mapping and EJB since 1999
- Frequent JDO presenter at JUGs, Java conventions and seminars
- Member, JDO Expert Group
- Member, EJB Expert Group
- Luminary, JDOcentral
- Co-author of *Bitter EJB* along with Bruce Tate, Mike Clark, and Bob Lee



Corporate Profile



- International Company Based in Austin, TX
 - Offices in London, California, Massachusetts
 - Through partners, reach is worldwide
- Founded in 2001 by MIT alums.
 - Core team has been together since 1997.
- Frustrated with trying to do Java object persistence with:
 - Proprietary Tools – Vendor Lock-in
 - Entity beans – difficult, slow, impose undesirable constraints on object model and development patterns
- Committed to Technical Quality and Innovation, Customer Support
- Leading JDO implementation
- Client base is diverse, both in terms of industry and size
 - 300+ customers
- Regular JDO Training courses throughout the world

SolarMetric's Role



JSR 220 (EJB 3)

- Actively contributing member
- Actively developing preview version of the EJB 3 specification

JSR 243 (JDO 2)

- Will continue active role on the development and release of JDO 2
- Will continue active development of Kodo's JDO bindings, for JDO 2 and for future JDO versions

Object / Relational Mapping

Object / Relational Mapping



- Object-oriented programming languages are vastly different than relational data languages.
- Object / relational “impedance mismatch” has plagued enterprise programmers for years.
 - Mapping
 - Remoteness of data
- Proprietary O/R mapping products exist for a variety of languages: Smalltalk, C++, Java, others.
- JDO and EJB3
 - designed for object/relational mapping and
 - designed to work inside and outside a container

Why not just JDBC?



- JDBC misses the “O” part of O/R mapping
 - Interface is not at an object level but rather at SQL (row and column) level
 - Not Java
 - Creates complexity especially when leveraging OO concepts e.g., inheritance, polymorphism
- JDBC is a low level API
 - Used as a building block by most O/R mapping tools
- Sadly, SQL is not portable
 - Many different dialects

Goals of an O/R specification



- Abstract API for object persistence
- No need to write persistence infrastructure
- Standard means of persisting objects. Low risk of vendor lock-in
- Portability between data stores

Requirements of an O/R Specification



- Persist objects whether simple or complex mapping is required
- Query those persisted objects
- Minimize visibility of O/R mapping APIs
- Connection management
- Transaction management
- Allow object model and data model to be optimized independently

Mapping Objects to RDBMS



- Many ways to “map” (describe the relationship) between an object model and a schema
 - Map directly to a column
 - Relationships between objects can be mimicked with foreign key relationships in schema
 - Collections of objects can be:
 - One-to-many
 - Many-to-many with a join table
- Schema and classes are not tied together

Most good O/R Solutions Have Tools To Help with Mappings

Kodo Development Workbench

File MetaData Schema Visualization Help Query

Explorer

MetaData Schema

MetaData

samples.ide

Passenger

- luggage
- name
- price

Editor: Visualization

Passenger

base

Version: version-number

Class Ind: in-class-name

luggage: Set<String>

name: String

price: double

collection

PASS_DATA

FULL_NAME (VARCHAR)

JDOCLASS (VARCHAR)

JDOID (BIGINT)

JDOVERSION (INTEGER)

PRICE (DOUBLE)

PASS_LUGGAGE

ELBMENT (VARCHAR)

JDOID (BIGINT)

Visualization JDOQL

Log

```
TRACE -- Generating metadata for type "class samples.ide.Passenger".
TRACE -- Using reflection for metadata generation.
INFO -- Parsing metadata resource "file:/D:/KOD0-J~1.0B2/samples/ide/package.jdo".
TRACE -- Parsing class "Passenger".
```

MetaData KodoWorkbench JDBC Schema Tool Enhance

Ready.

Benefits of these specs



There are already a number of persistence solutions out there. What does JDO and EJB3 UPS bring to the table?

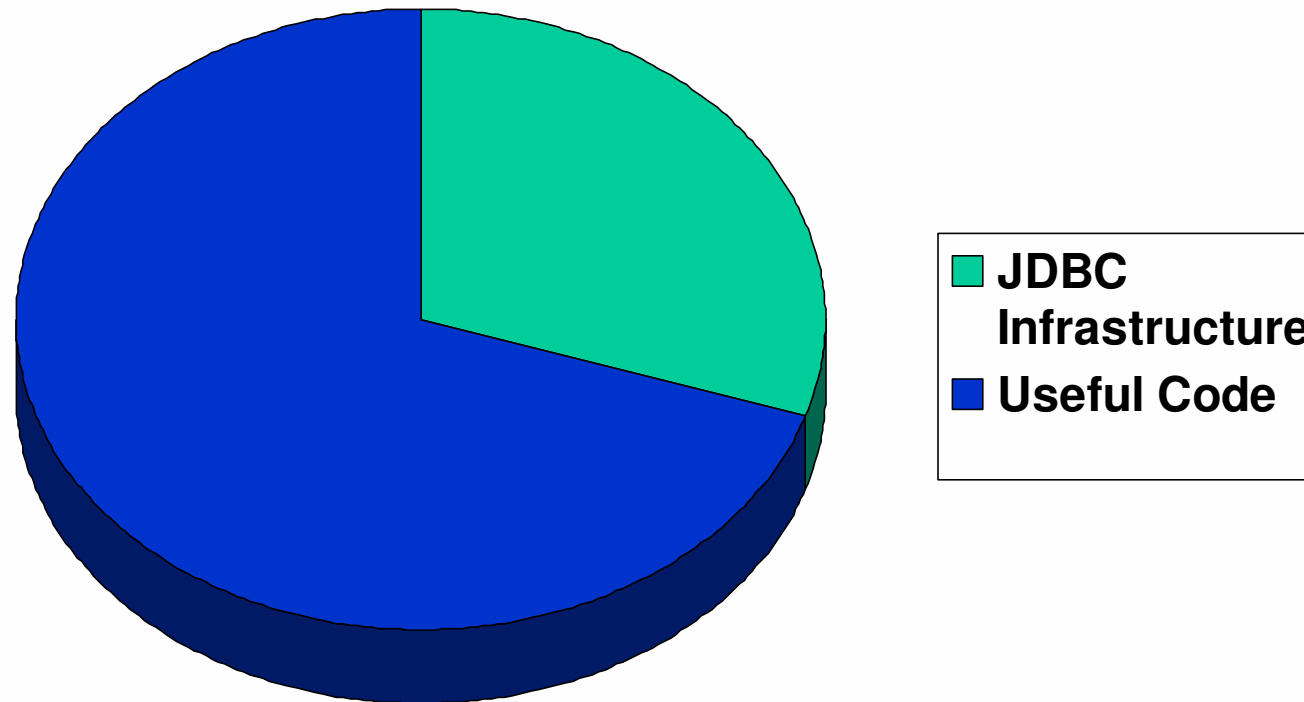
- Defer decisions
- Focus on core competencies
- Use Java to its fullest
- Project maintainability

Defer data store decisions



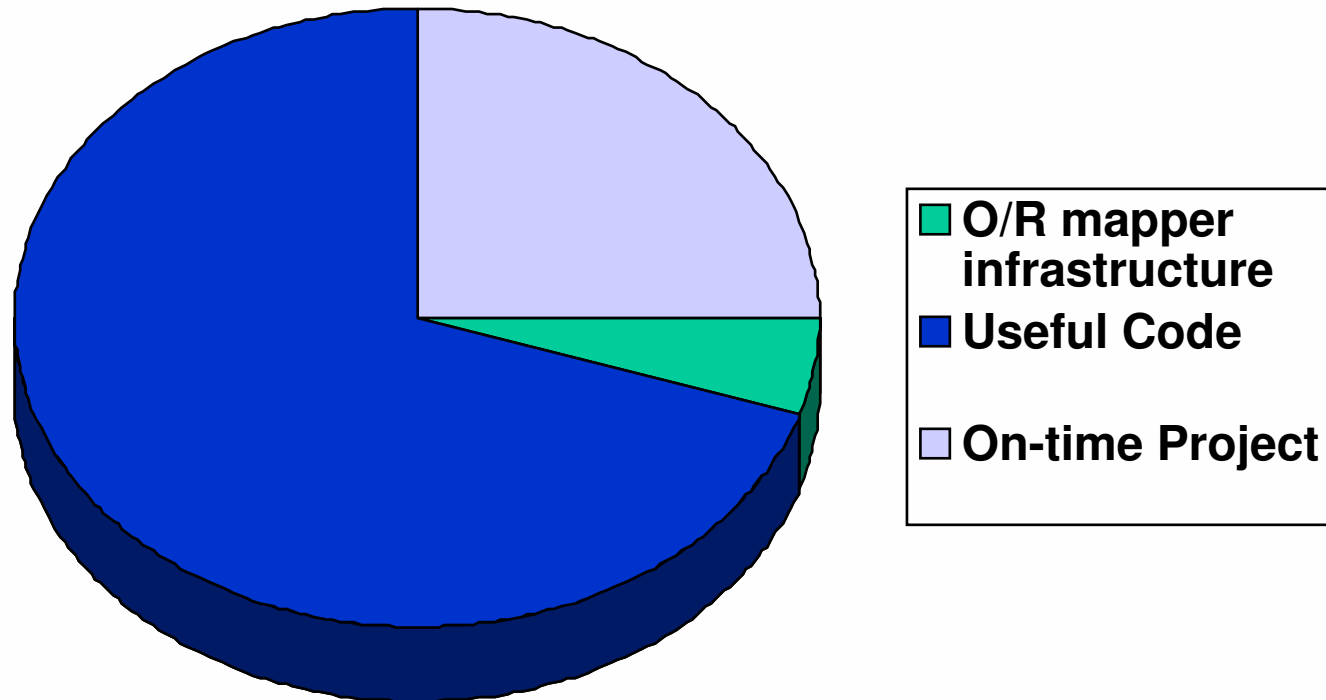
- Choice of persistence technology should not require any other architecture decisions
- Entity beans prior to EJB3: require the use of a container
- JDBC and proprietary object-relational mapping: require a relational database
 - often, results in dependencies to a particular SQL variant
- Changes to schema would cause pain and large amounts of code rewriting
- JDO and EJB3: seamlessly work with J2EE, but do not require it. Can be used with all sorts of data stores (relational, legacy EISs, J2ME devices, ...)

Focus on core competencies



- 20% to 40% of typical JDBC projects is persistence infrastructure
- JDBC code is repetitive and difficult, and therefore prone to errors

Focus on core competencies



- O/R mapper's infrastructure is typically closer to 5%
- No boilerplate code, so less room for cut-and-paste errors

Use Java to its fullest



O/R mapping is designed to integrate seamlessly with Java

- Polymorphism, both in queries and relations
 - Show me all the Vehicles in the warehouse
- Data encapsulation supported

Developers need not re-learn new rules and limitations.
Regular Java concepts behave exactly as expected.

Maintenance over time



- Most interaction happens through object model
- Direct use of APIs is minimal, so is readily understandable
 - query for objects
 - transaction demarcation (when used outside a JTA environment)
 - adding objects to database

Why O/R mapping?



To sum up, an O/R mapping framework enables you to:

- Make architecture and data store decisions as needed
- Do more “real work” in a given amount of time
- Do the same amount of “real work” in less time
- Write regular Java, without learning all sorts of restrictions
- More easily maintain your code base through its lifecycle

JDO 2 Status / Schedule

JDO 2 is an Evolution from JDO 1, Not a Revolution



- Maintain Java Data Objects 1.0 Compatibility
- Standardize Mapping to Relational DB
 - SQL as a supported query language
- Multi-tier Development Improvements
 - Disconnected Object Graphs
- Usability Improvements
- Better Object Modelling
- Richer Queries
 - Single-line format
 - Projections
 - Aggregates
 - Paging Query Results
- More Vendor Support
 - Remove PersistenceCapable requirement

JDO 2.0 schedule



- Public Draft currently available
- TCK, RI being developed in the Apache project
- JDO2 jars should be available in the next few weeks
- Final JDO2 specification (including TCK and RI) in four to eight months

Unnamed Persistence Specification (UPS) Status

Enterprise Java Beans 3

The Announcement



- <http://java.sun.com/j2ee/letter/persistence.html>

Historical background

- Core of EJB and JDO specs have included persistence
- Data persistence models have differed
- Sun is leading a community effort to have a single POJO persistence model

The Announcement



- <http://java.sun.com/j2ee/letter/persistence.html>

Summary of announcement

- Added 6 JDO members to JSR-220 (EJB3 Spec Team)
- All current JDO experts are listeners on JSR-220 list
- Work will start under JSR-220 but independent of EJB3
- Goal: Single O/R Mapping Framework for J2EE, J2SE
- Timeframe: J2EE 5.0 (January 2006 currently)

EJB 3



- Unnamed Persistence Specification (UPS) is part of EJB3
- Work in progress on a second Early Release Draft
- Heavy utilization of annotations
- Lots of mis-information out there
 - No existing product is UPS

Criteria to Consider When Choosing a Persistence Solution

Looking for a Persistence Solution? (1 of 3)



Criteria to Consider

- Standards based vs. Proprietary
 - Portability – avoid vendor lock-in
 - Specification created by experienced community members vs. 1 individual or 1 company with other goals
- Which standard?
 - JDO (27+ implementations)
 - EJB3
 - Both
- Total Cost of Ownership
 - Upfront Costs vs. Runtime Costs
 - Support Costs
 - Maintenance Costs
 - Training Costs

Looking for a Persistence Solution? (2 of 3)



Criteria to Consider

- Ease of Use
 - Tooling
- Ability to Optimize Scalability & Performance Trade-Offs
- Supported Mappings
 - Custom Mappings Available?
- Extensibility
- Datastore Flexibility
 - Relational
 - Non-relational (legacy, hierarchical, object, etc.)
- Architecture flexibility
 - In container – session beans, CMP, BMP
 - Outside container – servlets, Spring, JSP, client-server

Looking for a Persistence Solution? (3 of 3)



Criteria to Consider

- Performance and Scalability Issues
 - Lazy Loading
 - Dirty Field Tracking
 - Caching (pluggable caching with variety of invalidation strategies)
 - Minimize Round-Tripping
- Legacy Database Support
 - Stored Procedure Support
 - SQL Support
- Supporting Organization
 - Tied to another product?
 - Organization's thought leadership
 - Support response times
- Legal Issues
 - Indemnification

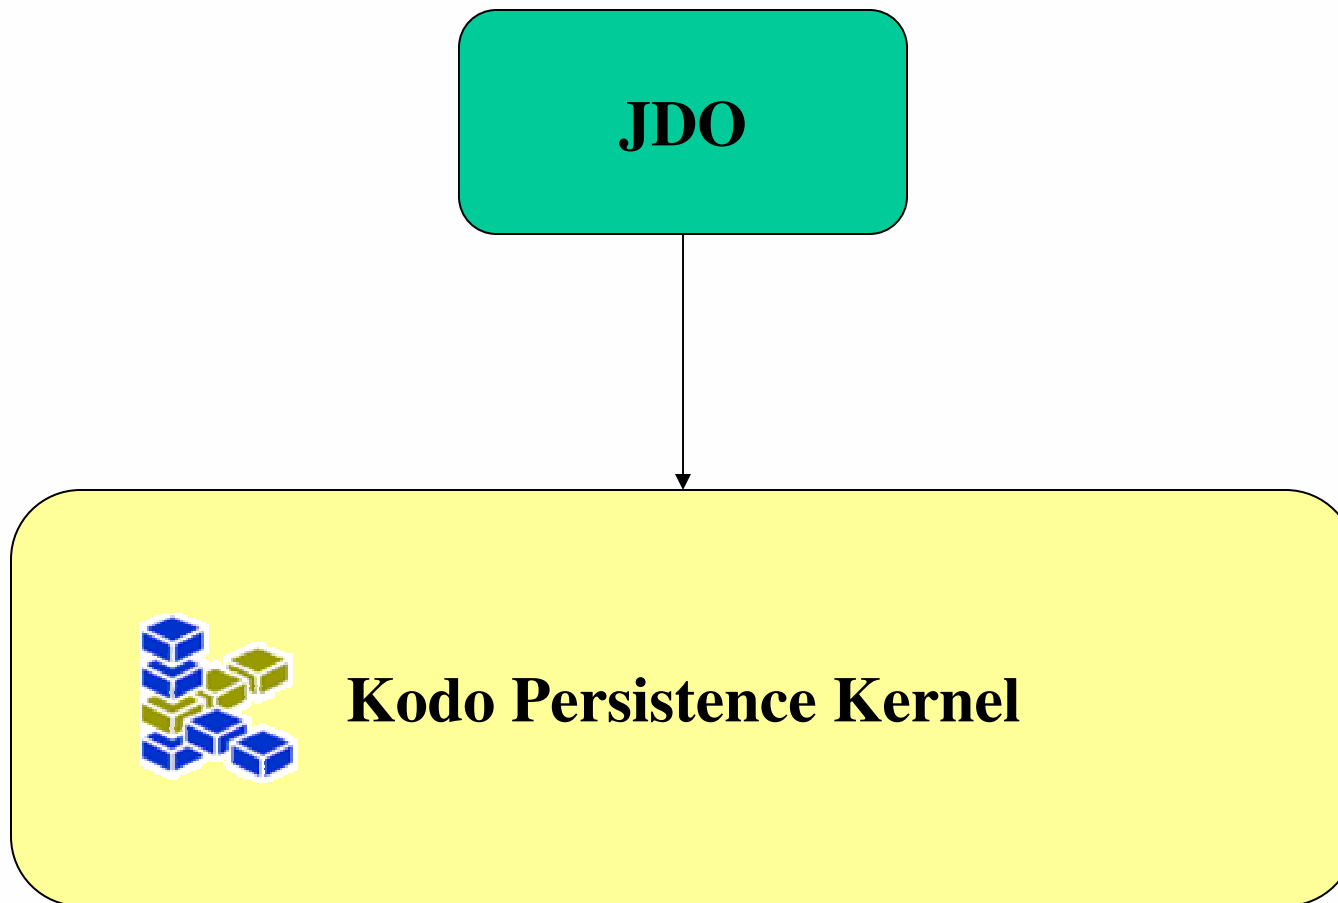
Kodo Product Suite

Kodo Roadmap

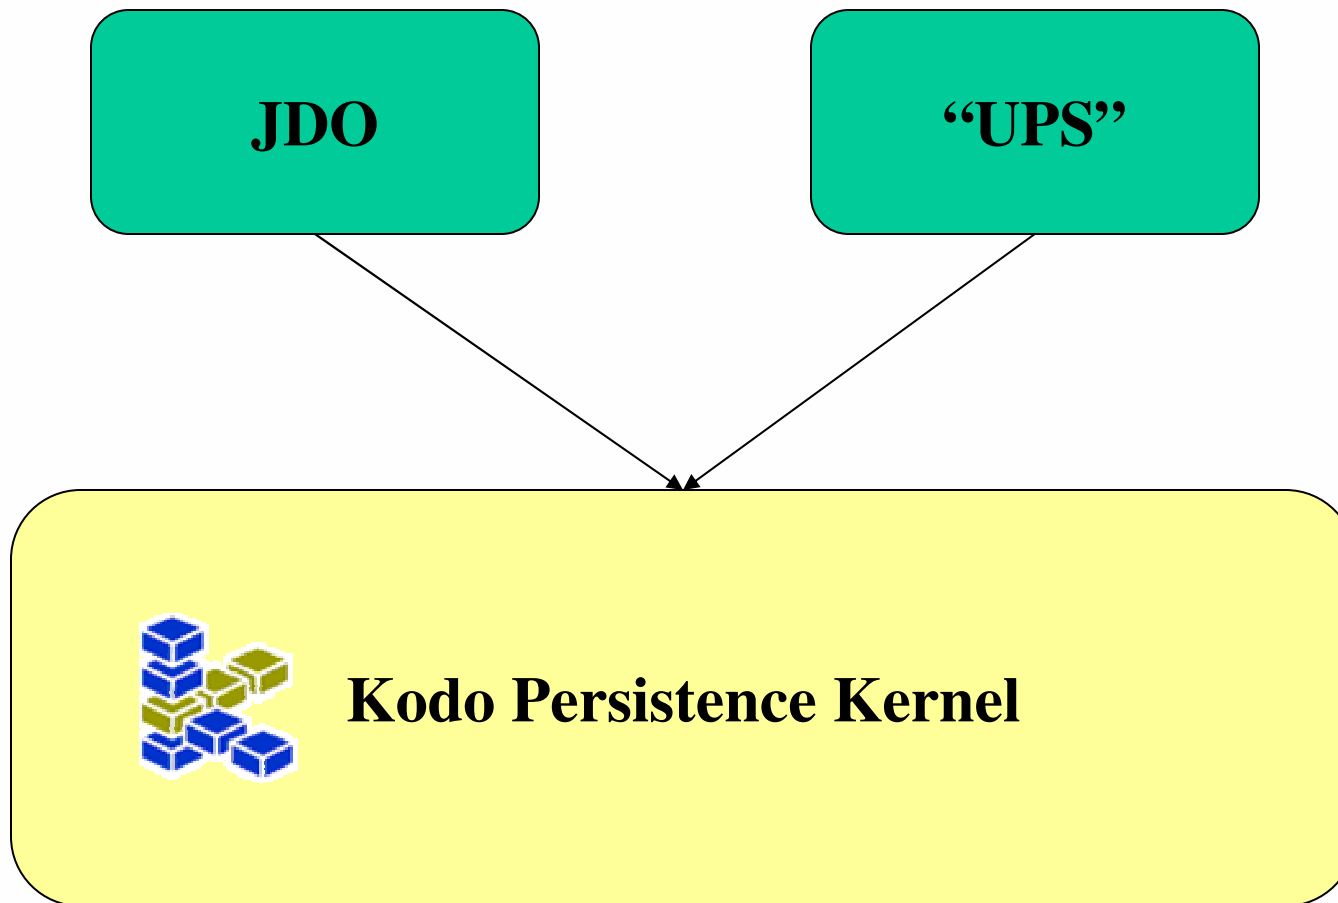


- Continue to build an underlying engine that promotes:
 - Performance
 - Scalability
 - Flexibility
 - Ease of Use
- Will offer multiple bindings to the underlying Kodo engine
 - Kodo JDO – JDO bindings
 - Kodo “UPS” – “UPS” (Unnamed Persistence Spec) bindings
- Interoperability between Kodo JDO and Kodo “UPS”
 - Will be possible to simultaneously use both APIs with same domain model

Current Situation – Kodo JDO



Future Situation – Kodo JDO and Kodo “UPS”



Kodo JDO 3.2



- Significant support for JDO 2 including:
 - JDO 2 Queries (single string JDOQL queries, implicit parameters / variables, named query support, subqueries)
 - More flexibility with detachment APIs, including automatic detachment on PM close
- Improvements to Kodo Development Workbench, Management Console, and Profiler
- Reverse Mapping Tool is accessible via a guided wizard
- Improved eager fetching
- Support for managed inverses
- Intersystems Caché database support



- Tools – Workbench and Management Console
- High-performance database cache
- Prepared statements, statement caching, statement batching
- Intelligent handling of large result sets
- SQL as query language
- flat, vertical, horizontal inheritance mappings
- Extensible architecture
- Reverse and forward engineering of database schema
- Schema evolution

Kodo Development Workbench



Kodo Development Workbench

File Metadata Schema Visualization Help Query

Explorer

Metadata Schema

samples.ide

Passenger

- luggage
- name
- price

Editor: Visualization

Passenger

base

Version: version-number

Class Ind: in-class-name

luggage: Set<String>

name: String

price: double

collection

PASS_DATA

FULL_NAME (VARCHAR)

JDOCLASS (VARCHAR)

JDOID (BIGINT)

JDOVERSION (INTEGER)

PRICE (DOUBLE)

PASS_LUGGAGE

ELEMENT (VARCHAR)

JDOID (BIGINT)

Visualization JDOQL

Log

```
TRACE -- Generating metadata for type "class samples.ide.Passenger".
TRACE -- Using reflection for metadata generation.
INFO -- Parsing metadata resource "file:/D:/KOD0-J-1.0B2/samples/ide/package.jdo".
TRACE -- Parsing class "Passenger".
```

Metadata KodoWorkbench JDBC Schema Tool Enhance

Ready.

Kodo JDOQL Editor



Kodo Development Workbench

File MetaData Schema Visualization Help Query

Explorer

MetaData Schema

MetaData

- samples.ide
 - Passenger
 - luggage
 - name
 - price

Editor: JDOQL

Candidates

samples.ide.Passenger

☒ Include Subclasses

Filter

Filter Language: javax.jdo.query.JDOQL

price>100

Ordering

field	direction

Variables

Parameters

Aggregates and Projections

Grouping

Fetch Configuration

Query range

Imports

Query Results: Passenger

2 records found. Execution time: 0 milliseconds.

	name	price
Passenger: 2	Patrick Linskey	135.12
Passenger: 4	Steve Kim	187.12

```
SELECT t0.JDOID, t0.JDOCLASS, t0.JDOVERSION,
       t0.FULL_NAME, t0.PRICE
FROM PASS_DATA t0
WHERE (t0.PRICE > 100)
```

☐ open results in new tabs

Visualization JDOQL

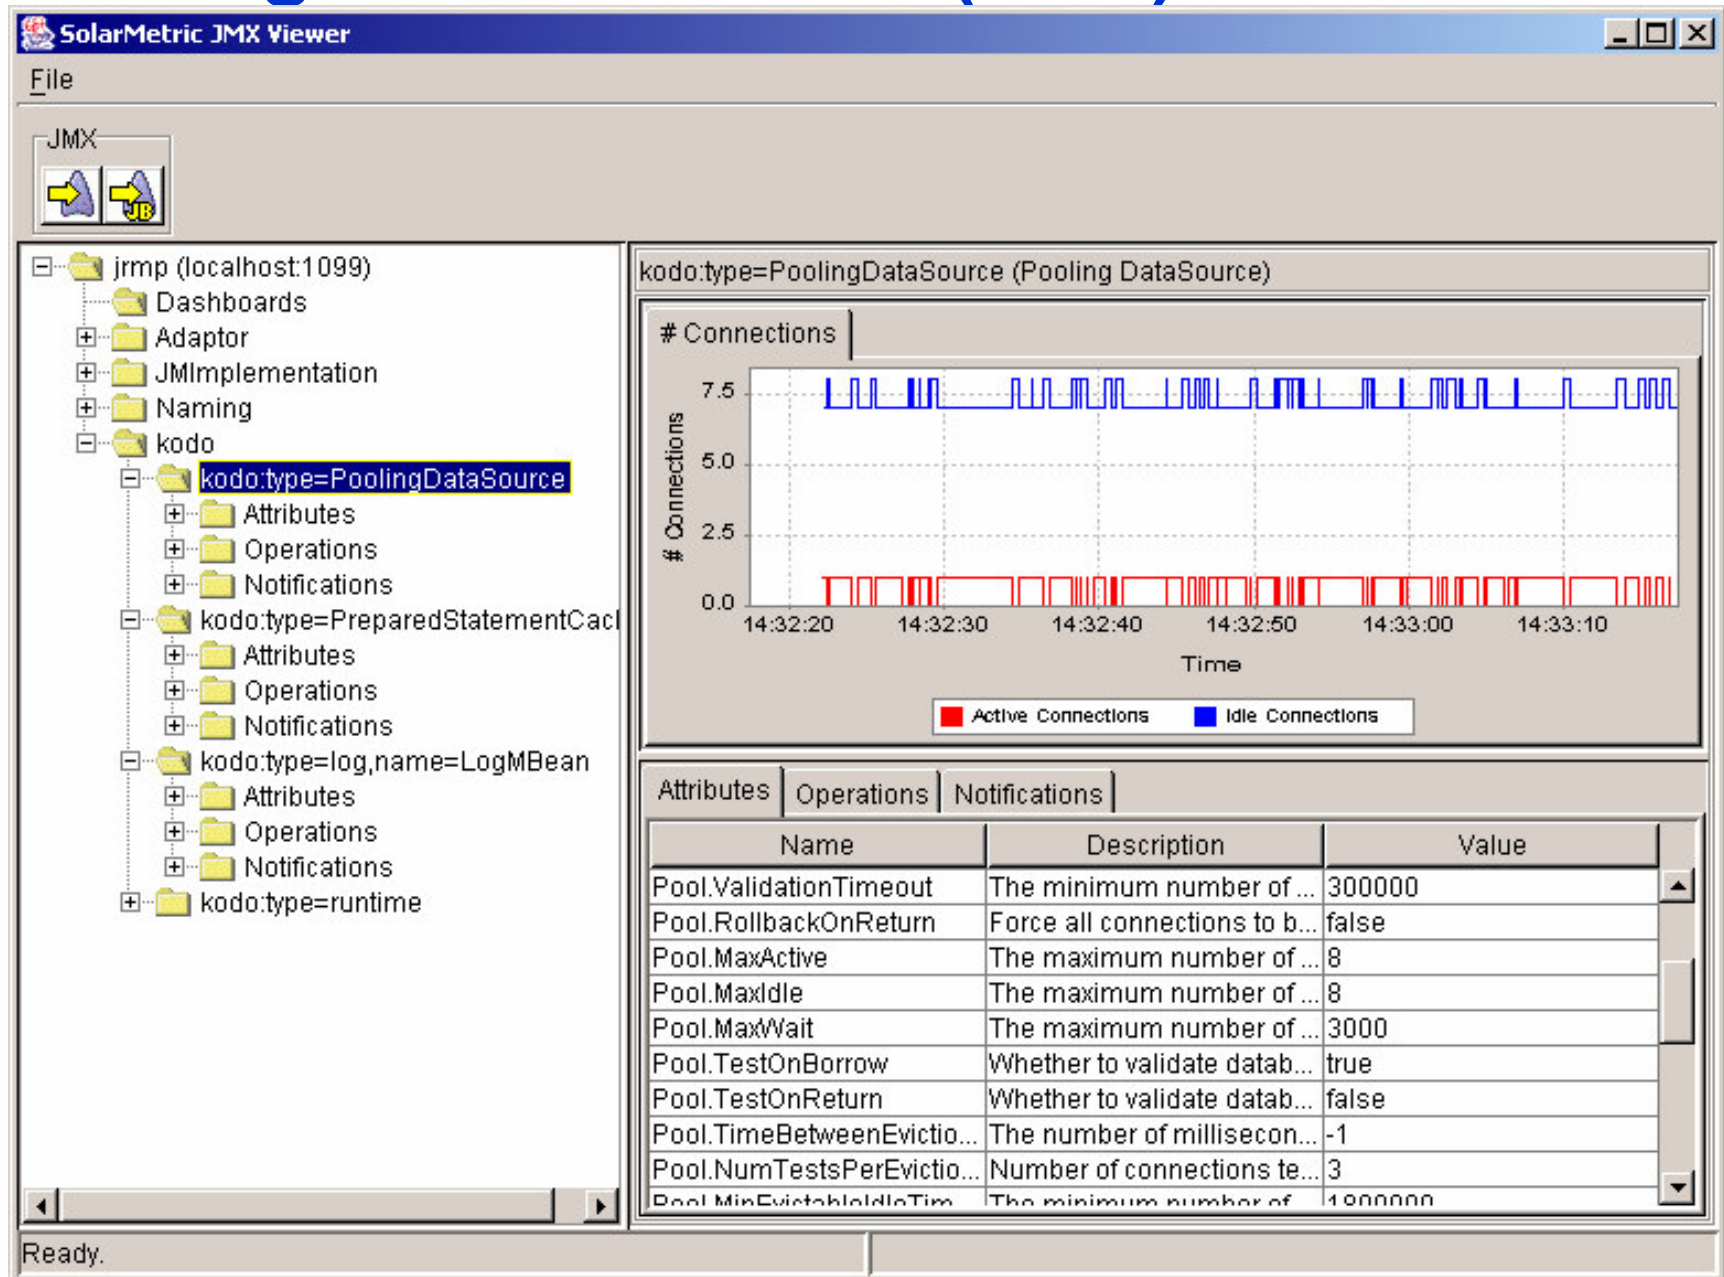
Log

```
INFO -- Parsing metadata resource "file:/D:/KOD0-J-1.0B2/samples/ide/package.jdo".
TRACE -- Parsing class "Passenger".
TRACE -- Generating metadata for type "class samples.ide.Passenger".
```

MetaData KodoWorkbench JDBC Schema Tool Enhance

Ready.

Management Console (JMX)



Profiler



Kodo Profiling Console

File

Profiling

fast slow never

SQL

Profiling Agent

- kodo.profile.KodoProfilingAgentImpl@31f2a7
 - PersistenceManager: 0 ms
 - PersistenceManager: 0 ms
 - PM.newQuery (Class cls, String filter): 210 ms
 - PM.newQuery (): 210 ms
 - Query.setClass (Class pcClass): 0 ms
 - Query.setFilter (String filter): 0 ms
 - Query.declareParameters (String parameters): 0 ms
 - Query.execute (...): 96,126 ms
 - query.close (Object queryResult): 70 ms
 - loadField: 420 ms
 - loadField: 1,280 ms
 - loadField: 350 ms

Node: PersistenceManager

samples.management.QueryThreadSinglePM.<init> line: 36

Field Name	Total	Fetchd (% Used)	Unfetchd (% U...
PersistenceManager			
samples.management.Dog			
name	51	51 (100%)	0 (-%)
price	51	0 (-%)	51 (106%)
collar	51	0 (-%)	51 (108%)
toyName	51	0 (-%)	51 (108%)
samples.management.Collar			
dog			

Consider putting this field in the initial fetch configuration.

PersistenceManager

- samples.management.Dog: name == dogname
 - executed: 100
 - avg. accessed: 0.8 (80 / 100)
 - avg. % accessed: 100.0% (100 samples)
 - avg. size: 0.8 (100 samples)

Statistics

Name	Value
Sum	0ms
Average	0ms
Minimum	0ms

Ready.

Supported Databases



Relational:

- Oracle 8i/9i/10
- Microsoft SQL Server
- IBM DB2
- Sybase Adaptive Server Enterprise
- Informix IDS
- MySQL
- PostgreSQL
- Firebird
- Interbase
- Hypersonic SQL
- JDataStore
- FoxPro
- MS Access
- Pointbase
- Cloudscape
- InstantDB
- Empress
- **Extensible database support**

Non-Relational:

- Legacy EIS (CICS, Unisys, others) (ask for details)
- LDAP (coming soon)
- BerkeleyDB (coming soon)
- InterSystems Caché

Kodo JDO IDE Support



- All IDEs with Ant Support
- Tight integration with:
 - Borland JBuilder 7 or newer
 - Sun One Studio / NetBeans
 - Eclipse / WSAD / WSED
- Standalone Development Workbench

Kodo JDO Supported App Servers

- WebLogic 6.2 – 8.1
- WebSphere 5
- JBoss 3.0 – 3.2
- SunONE
- Trifork
- JRun 4
- Borland Enterprise Server

More Info



- Web site: <http://www.solarmetric.com>
- On-site and off-site Kodo JDO training available
- Professional Services available
- Email contact: info@solarmetric.com
- Speaker
Patrick Linskey: pcl@solarmetric.com
- Advanced Webinar – December 9 – 10:00 AM Pacific
 - Covers Performance and Scalability
 - Using JDO with J2EE
 - Other Advanced Topics
 - Register at <http://www.solarmetric.com>

Backup Slides

What's New In JDO 2?

Goals for Java Data Objects 2.0



- Maintain JDO 1.0 Compatibility
- Standardize Mapping to Relational DB
 - SQL as a supported query language
- Multi-tier Development Improvements
 - Disconnected Object Graphs
- Usability Improvements
- Better Object Modelling
- Richer Queries
 - Single-line format
 - Projections
 - Aggregates
 - Paging Query Results
- More Vendor Support
 - Remove PersistenceCapable requirement

Standardize Mapping to RDBMS

- Simple Mapping
 - Class \Leftrightarrow Table
 - Field \Leftrightarrow Column
 - 1-1 relationship \Leftrightarrow unique foreign key
 - 1-many relationship \Leftrightarrow foreign key
 - many-many relationship \Leftrightarrow join table
- Complex Mapping
 - Inheritance strategies
 - Multiple tables per class
 - List, Map, Embedded

Multi-tier Development



- Disconnected Object Graphs
 - Disconnect objects from PersistenceManager
 - Close PersistenceManager
 - Modify disconnected objects
 - send to different tier, or
 - apply changes directly to objects
 - Begin new transaction
 - Reconnect objects to PersistenceManager
 - Commit transaction (applies changes)
- Disconnected Objects keep original identity and version

Usability (Bootstrap)



- How do you find a PersistenceManagerFactory?
- Java Data Objects 1.0 Properties:

```
javax.jdo.PersistenceManagerFactoryClass:\
com.sun.jdori.fostore.FOStorePMF
javax.jdo.option.ConnectionURL:\
fostore:/shared/databases/jdo/dbdir
javax.jdo.option.ConnectionUserName:craig
javax.jdo.option.ConnectionPassword:faster
javax.jdo.option.Optimistic:true
javax.jdo.option.RetainValues:true
javax.jdo.option.NontransactionalRead:true
```


Bootstrap



- JDO 1.0 Properties instance:
 - getPersistenceManagerFactory (Properties props)
- JDO 2.0 Properties loaded from:
 - stream
 - getPersistenceManagerFactory (InputStream props)
 - File (use to get a stream)
 - getPersistenceManagerFactory (File propsFile)
 - resource name (use to get a stream)
 - getPersistenceManagerFactory (String propsResource)
- JDO 2.0 JNDI lookup
 - getPersistenceManagerFactory (String name, Context c)

Usability



- JDO 2.0 Transaction adds:
 - boolean getRollbackOnly();
 - void setRollbackOnly();
- JDO 2.0 Query can be defined in metadata and accessed by name

Usability



- ThreadLocal JDOHelper methods
 - getThreadLocalPersistenceManager
 - setThreadLocalPersistenceManager
 - getThreadLocalPersistenceManagerFactory
 - SetThreadLocalPersistenceManagerFactory
- Allows business delegates to be independent of their environment
 - Web server
 - App server
 - Two tier

Better Object Modeling



- Persistent Abstract Class Support
 - map abstract classes to tables
 - map properties, fields to columns
 - create new implementation instances
- Persistent Interface Support
 - map interfaces to tables
 - map properties to columns
 - create new implementation instances
 - use interfaces in queries

Richer Queries



- Projections
- Aggregates
- More String expressions
- More numeric, Map functions
- Paging query results
- Native SQL
- User-defined Result Class

Preserve Expressability in SQL

Projections



- JDO 1.0: query results are a subset of instances in the candidate collection
- JDO 2.0: query results can be projections of:
 - candidate collection instances;
 - variable instances;
 - fields;
 - parameters;
 - combinations of the above.

Projections: Example



```
Query q = pm.newQuery (Employee.class,  
    "dept.name.startsWith(deptName)");  
q.declareParameters ("String deptName");  
q.setResult("name, salary, boss");  
Collection names = (Collection)  
q.execute("R&D");
```

Bill Jones	12546	Employee@18d8788
Sam Adams	15948	Employee@18d8788
Will Clinton	50453	Employee@1867480
Westlake Clerk	18737	Employee@18720
Harvey Dean	14657	Employee@18720
Ed Muskrat	13009	Employee@18720

Aggregates



- Query Results can be aggregates
- Aggregates have standard semantics
 - min
 - max
 - avg
 - sum
 - count
- Group by projected fields
- “Having” allows conditional filtering

Aggregates: Example



```
Query q = pm.newQuery (Employee.class,  
    "dept.name.startsWith(deptName)");  
q.declareParameters ("String deptName");  
q.setResult("dept.name, min(salary), max(salary)");  
q.setGrouping("dept.name having count(dept.name) > 1");  
Collection names = (Collection) q.execute("R&D");  
for (Object[] i : (Collection<Object[]>) results) {  
    println(i[0], i[1], i[2]);  
}
```

R&D Santa Clara	12546	15948
R&D Burlington	13009	18737

String Expressions



- `toLowerCase()`, `toUpperCase()`
- `indexOf(String)`, `indexOf(String, int)`
- `matches(String pattern)`
 - pattern is a subset of regular expressions:
 - `(?i)` global case-insensitive
 - `.` match any one character
 - `.*` match any number (0 to n) of characters
 - pattern is literal or parameter only
- `substring(int)`, `substring(int, int)`

Other Query Filter Methods



- `Math.abs(numeric expression)`
- `Math.sqrt(numeric expression)`
- `Map.containsKey(Object)`
- `Map.containsValue(Object)`

Paging Query Results



- Improves performance for some applications
 - Skips already-returned results
 - Limits number of results
- `Query.setRange (int fromIncl, int toExcl);`
- default
 - `fromIncl = 0`
 - `toExcl = Integer.MAX_VALUE`

User-Defined Result Class



- JDO 1.0 query results are of type Collection
- JDO 2.0 allows user to specify result class:
 - Primitive wrapper (unique results)
 - Collection<Primitive wrapper>
 - Object[] (projected or aggregate unique results)
 - Collection<Object[]>
 - User-defined Class (unique results)
 - Collection<User-defined Class>

ResultClass: Example



```
class Info {  
    public String name;  
    public Float salary;  
    public Employee reportsTo;  
}
```

```
class Employee {  
    private String name;  
    private float salary;  
    private Department dept;  
    private Employee boss;  
}
```

```
Query q = pm.newQuery (Employee.class,  
    "dept.name == deptName");  
q.declareParameters ("String deptName");  
q.setResult("name, salary, boss as reportsTo");  
q.setResultClass(Info.class);  
Collection results = (Collection)  
q.execute("R&D");  
for (Info i : (Collection<Info>) results) {  
    println(i.name, i.salary, i.reportsTo.name);  
}
```

JDO Code Examples

Employee.java



```
package kodo.example;

import java.util.*;

public class Employee extends Person
{
    private float salary;
    private Company company;
    private Set projects = new HashSet();

    public Employee (String firstName,
                     String lastName) {
        super (firstName, lastName);
    }

    public void giveRaise (float percent) {
        salary *= 1 + percent;
    }

    public Collection getProjects () {
        return projects;
    }
}
```

```
<?xml version="1.0" encoding="UTF-8"?>
<jdo>
  <package name="kodo.example">
    <class name="Employee"
      persistence-capable-superclass="Person">

      <field name="projects">
        <collection element-type="Project" />
      </field>

    </class>
  </package>
</jdo>
```


JDO Usage Example



```
import javax.jdo.*;

public class MyPersistenceCode
{
    public static void main (String[] args)
    {
        // configure system
        PersistenceManagerFactory pmf =
            JDOHelper.getPersistenceManagerFactory (System.getProperties());
        PersistenceManager pm = pmf.getPersistenceManager();

        // business code
        Employee emp = new Employee ("Marc", "Prud'hommeaux");
        emp.setCompany (new Company ("SolarMetric, Inc.));
        emp.getProjects ().add (new Project ("Kodo"));
        emp.giveRaise (.10F);

        // persistence code
        pm.currentTransaction().begin();
        pm.makePersistent (emp);
        pm.currentTransaction ().commit ();
        pm.close ();
        pmf.close ();
    }
}
```

JDOQL Examples



- Basic Query:

```
String filter = "salary > 30000";  
Query q = pm.newQuery (Employee.class, filter);  
Collection emps = (Collection) q.execute ();
```

- Basic Query with Ordering:

```
String filter = "salary > 30000";  
Query q = pm.newQuery (Employee.class, filter);  
q.setOrdering ("salary ascending");  
Collection emps = (Collection) q.execute ();
```

JDOQL Examples



- Query with Relation Navigation and Parameters:

```
String params = "float min, float max";  
String filter = "company.revenue > min"  
    + " && company.revenue <= max";  
Query q = pm.newQuery (Employee.class, filter);  
q.declareParameters (params);  
Collection emps = (Collection) q.execute  
    (new Float (500000F), new Float (1000000F));
```

JDOQL Examples



- Query with Multi-value Navigation:

```
String vars = "Project p";  
String filter = "projects.contains (p)"  
    + " && p.name == \"Kodo\"";  
Query q = pm.newQuery (Employee.class, filter);  
q.declareVariables (vars);  
Collection emps = (Collection) q.execute ();
```

Native SQL



- JDO 1.0 does not recognize SQL
- JDO 2.0 supports:
 - obtaining `java.sql.Connection` from `PersistenceManager`
 - defining SQL queries
 - may improve performance (good)
 - makes query non-portable (bad)
 - may support SQL-specific constructs (?)
 - may allow SQL DBA more control (!)

More Vendor Support



- Remove PersistenceCapable requirement
 - Remove pre-processor, post-processor requirement
 - Allow non-enhancer versions of JDO implementation
 - All other compliance requirements remain