

Exploring Terracotta

Alex Miller

<http://terracotta.org>

<http://tech.puredanger.com>



Network Attached Memory

Open Source

Virtual Heap

Clustering

Persistent Heap

Grid

Scalability

Java Memory Model

Agenda

- What is it?
- How do I use it?
- What's it good for?
- Q&A

Clustered Heap

- Virtual
- Shared
- Persistent
- Large (1 TB)

Clustered Synchronization

- synchronized
- wait/notify
- java.util.concurrent
- Java Memory Model

Transparency

- No API, no JAR, just Java
- No serialization
- External declarative configuration

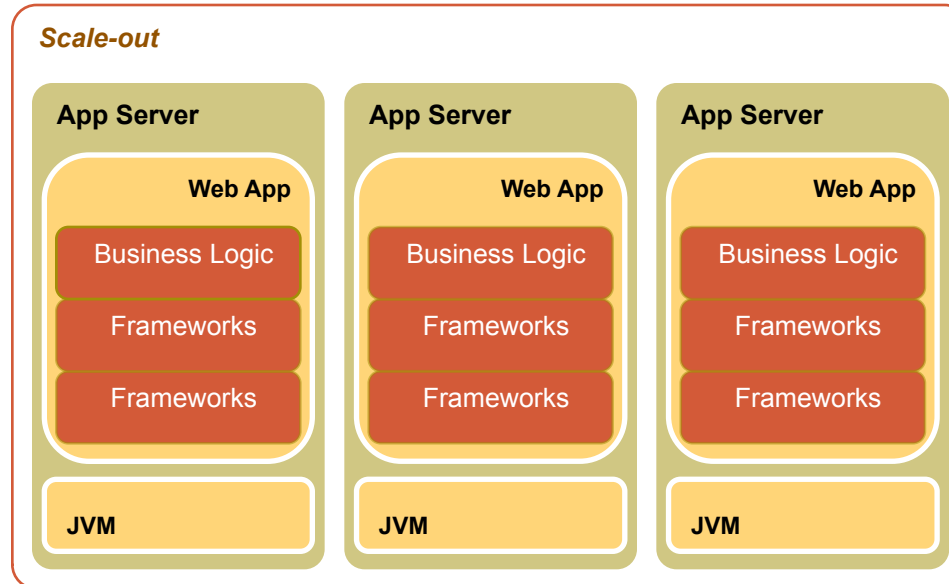
Enterprise-ready

- Scalability
- High-availability
- Monitoring
- Management

Terracotta approach

- Today's Reality

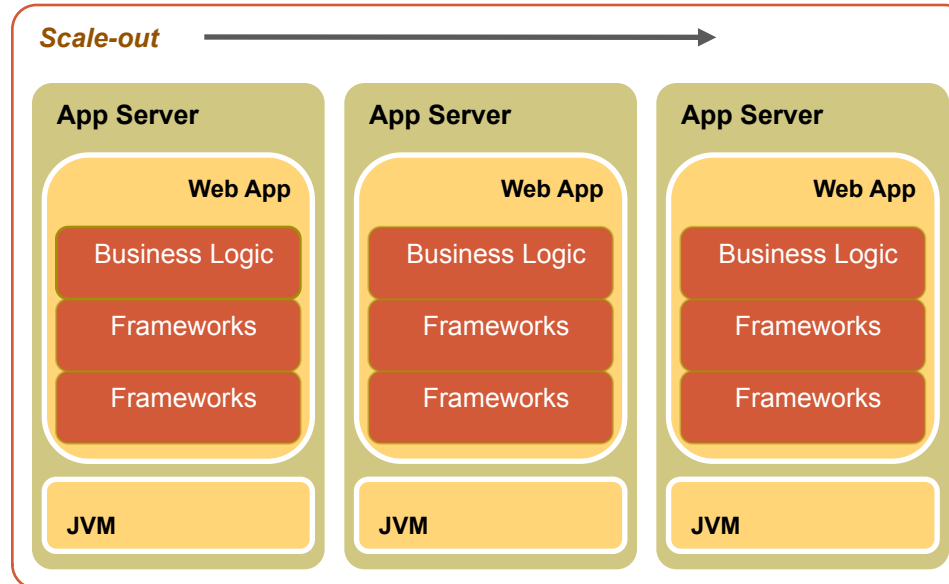
- Scale out is complex
- Requires custom Java code



Terracotta approach

● Today's Reality


- Scale out is complex
- Requires custom Java code



Terracotta approach

● Today's Reality

- Scale out is complex
- Requires custom Java code



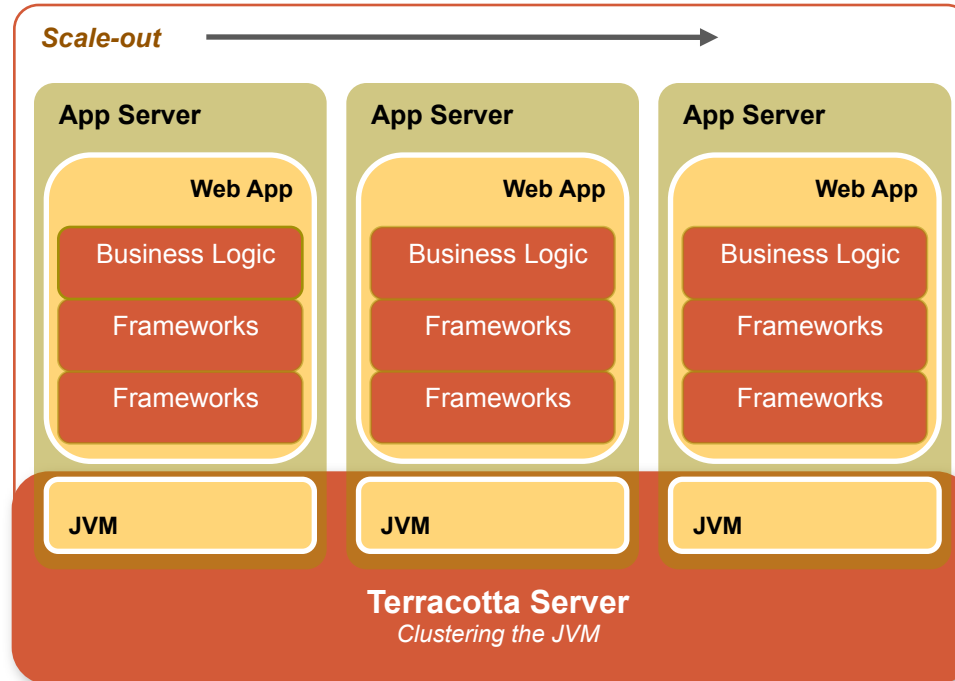
Terracotta approach

- Today's Reality

- Scale out is complex
- Requires custom Java code

- Our different approach

- Cluster the JVM
- Eliminate need for custom



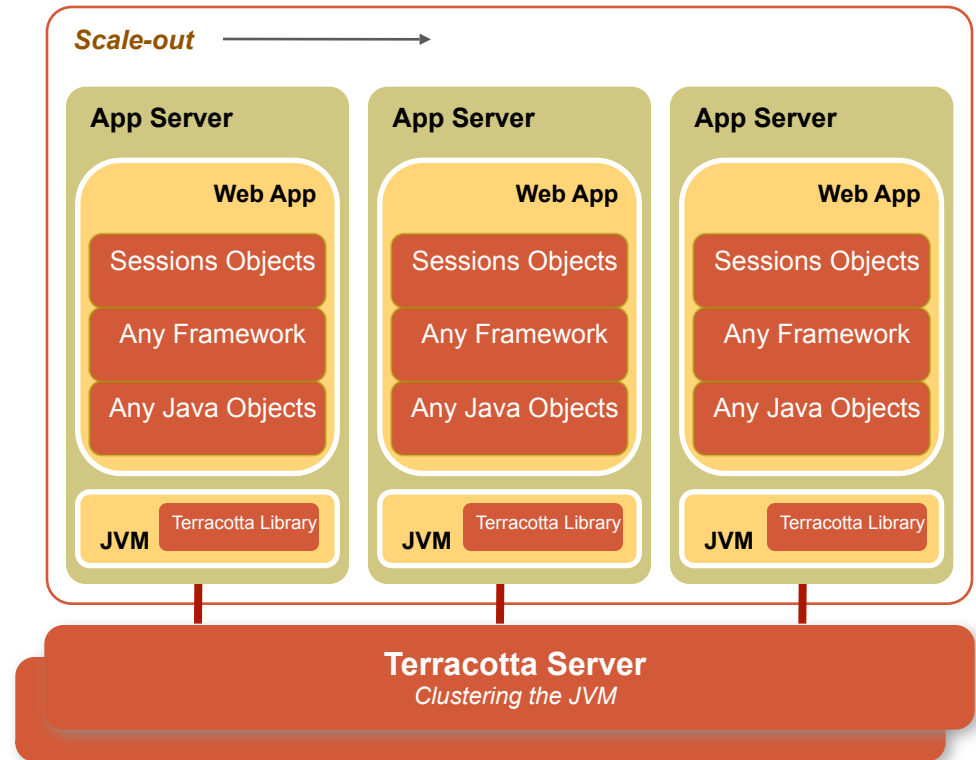
The architecture of Terracotta

- Local JVM Client

- Transparent
- Pure Java Libraries

- Terracotta Server

- 100% Pure Java
- HA Active / Passive Pair



The architecture of Terracotta

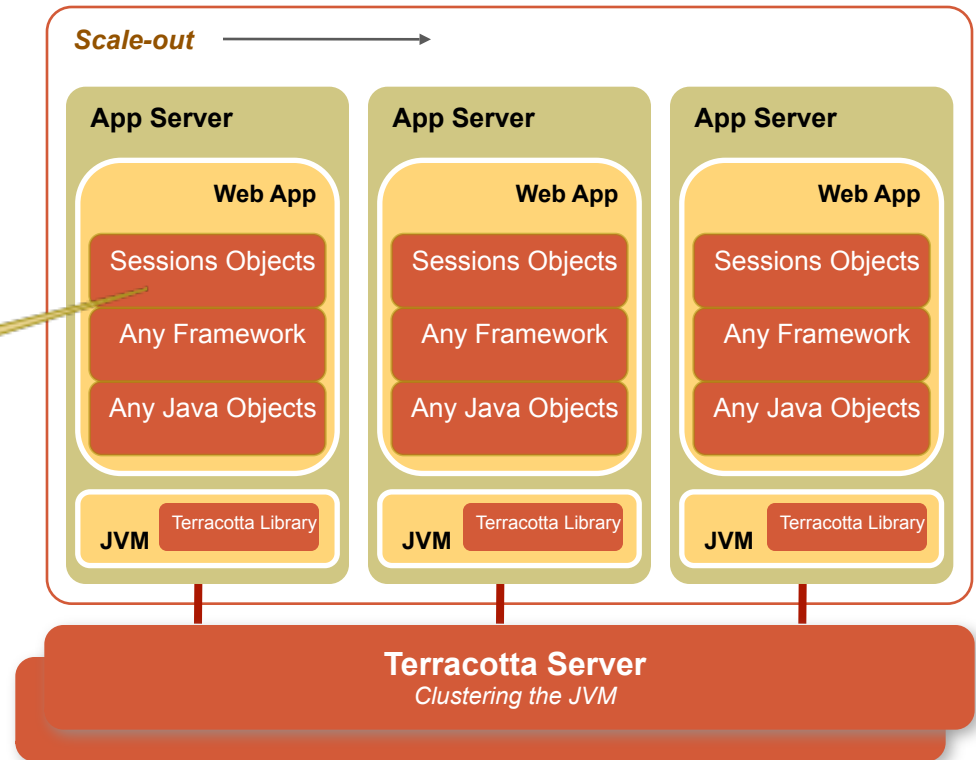
● Local JVM Client

- Transparent
- Pure Java Libraries

Local JVM Client

● Terracotta Server

- 100% Pure Java
- HA Active / Passive Pair



The architecture of Terracotta

● Local JVM Client

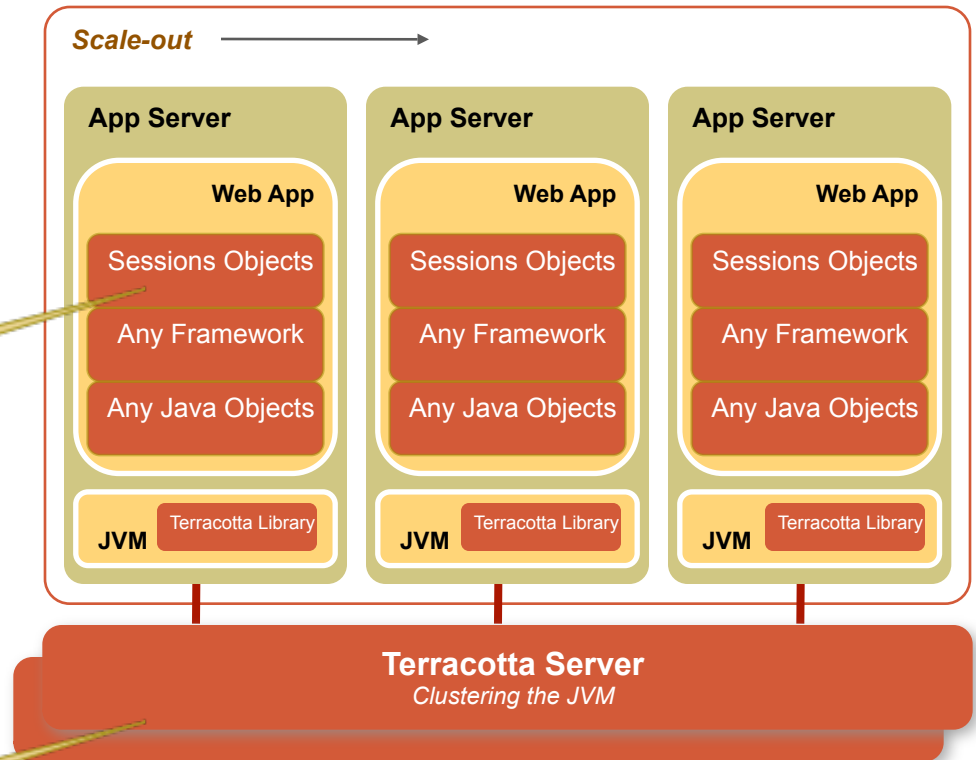
- Transparent
- Pure Java Libraries

Local JVM Client

● Terracotta Server

- 100% Pure Java
- HA Active / Passive Pair

Terracotta Server



The architecture of Terracotta

● Local JVM Client

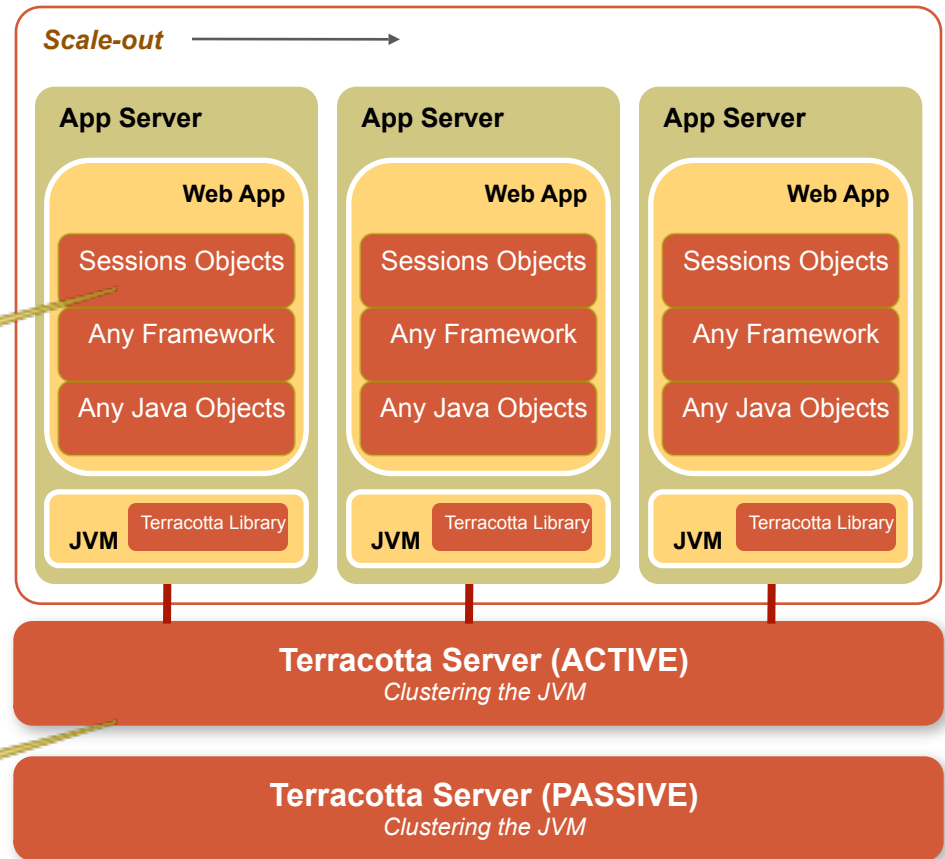
- Transparent
- Pure Java Libraries

Local JVM Client

● Terracotta Server

- 100% Pure Java
- HA Active / Passive Pair

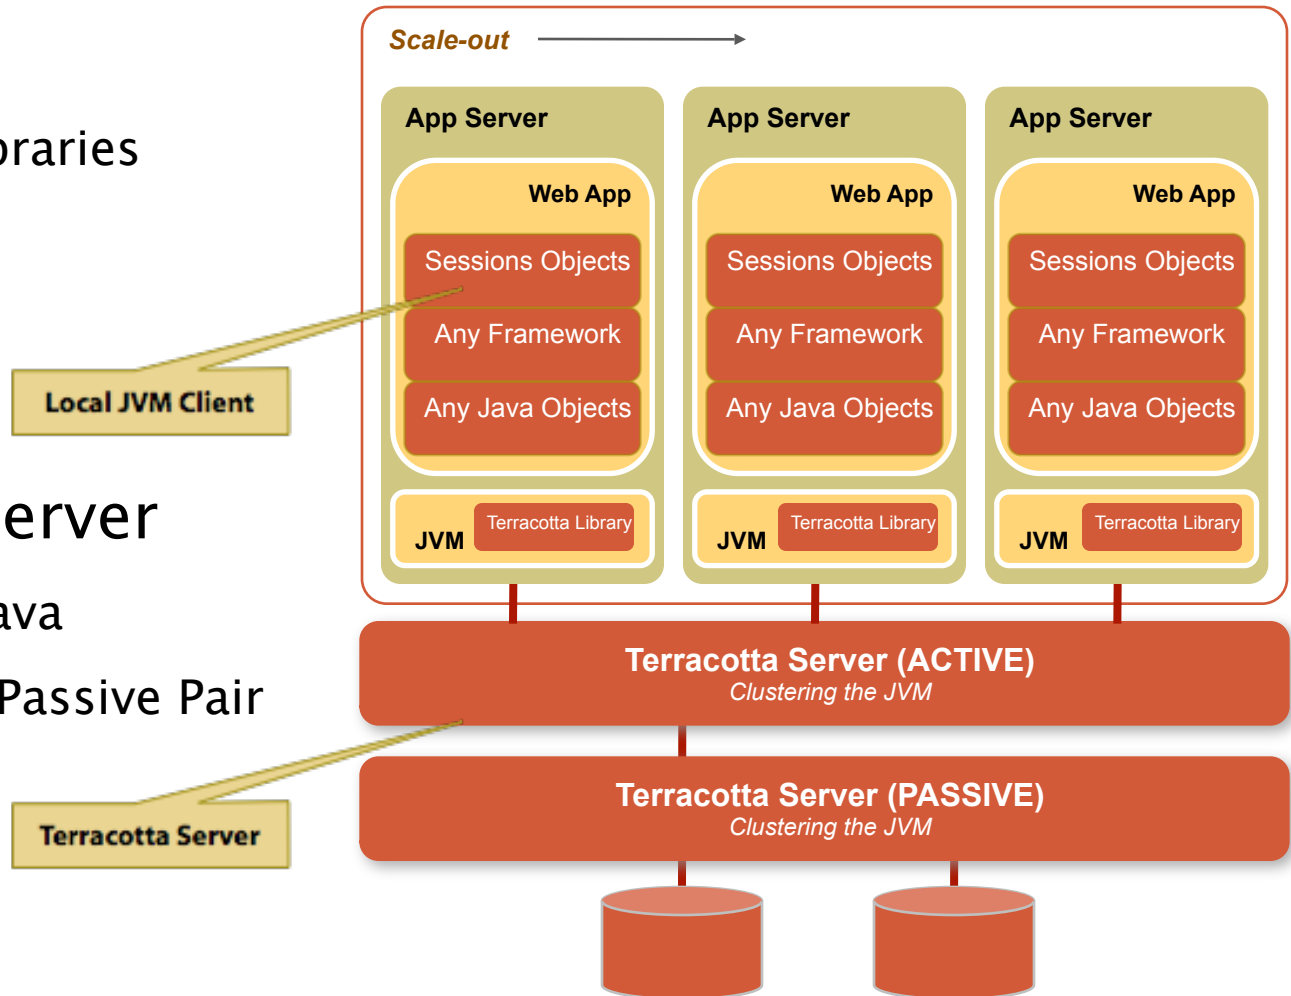
Terracotta Server



The architecture of Terracotta

● Local JVM Client

- Transparent
- Pure Java Libraries



● Terracotta Server

- 100% Pure Java
- HA Active / Passive Pair

Terracotta Server

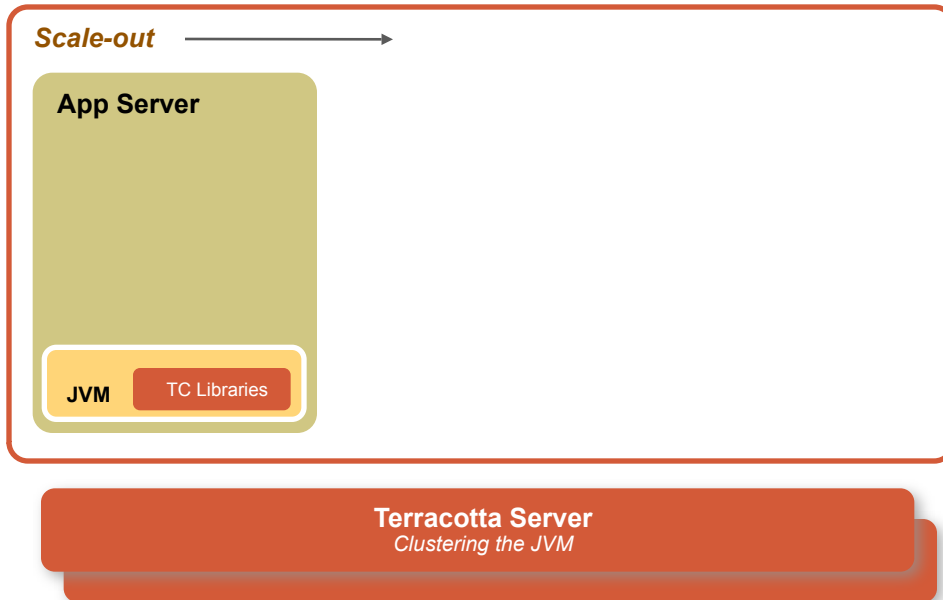
Network-Attached Memory

Scale-out →

Terracotta Server
Clustering the JVM

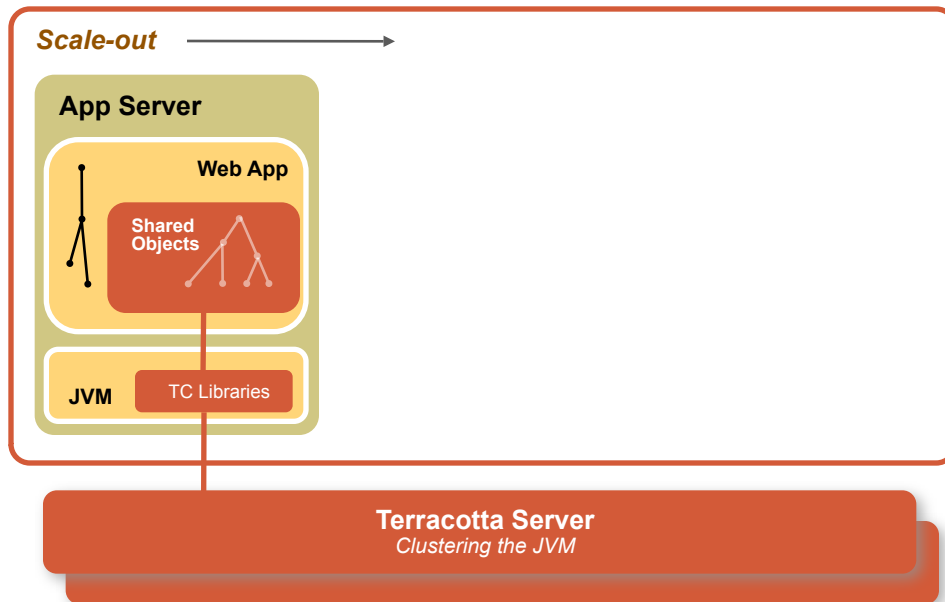
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



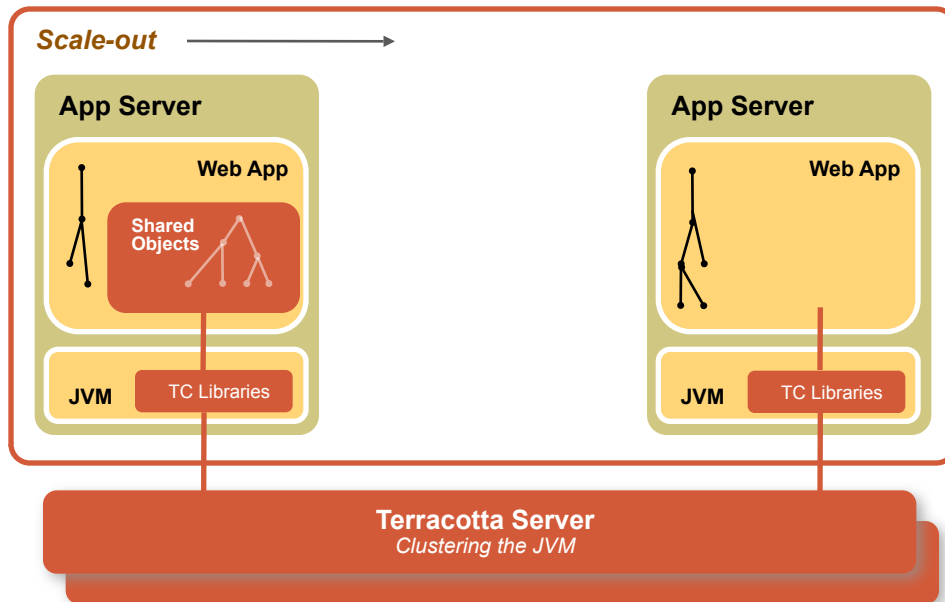
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



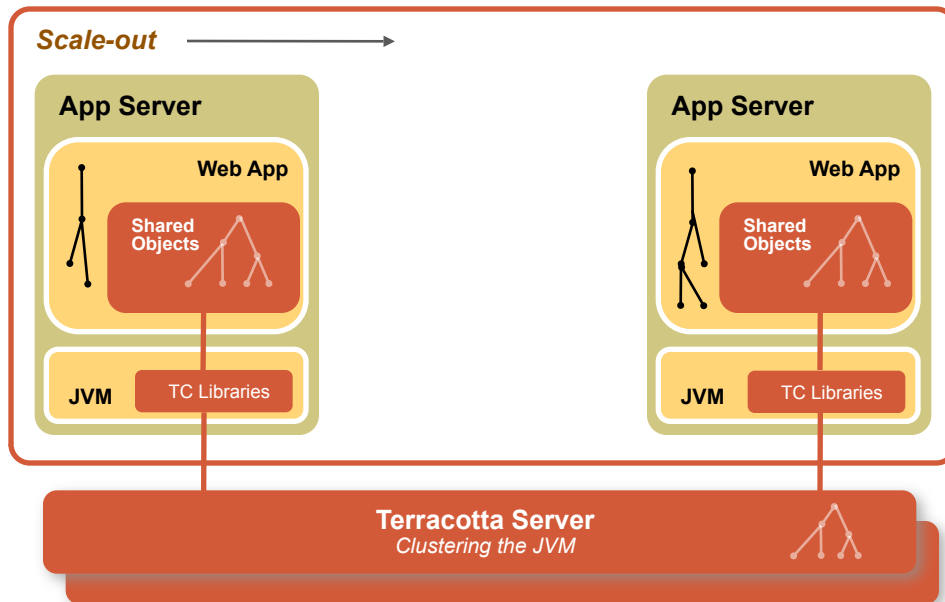
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



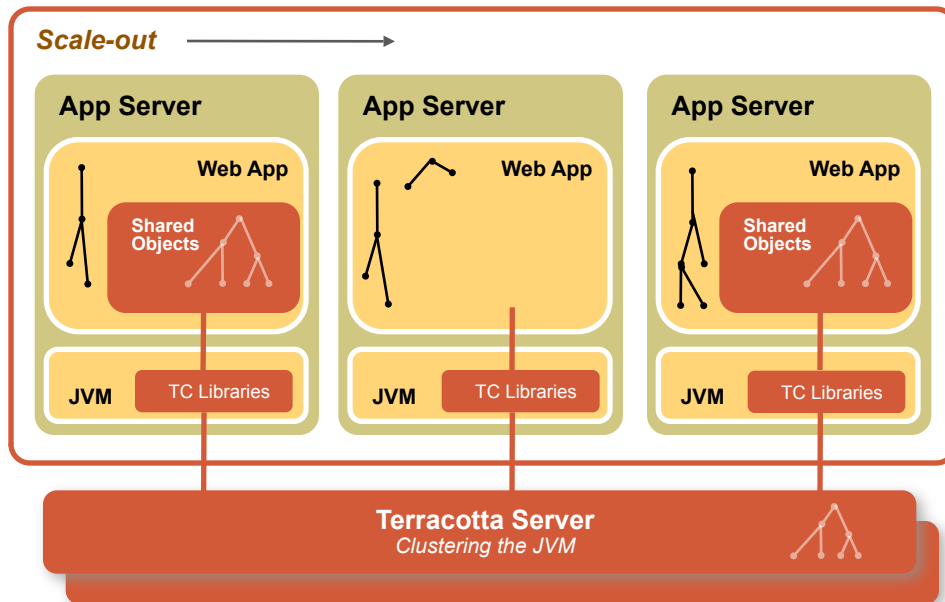
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



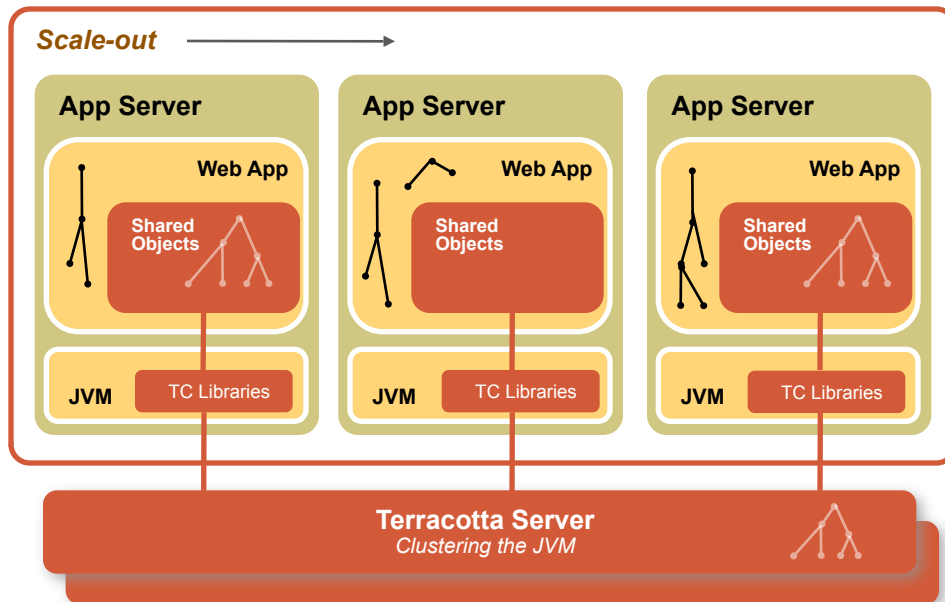
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



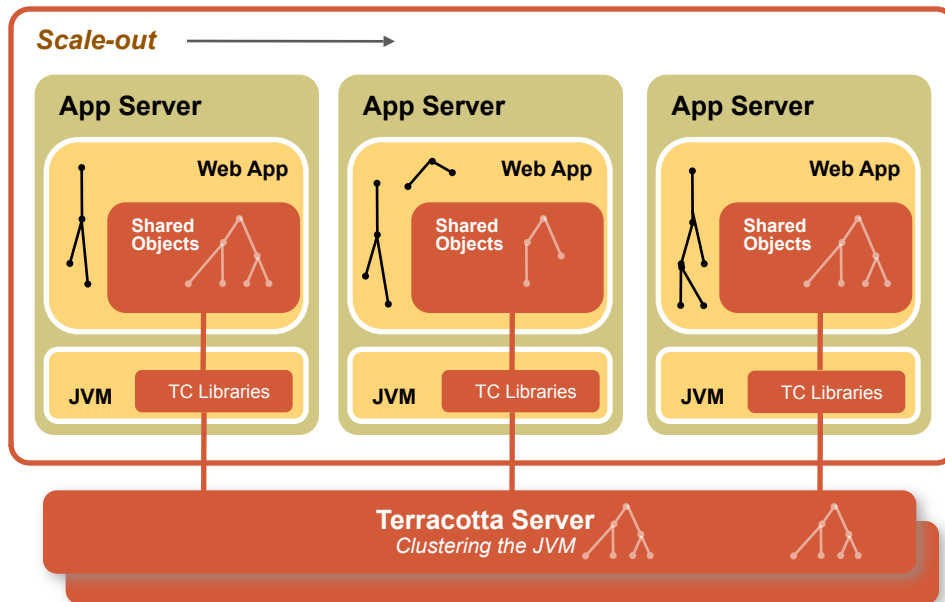
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



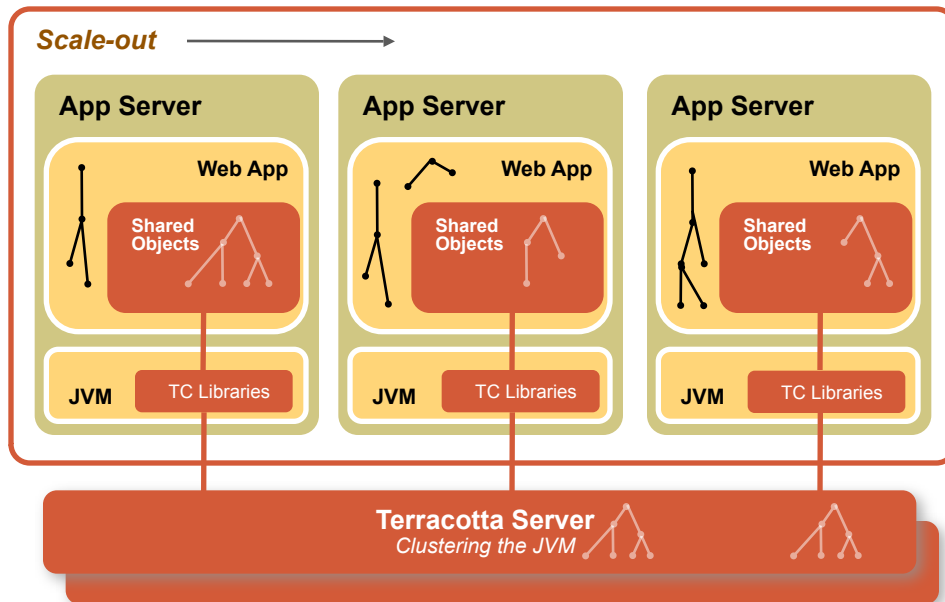
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



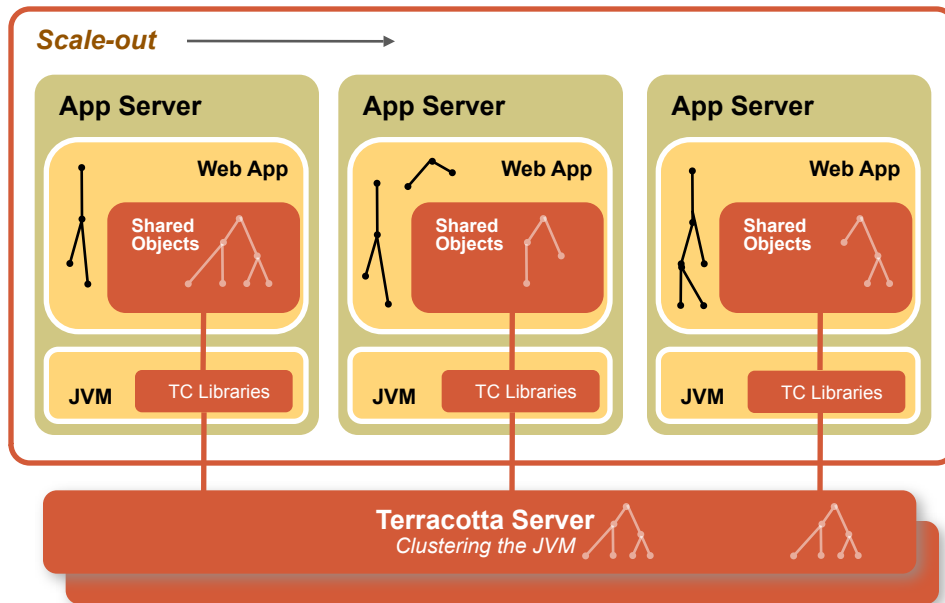
- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Network-Attached Memory



- Management Console
 - Runtime visibility
 - Data introspection
 - Cluster monitoring

- Heap Level Replication
 - Declarative
 - No Serialization
 - Fine Grained / Field Level
 - GET_FIELD – PUT_FIELD
 - Only Where Resident
- JVM Coordination
 - Distributed Synchronized Block
 - Distributed wait()/notify()
 - Fine Grained Locking
 - MONITOR_ENTRY – MONITOR_EXIT
- Large Virtual Heaps

Agenda

- What is it?
- How do I use it?
- What's it good for?
- Q&A

A simple counter...

```
public class Main {  
    // Create single static Main instance  
    public static Main instance = new Main();  
  
    // Main - just count the instance twice  
    public static void main(String[] args) throws InterruptedException {  
        for (int i = 0; i < 100; i++) {  
            instance.count();  
            Thread.sleep(1000);  
        }  
    }  
  
    // Main state - a simple counter  
    private int counter = 0;  
  
    // Synchronize access to make Main thread-safe  
    public synchronized void count() {  
        counter++;  
        System.out.println("Counter is: " + counter);  
    }  
}
```

```
<?xml version="1.0" encoding="UTF-8"?>
<con:tc-config xmlns:con="http://www.terracotta.org/config">
  <servers>
    <server host="%i" name="localhost">
      <dso-port>9510</dso-port>
      <jmx-port>9520</jmx-port>
      <data>terracotta/server-data</data>
      <logs>terracotta/server-logs</logs>
    </server>
  </servers>
  <clients>
    <logs>terracotta/client-logs</logs>
  </clients>
  <application>
    <dso>
      <instrumented-classes>
        <include>
          <class-expression>Main</class-expression>
        </include>
      </instrumented-classes>
      <roots>
        <root>
          <field-name>Main.instance</field-name>
        </root>
      </roots>
      <locks>
        <autolock>
          <method-expression>void Main.count()</method-expression>
          <lock-level>write</lock-level>
        </autolock>
      </locks>
    </dso>
  </application>
</con:tc-config>
```

Coordination

```

public class Coordination {
    public static final Coordination instance = new Coordination();

    public static void main(String[] args) throws Exception {
        instance.run();
    }

    public AtomicInteger    counter  = new AtomicInteger(0);
    public String           msg;

    public void run() throws Exception {
        // ...choose a role and call a method below forever
    }

    private synchronized void getInput() throws IOException {
        System.out.print("Enter a message> ");
        System.out.flush();
        msg = new BufferedReader(new InputStreamReader(System.in)).readLine();
        notify();
    }

    private synchronized void printInput() throws InterruptedException {
        while (msg == null) {
            wait();
        }
        System.out.println(msg);
        msg = null;
    }
}

```


Shared Map

```
public class SharedData {  
  
    private final Map<String, Date> data = new HashMap<String, Date>();  
  
    public synchronized void addName(String name) {  
        data.put(name, new Date());  
    }  
  
    public synchronized void removeName(String name) {  
        data.remove(name);  
    }  
  
    public synchronized Date getJoined(String name) {  
        return data.get(name);  
    }  
  
    public synchronized Set<String> getNames() {  
        return data.keySet();  
    }  
}
```

```

public class Main {
    public static void main(String arg[]) throws IOException {
        new Main().run();
    }

    private final SharedData data = new SharedData();

    public void run() throws IOException {
        while(true) {
            String[] command = readCommand();
            if(command[0].equals("add")) {
                data.addName(command[1]);
            } else if(command[0].equals("remove")) {
                data.removeName(command[1]);
            } else if(command[0].equals("list")) {
                for(String name : data.getNames()) {
                    System.out.println("  " + name + " : " + data.getJoined(name));
                }
            } else if(command[0].equals("quit")) {
                return;
            }
        }
    }

    private String[] readCommand() throws IOException { ... }
}

```

CyclicBarrier

```
public class Node {  
    public static void main(String[] args) throws Exception {  
        int parties = Integer.parseInt(args[0]);  
        Node node = new Node(parties);  
        node.run();  
    }  
  
    private final CyclicBarrier barrier;  
  
    public Node(int parties) {  
        barrier = new CyclicBarrier(parties);  
    }  
  
    public void run() throws Exception {  
        System.out.println("Waiting for other node to join...");  
        int joined = barrier.await();  
        System.out.println("Started: " + joined);  
  
        barrier.await();  
        System.out.println("Ended: " + joined);  
    }  
}
```

Agenda

- What is it?
- How do I use it?
- What's it good for?
- Q&A

Use-cases

● **Relieving Database Overload**

- Distributed Caching
- Hibernate Clustering
- HTTP Session Clustering

● **Simplifying Application Architecture and Development**

- Virtual Heap for Large Datasets
- Clustering OSS Frameworks (Spring, Struts, Lucene, Wicket, EHCache etc.)
- Master/Worker – Managing Large Workloads
- POJO Clustering
- Messaging, Event-based Systems and Coordination-related Tasks

HTTP Session Clustering without Java Serialization


HTTP Session Clustering – Benefits

- Terracotta Sessions gives you:

- Near-Linear Scale
- No java.io.Serializable
- Large Sessions – MBs
- Higher Throughput

- Supported Platforms:

- Jetty,
- JBoss 4.x,
- Tomcat 5.0 & 5.5,
- WebLogic 8.1, WebLogic 9.2,
- WebSphere CE, Geronimo Alpha, WebSphere 6.1



HTTP Session Clustering – DummyCart.java

```
package demo.cart;

import java.util.*;

public class DummyCart {
    private List items = new ArrayList();

    public List getItems() {
        return Collections.unmodifiableList(items);
    }

    public void addItem(String name) {
        items.add(name);
    }

    public void removeItem(String name) {
        items.remove(name);
    }
}
```

HTTP Session Clustering – carts.jsp

```
<%@ page import="java.util.Iterator" %>
<html>
  <jsp:useBean id="cart" scope="session" class="demo.cart.DummyCart" />
  <% String submit = request.getParameter("submit");
     String item = request.getParameter("item");
     if (submit != null && item != null) {
       if (submit.equals("add")) {
         cart.addItem(item);
       } else if (submit.equals("remove")) {
         cart.removeItem(item);
       }
     }
  %>
  <body>
    <p>You have the following items in your cart:</p>
    <ol> <%
      Iterator it = cart.getItems().iterator();
      while (it.hasNext()) {
        %> <li> <% out.print(it.next()); %> </li> <%
      } %>
    </ol>

    <form method="get" action="<%=request.getContextPath()%>/cart.jsp">
      <p>Item to add or remove: <input type="text" name="item" /> </p>
      <input type="submit" name="submit" value="add" />
      <input type="submit" name="submit" value="remove" />
    </form>
  </body>
</html>
```

HTTP Session Clustering – tc-config.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<tc:tc-config xmlns:tc="http://www.terracotta.org/config">
  <application>
    <dso>
      <web-applications>
        <web-application>cart</web-application>
      </web-applications>
      <instrumented-classes>
        <include>
          <class-expression>demo.*</class-expression>
        </include>
      </instrumented-classes>
    </dso>
  </application>
</tc:tc-config>
```

Demo

Cluster Spring and other OSS frameworks

Glimpse of supported integrations

- 'Supported' by Terracotta means:
 - Just include a 'Configuration Module' (OSGi bundle) in your config to cluster your application (a one liner)
 - Plugs in underneath without any setup
 - Technically, Terracotta supports all integrations as long as it runs



Example Configuration Module

– EHCache clustering

```
<?xml version="1.0" encoding="UTF-8"?>
<tc:tc-config xmlns:tc="http://www.terracotta.org/config">
  <clients>
    <modules>
      <module name="clustered-ehcache-1.3" version="1.0.0"/>
    </modules>
  </clients>
  <application>
    <dso>
      <instrumented-classes>
        <include>
          <class-expression>tutorial.*.*</class-expression>
        </include>
      </instrumented-classes>
    </dso>
  </application>
</tc:tc-config>
```


Example – Spring clustering

Terracotta config

```
<spring>
  <application name="tc-jmx">
    <application-contexts>
      <application-context>
        <paths>
          <path>*/applicationContext.xml</path>
        </paths>
        <beans>
          <bean name="clusteredCounter"/>
          <bean name="clusteredHistory"/>
        </beans>
      </application-context>
    </application-contexts>
  </application>
</spring>
```

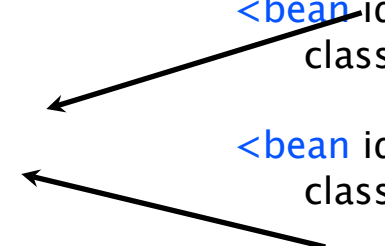
Spring config

```
<bean id="localCounter"
      class="demo.jmx.Counter"/>

<bean id="clusteredCounter"
      class="demo.jmx.Counter"/>

<bean id="localHistory"
      class="demo.jmx.HistoryQueue"/>

<bean id="clusteredHistory"
      class="demo.jmx.HistoryQueue"/>
```



- Terracotta can declaratively cluster Spring beans (Singleton + Session and Custom scoped) with zero code changes
- Can also cluster Spring ApplicationContext events, JMX State and Spring Web Flow

Distributed client-side events and data

Distributed client-side – overview

- Terracotta works for all Java apps, not just server-side
- Method invocations can be distributed : DMI

`<distributed-methods>`

`<method-expression>`

`void com.MyClass.somethingHappened(String, int)`

`</method-expression>`

`</distributed-methods>`

- Works out-of-the-box for Swing events and listeners

Distributed client-side – TableDemo.java

```
class TableDemo extends JFrame {
    private DefaultTableModel model;

    private static Object[][] tableData = {
        { " 9:00", "", "", "" }, { "10:00", "", "", "" }, { "11:00", "", "", "" },
        { "12:00", "", "", "" }, { " 1:00", "", "", "" }, { " 2:00", "", "", "" },
        { " 3:00", "", "", "" }, { " 4:00", "", "", "" }, { " 5:00", "", "", "" }
    };

    TableDemo() {
        super("Table Demo");
        setSize(350, 220);
        setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);
        Object[] header = { "Time", "Room A", "Room B", "Room C" };
        model = new DefaultTableModel(tableData, header);
        JTable schedule = new JTable(model);
        getContentPane().add(new JScrollPane(schedule), java.awt.BorderLayout.CENTER);
    }

    public static void main(String[] args) {
        new TableDemo().setVisible(true);
    }
}
```

Distributed client-side – tc-config.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<tc:tc-config xmlns:tc="http://www.terracotta.org/config">
  <application>
    <dso>
      <instrumented-classes>
        <include>
          <class-expression>demo.*.*</class-expression>
        </include>
      </instrumented-classes>
      <roots>
        <root>
          <field-name>demo.TableDemo.model</field-name>
        </root>
      </roots>
    </dso>
  </application>
</tc:tc-config>
```

Demo

Fine-grained Distributed Caches

Distributed Caches – Benefits

- Simple : pick the cache and just declare the module
 - EHCache, JBoss TreeCache, OSCache
 - Java Collection (e.g. `java.util.HashMap`, `java.util.HashSet`, `java.util.TreeMap`, ...)
- Fast:
 - Field Level Delta changes
 - Replicate only where needed
 - Appliance-like design that can optimize itself
- Big:
 - Virtual heap that exceeds limitations of single JVM
 - Lazy-loading with on-demand faulting and flushing of data
 - Coherent

Distributed EHCache – EHCache.java

```
package tutorial;
import net.sf.ehcache.*;

public class EHCache {
    private CacheManager cacheManager = CacheManager.create("ehcache.xml");
    private Cache cache;

    public EHCache() {
        this.cache = cacheManager.getCache("TestCache");
        if (null == this.cache) {
            Cache cache = new Cache("TestCache", 1000, false, false, 120, 120);
            cacheManager.addCache(cache);
            this.cache = cache;
        }
    }

    public void doCache() {
        cache.put(new Element(new java.util.Date(), System.currentTimeMillis()));
        for (Object key : cache.getKeys()) {
            System.out.println(cache.get(key));
        }
    }

    public static void main(String[] args) {
        new EHCache().doCache();
    }
}
```

Distributed Ehcache – tc-config.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<tc:tc-config xmlns:tc="http://www.terracotta.org/config">
  <clients>
    <modules>
      <module name="clustered-ehcache-1.3" version="1.0.0"/>
    </modules>
  </clients>
  <application>
    <dso>
      <instrumented-classes>
        <include>
          <class-expression>tutorial.*.*</class-expression>
        </include>
      </instrumented-classes>
    </dso>
  </application>
</tc:tc-config>
```

Demo

Agenda

- What is it?
- How do I use it?
- What's it good for?
- Q&A

Wrapping up

Wrapping up

- Terracotta is Network-Attached Memory for the JVM
- Turns Scalability and High-Availability into a **deployment artifact**
- Keep the simplicity of POJO-based development
 - get **Scale-Out with Simplicity**
- Makes mission-critical applications **simpler to:**
 - **Write**
 - **Understand**
 - **Test**
 - **Maintain**
- Endless possibilities for clustering and distributed programming – these were just a few
- Be creative, use your imagination and have fun...

Questions?



<http://terracotta.org>

<http://tech.puredanger.com>