


Node.js - the core

Mark Volkmann
mark@ociweb.com
Object Computing, Inc.
April 12, 2012



Overview

- “Node’s goal is to provide an easy way to build scalable network programs.”
 - <http://nodejs.org/#about>
- A full programming environment, not just for building “servers”
- “The official name of Node is “Node”.
The unofficial name is “Node.js” to disambiguate it from other nodes.”
 - <https://github.com/joyent/node/wiki/FAQ>
- Event-based rather than thread-based; can use multiple processes
- Assumes most time consuming operations involve I/O
 - invoked asynchronously; non-blocking
 - a callback function is invoked when they complete
- Created by Ryan Dahl at Joyent passed control of the project to Isaac Schlueter on 1/30/12
- Runs on top of Chrome V8 (see next slide)
- Implemented in C++ and JavaScript
- Supported on Linux, Mac OS X and Windows



a cartoon from substack

Should You Use It?

- Reasons To Use

- application can benefit from asynchronous, non-blocking I/O
- application is not compute-intensive
- V8 engine is fast enough
- prefer callback or actor models of concurrency
 - over thread-based approach with synchronized access to mutable state
- same language on client and server
- like dynamically typed languages
- large number of JavaScript developers

- Some issues being addressed

- finding packages - there are a large number of them and finding the best ones isn't easy enough
- debugging - stack traces from asynchronously executed code are incomplete
- event loop - sometimes difficult to determine why a program isn't exiting
 - typically due to open connections



Multiple Threads & Processes

- Node uses multiple threads internally
 - to simulate non-blocking file I/O
- You can't create new threads
 - unless you use "Threads A GoGo"
 - <https://github.com/xk/node-threads-a-gogo>
 - "provides an asynchronous, evented and/or continuation passing style API for moving blocking/longish CPU-bound tasks out of Node's event loop to **JavaScript threads that run in parallel** in the background and that use **all the available CPU cores automatically**; all **from within a single Node process**"
- Can use multiple, cooperating processes
 - see "Child Processes" core module
 - processes created with `fork` function can emit and listen for messages
 - see "Clusters" core module
 - "easily create a network of processes that all share server ports"



Chrome V8

- From Google
- Used by Chrome browser and Node.js
- Implemented in C++
- Currently supports ECMAScript 5
- Node adopts the JavaScript syntax supported by V8
 - so will support ES6 when V8 supports it




Where To Look For Functionality

1. JavaScript

- core classes: `Arguments`, `Array`, `Boolean`, `Date`, `Error`,
`Function`, `Global`, `JSON`, `Math`, `Number`, `Object`, `RegExp`, `String`

2. Core Modules

- included with Node
- <http://nodejs.org/docs/latest/api/>
- view source at <https://github.com/joyent/node>
 - JavaScript is in `lib` directory
 - C++ code is in `src` directory



3. Userland Modules (third party)

- typically installed using NPM tool
- <http://search.npmjs.org/>
- 8802 NPM packages on 4/12/12

4. Write yourself

Packages have JavaScript APIs,
but can be partially implemented in C++.

Event Loop


- When a Node program starts, it automatically starts an event loop
 - `node name.js`
- The currently running function, or the main script, can add function calls to the event queue
 - one way is by passing a function to `process.nextTick`
- When the currently running function completes
 - next function in event queue is removed from queue and run
- Most asynchronous functions, such as those that perform I/O
 - take a callback function as an argument
 - add a call to that function to the event queue when their work completes
- Program ends when event queue is empty
 - and there are no open network connections



U.K. roller coaster
(1930 to 2007)

Synchronous vs. Asynchronous

- Asynchronous functions
 - preferred over synchronous in most cases, especially when time to complete is long or unpredictable
 - take a callback function, typically as last argument
 - invoke it, passing an error description as first argument and possibly additional arguments
- Synchronous functions
 - can make application unresponsive if long running
 - do not take a callback function
 - if an error occurs, throw an error description
 - either a string or an Error object
 - throwing an Error is preferred because when strings are thrown, no stacktrace is available



Callbacks

- Functions passed to asynchronous functions
 - often anonymous
- Invoked any number of times,
but often just once when operation completes
- Parameter that accepts callback
 - by convention, last parameter
 - by convention, named `cb` or `callback`
- Callback parameters
 - typically an object describing an error, if any, and a result
 - by convention, error is first argument and is named `err`
- Some libraries require following these conventions
 - ex. Async.js
- See example on next slide



Callback Example

```
var fs = require('fs');           JavaScript

function readObject(filePath, cb) {
  fs.readFile(filePath, function (err, buf) {
    var obj = null;
    if (!err) {
      try {
        obj = JSON.parse(buf); // can throw
      } catch (e) {
        err = e;
      }
    }
    cb(err, obj);
  });
}

readObject('demo.json', function (err, obj) {
  if (err) {
    console.error(err);
  } else {
    console.log(obj);
  }
});
```

```
{                                              demo.json
  "name": "Mark Volkmann",
  "address": {
    "street": "644 Glen Summit",
    "city": "St. Charles",
    "state": "Missouri",
    "zip": 63304
  },
  "hobby": "running"
}
```

```
fs = require 'fs'           CoffeeScript

readObject = (filePath, cb) ->
  fs.readFile filePath, (err, buf) ->
    if !err
      try
        obj = JSON.parse(buf) # can throw
      catch e
        err = e
      cb err, obj

readObject 'demo.json', (err, obj) ->
  if err
    console.error err
  else
    console.log obj
```

Node Globals

(other than standard JavaScript globals)

- Variables defined outside functions
 - are global in browsers
 - are local to current module in Node
- Node global variables
 - `console` - used to write to stdout and stderr
 - `global` - object that holds most global properties and functions
 - can use to share properties across modules; values can be functions
 - `process` - has methods that get info about and interact with current process; extends `EventEmitter`

more on this later
 - `require` - has property `cache` (see next slide)
 - `Buffer` - constructor function for creating objects that read and write data, especially binary data
- Node global functions
 - `require`, `setTimeout`, `clearTimeout`, `setInterval`, `clearInterval`



Node Local Variables

- Node variables that are local to current module

- **`__dirname`**

- full path to directory that contains the module source file

- **`__filename`**

- full path to source file name that defines the module

- ★ • **`module`**

- object that holds information about the current module
- shared by all instances of current module
- main property of interest is **`exports`**

- ★ • **`exports`**

- object used to export properties from a module; values can be functions
- same as `module.exports`

- **`require.cache`**

- a property on the `require` function
- holds required modules so each is only loaded once
- delete a property to allow a module to be reloaded by a subsequent call to `require`

- property is full path to module, ex. `delete require.cache[__dirname + '/mymodule.js'];`



The `require` function has other properties, but they are rarely used directly. They include: `extensions`, `main`, `registerExtension` and `resolve`.

console Methods

similar to methods supported in browsers

- ★ • **console.log(args)** - writes to stdout with a newline
 - first arg can be a string containing formatting directives
 - if not, `util.inspect` is called on each argument (returns string representation of object)
 - formatting directives: `%s` - String, `%d` - Number, `%j` - JSON, `%%` - single percent sign
- **console.info** - same as `console.log`
- **console.warn** - same as `console.log`, but writes to stderr
- **console.error** - same as `console.warn`
- **console.dir(obj)** - writes result of `util.inspect(obj)` to stdout
- **console.time(label)** - marks start time
- **console.timeEnd(label)** - marks end time and outputs label and duration
- ★ • **console.trace** - writes stack trace to stderr
- **console.assert(boolean, msg)**
 - same as `assert.ok();` throws `AssertionError` with `msg` if false

multiple arguments
are output with a
space between each




Process Properties

- ★ • **process.argv** - array containing '`node`', main script absolute file path, and command-line arguments
- ★ • **process.env** - object containing environment variables
 - **process.pid** - process id
 - **process.stdin** - non-blocking readable stream; paused by default
 - emits '`data`' event when return key is pressed and '`end`' event when ctrl-d is pressed

see upcoming example
 - **process.stdout** and **process.stderr** - blocking, writable streams
 - important so output from asynchronous functions isn't interspersed
 - **process.title** - get/set name displayed by `ps` command; defaults to "node"

doesn't work on Mac OS X
 - **process.version** - Node version
 - and more



Process Methods

- `process.chdir(directory)` - changes current working directory
- `process.cwd()` - returns current working directory
- ★ • `process.exit(code)` - exits process with given status code
- `process.memoryUsage()`
 - returns object with `heapTotal` and `heapUsed` properties
- ★ • `process.nextTick(function)`
 - places given function at end of event loop queue so it runs in next iteration of event loop
 - one way to break up a long running function that avoids blocking event loop
- `process.uptime()`
 - returns number of seconds (integer, not float) process has been running
- more

see upcoming example



Process Events

- **exit** - process is about to exit; event loop is no longer running
- **uncaughtException** - error has bubbled to top
 - if a listener is registered, uncaught exceptions will not cause a stack trace to print and program to exit
- POSIX signals - ex. SIGINT emitted when ctrl-c is pressed

```
process.on(event-name, function () {  
  ...  
});
```

for more detail on listening for events,
see slides on **EventEmitter** later



Buffers



- For reading and writing data, including binary data
 - some read and write functions in the “file system” module work with **Buffer** objects
- Must specify encoding when converting between strings and **Buffer** objects
 - `'ascii', 'base64', 'binary', 'hex', 'ucs2', 'utf8'`
- To create a **Buffer**
 - `new Buffer(size-in-bytes)`
 - `new Buffer(array)`
 - `new Buffer(string, encoding='utf8')`
- **Buffer** properties
 - `length` in bytes

“If you pass a **Buffer** to a function,
it's no longer your buffer!
Reading from it or writing to it
at that point is entering the
territory of undefined behavior.”
Issac Schlueter on
Node.js mailing list, 3/15/12

Buffer Writing Methods ...

- **buffer[index] = value;**
 - sets a given byte
- **buffer.write(string, offset=0, length=buffer.length, encoding='utf8')**
 - length is the number of bytes to write
 - if not enough room, will write as many bytes as will fit
 - returns number of bytes written
- **buffer.writetype(value, offset, noAssert=false)**
 - where type is `Int8`, `Int16BE`, `Int16LE`, `Int32BE`, `Int32LE`,
`UInt8`, `UInt16BE`, `UInt16LE`, `UInt32BE`, `UInt32LE`,
`FloatBE`, `FloatLE`, `DoubleBE`, `DoubleLE`
 - when `noAssert` is `true`, it doesn't verify that there is enough space from the offset to the end of the buffer to write the type
 - no return value since the number of bytes written is known from the method name




LE = Little Endian
BE = Big Endian

... Buffer Writing Methods

- **buffer.copy(targetBuffer, targetStart=0, sourceStart=0, sourceEnd=buffer.length)**
 - copies data from one buffer (the method receiver) to another
- **buffer.fill(value, offset=0, end=buffer.length)**
 - **value** is used for each byte
 - **value** should be an integer (0 to 255) or a string (only first byte is used)
 - if only **value** is specified, the entire buffer is filled

Buffer Reading Methods

- **buffer[index]**
 - returns a given byte
- **buffer.toString(encoding, start=0, end=buffer.length)**
- **buffer.readtype(offset, noAssert=false)**
 - where type is `Int8`, `Int16BE`, `Int16LE`, `Int32BE`, `Int32LE`,
`UInt8`, `UInt16BE`, `UInt16LE`, `UInt32BE`, `UInt32LE`,
`FloatBE`, `FloatLE`, `DoubleBE`, `DoubleLE`
 - when `noAssert` is `true`, it doesn't verify that there are enough bytes from the offset to the end of the buffer to read the type
 - returns a `Number`



LE = Little Endian
BE = Big Endian

Other Buffer Methods/Functions

- Other **Buffer** methods
 - **buffer.slice(start, end=buffer.length)**
 - returns a new buffer that shares memory with the receiver
 - **start** is the offset and **end** is the length of the new buffer
- **Buffer** functions
 - **Buffer.byteLength(string, encoding='utf8')**
 - returns byte length of a given string which isn't always the same as **string.length**
 - **Buffer.isBuffer(obj)**
 - determines if an object is a **Buffer**

Buffer Example

```
var buf = new Buffer(100);

buf.writeInt16BE(12345, 0);


buf.writeFloatLE(Math.PI, 16);

var number = buf.readInt16BE(0);
console.log('number =', number);

var pi = buf.readFloatLE(16);
console.log('pi =', pi);
```

Modules

- Defined by a single JavaScript file
 - may “require” others that are their own modules
- Top-level variables and functions defined in them are local to the module
 - not global in the entire runtime like in a browser environment
 - not visible to other modules unless exported
- Each module has its own local variable named “`module`” that refers to an object with these properties
 - ★ • `exports` - initially set to `{ };` see next slide
 - `parent` - module object of module that required this one
 - `filename` - full path to file that defines this module
 - `loaded` - false until first `require` of the module completes; defaults to `false`
 - `paths` - array of filepaths that would be searched to find this module
 - `exited` - no longer used
 - `children` - no longer used



Defining Modules

- A module can expose functions to other modules by exporting them
 - not visible outside module if not exported
- To export many functions
 - `exports.name = some-function;`
 - repeat to export additional things
- To export a single function
 - `module.exports = some-function;` can be a constructor function
 - replaces the default `exports` object
 - exports only one thing from the module
 - not used in conjunction with previous kind of exports
- Should also create `package.json` and `README.md`
 - used by npm
 - used by GitHub

can also export non-function values,
including objects and arrays,
but that isn't as common

A **Node package** is a collection of one or more JavaScript modules, optional C++ source files, optional shell scripts and a `package.json` file that describes the contents of the package and identifies the main module (or uses `index.js` by default).

Using Modules


- **`var name = require('module-name');`**
 1. searches core modules
 2. searches directories listed in `NODE_PATH` environment variable
 - delimited with : in Linux and Mac or ; in Windows
 3. searches upward in path for “node_modules” subdirectories
- **`var name = require('module-path');`**
 1. only reads from specified path; typically start with ./ or ../
- Object returned is typically
 - an object with many properties that are the exported functions
 - a constructor function
 - a single, non-constructor function
- Caches result
 - subsequent requires for same module return cached object without re-reading the file that defines the module
 - unless `require.cache` property matching full path to module is deleted

searches for specified name,
then tries these file extensions:
`.js, .json, .node`

for more detail, see <http://nodejs.org/docs/latest/api/modules.html>

REPL

- Provides a Read-Eval-Print-Loop
 - launched from a terminal window by entering "node"
 - result of last expression is held in variable `_`
- Other than entering standard JavaScript code, the following REPL commands are supported
 - `.help` - lists these commands
 - `.break` - discards a partially entered multi-line expression (ctrl-c does same)
 - `.exit` - exits REPL (ctrl-d does same)
 - `.save {file-path}` - saves every line entered in REPL to specified file
 - `.load {file-path}` - loads a JavaScript file, even if it has already been loaded; picks up changes



Please Use A Lint Tool!


- Find coding errors and style violations, including incorrect indentation
- JSLint
 - from Douglas Crockford
 - very strict and opinionated - "Warning! JSLint will hurt your feelings."
 - <http://jslint.com/>
 - nodelint is an npm module that allows JSLint to be run from command line
 - <https://github.com/tav/nodelint>
- JSHint
 - a fork of JSLint from Anton Kovalyov, Wolfgang Kluge and Josh Perez
 - more configurable, so less opinionated
 - <http://www.jshint.com/>
 - node-jshint is an npm module that allows JSHint to be run from command line
 - <https://github.com/jshint/node-jshint>
 - `npm install -g jshint`



for more detail, see <http://nodejs.org/docs/latest/api/modules.html>

Lint Tool Editor Integration

- Highlights errors/violations as you type!
- Emacs
 - <https://github.com/daleharvey/jshint-mode>
- Vim
 - jslint.vim - <https://github.com/hallettj/jslint.vim>
 - jshint.vim - <https://github.com/manalang/jshint.vim>
- Sublime
 - sublime-jslint - <https://github.com/fbzhong/sublime-jslint>
 - sublime-jshint - <https://github.com/uipoet/sublime-jshint>
 - Sublime Linter - <http://rondevera.github.com/jslintmate/>



```
'use strict';

function avg(arr) {
    var sum = arr.reduce(function (a, b) [
        return a + b;
    ]);
    return sum / arr.length;
}

console.log(avg([1, 2, 3, 4]));
~  
~  
~  
Expected '}' to have an indentation at 6,4
```

Core Modules

Overview

- Core modules are “modules and bindings that are compiled into Node”
- “In general, Node is based on the philosophy that it should not come with batteries included.”
- “One goal of Node's minimal core library is to encourage people to implement things in creative ways, without forcing their ideas onto everyone.”
- “With a tiny core and a vibrant user space, we can all flourish and experiment without the onerous burden of having to always agree”
- See links in API doc at <http://nodejs.org/docs/latest/api/>

Table Of Contents

- Utilities - 4
- Assertion Testing - 7
- OS - 10
- Readline - 13
- TTY - 17
- Events - 19
- Path - 23
- File System - 25
- Streams - 37
- Zlib - 43
- String Decoder - 45
- Net - 46
- Datagram (UDP) - 55
- DNS - 60
- HTTP - 62
- URL - 66
- Query Strings - 69
- Crypto - 70
- TLS/SSL - 71
- HTTPS - 76
- VM - 78
- Child Processes - 82
- Cluster - 86

Utilities ...

```
var util = require('util');
```

- **util.debug(string)** - writes to stderr preceded by "DEBUG: "
- **util.log(string)** - writes to stdout preceded by timestamp and " - "
- **util.format(fmt-string, args)**
 - returns a formatted string
 - formatting directives: %s - String, %d - Number, %j - JSON, %% - single percent sign
 - excess arguments are converted to strings using **util.inspect(arg)**
- **util.inspect(object, hidden=false, depth=2)**
 - returns string representation of an object
 - includes non-enumerable properties only if **hidden** is true
 - traverses objects to default or specified **depth**; pass **null** for infinite



... Utilities ...

- `util.isArray(value)` - determines if an object is an **Array**
 - in ES5, can use `Array.isArray(value)`
- `util.isDate(value)` - determines if an object is a **Date**
- `util.isError(value)` - determines if an object is an **Error**
- `util.isRegExp(value)` - determines if an object is a **RegExp**
- Use `typeof` operator for other tests
 - `typeof(value) === 'boolean' | 'number' | 'string' | 'object' | 'function' | 'undefined'`

Why not just use the `instanceof` operator in place of these?

JavaScript's `instanceof` operator doesn't work across contexts, including those created with Node's "vm" module and created in different browser windows or frames.

The `util.is*` functions provide a more reliable way to determine if an object is of one of these fundamental types.

Here's the implementation of `util.isDate`:

```
function isDate(d) {
  return typeof d === 'object' &&
    objectToString(d) === '[object Date]';
}
```

... Utilities

- `util.inherits(ctor, superCtor)(obj)`

- inherits prototype methods from one constructor into another
- prototype of `ctor` is set to a new object created from `superCtor`
- adds `super_` property to `ctor`

```
var util = require('util');

function MySuper() {}

MySuper.prototype.foo = function () {
  console.log('MySuper foo entered');
};

function MySub() {
  MySuper.call(this);
}

util.inherits(MySub, MySuper);

MySub.prototype.foo = function () {
  MySub.super_.prototype.foo();
  console.log('MySub foo entered');
};

var sub = new MySub();
sub.foo();
```

can pass arguments to superclass ctor here

calls superclass method


Output:

MySuper foo entered
MySub foo entered

Assertion Testing ...

```
var assert = require('assert');
```

- Basic assertions that throw an `Error` if a condition isn't met
- Used by some unit test frameworks
- Actual and expected values are specified in opposite order of many other testing APIs
- Call functions on this `assert` object that are listed on the next slide
 - ex. `assert.equal(score, 100, 'perfect score');`



... Assertion Testing

- Functions

- `ok(value, [message])` or `assert(value, [message])`
 - verifies that `value` is truthy; `value` can be a boolean condition
- `equal(actual, expected, [message])` - uses `==`, so performs type conversions
- `notEqual(actual, expected, [message])` - uses `!=`, so performs type conversions
- `deepEqual(actual, expected, [message])` - also compares nested properties and array elements
- `notDeepEqual(actual, expected, [message])` - ditto
- `strictEqual(actual, expected, [message])` - uses `===`, so no type conversions
- `notStrictEqual(actual, expected, [message])` - uses `!==`, so no type conversions
- `throws(fn, [error], [message])` - succeeds if `fn` throws any error or a specified one
- `doesNotThrow(fn, [error], [message])` - succeeds if `fn` does not throw any error or a specified one
- `ifError(value)` - throws if value is truthy; useful for testing first parameter in callbacks
- `fail(actual, expected, message, operator)`
 - throws `AssertionError` with `message`, ignoring other arguments
 - if `message` is null, the error message `actual + ' ' + operator + ' ' + expected` and `message` isn't used

see Node.js
issue #2993

Assertion Examples

```
var assert = require('assert');
var fs = require('fs');

assert(1 < 2, 'math works');

var actual = [1, [2, 3], 4];
var expected = [1, [2, 3], 4];
assert.deepEqual(actual, expected);

assert.throws(
  fs.readFileSync.bind(null, '/does/not/exist'),
  Error);

assert.doesNotThrow(
  function () {
    fs.readFileSync('demo.js');
  },
  Error);

console.log('calling fs.readFile');
fs.readFile('/does/not/exist', function (err, data) {
  assert.ifError(err);
  console.log('data =', data);
});

assert.fail(null, null, 'did not expect to be here');
```

OS

```
var os = require('os');
```

- Retrieves information about the operating environment
 - processor architecture (ex. x64 or ia32 which are specific Intel processor architectures)
 - host name
 - load average over last 1, 5 and 15 minutes
 - OS platform (ex. 'darwin')
 - OS type (ex. 'Darwin')
 - OS release number
 - uptime in seconds
 - free and total memory in bytes
 - information about each network interface
 - information about each CPU
- Get number of processors with
`os.cpus().length`



OS Example

```
var os = require('os');

console.log('arch =', os.arch());
console.log('hostname =', os.hostname());
console.log('loadavg =', os.loadavg()); // 1, 5 and 15 minute load averages
console.log('platform =', os.platform());
console.log('release =', os.release());
console.log('type =', os.type());
console.log('uptime =', os.uptime(), 'seconds');

console.log('\nfreemem =', os.freemem(), 'bytes');
console.log('totalmem =', os.totalmem(), 'bytes');
var pctFree = os.freemem() / os.totalmem() * 100;
console.log('% free =', pctFree.toFixed(2) + '%');

// Returns object where keys are interface names and
// values are arrays of objects, 1 per address for the interface,
// that have address, family and internal properties.
console.log('\nnetworkInterfaces =', os.networkInterfaces());

// Returns array of objects, 1 per CPU,
// that have model, speed (in MHz) and times
// (# of CPU ticks spent in user, nice, sys, idle and irq) properties.
console.log('\ncpus =', os.cpus());
```

percentages with values between 0 and 1

user - milliseconds executed at user level without nice priority
nice - milliseconds executed at user level with nice priority
system - milliseconds executed at system level
idle - milliseconds doing nothing
irq - milliseconds servicing interrupts (includes waiting on I/O?)

OS Example Output

```
arch = ia32
hostname = Mark-Volkmanns-MacBook-Pro-17.local
loadavg = [ 0.8515625, 0.67724609375, 0.64111328125 ]
platform = darwin
release = 11.2.0
type = Darwin
uptime = 15954 seconds

freemem = 3607273472 bytes
totalmem = 8589934592 bytes
```

```
networkInterfaces = { lo0:
  [ { address: 'fe80::1', family: 'IPv6', internal: true },
    { address: '127.0.0.1', family: 'IPv4', internal: true },
    { address: '::1', family: 'IPv6', internal: true } ],
  en0:
  [ { address: 'fe80::5ab0:35ff:fef3:d095', family: 'IPv6', internal: false },
    { address: '192.168.0.5', family: 'IPv4', internal: false } ],
  en1:
  [ { address: 'fe80::5ab0:35ff:fe6a:23e4', family: 'IPv6', internal: false },
    { address: '192.168.0.6', family: 'IPv4', internal: false } ] }
```

```
cpus = [ { model: 'MacBookPro6,1',
  speed: 2660,
  times: { user: 730150, nice: 0, sys: 569300, idle: 14654330, irq: 0 } },
{ model: 'MacBookPro6,1',
  speed: 2660,
  times: { user: 279490, nice: 0, sys: 126060, idle: 15548110, irq: 0 } },
{ model: 'MacBookPro6,1',
  speed: 2660,
  times: { user: 985960, nice: 0, sys: 616860, idle: 14350840, irq: 0 } },
{ model: 'MacBookPro6,1',
  speed: 2660,
  times: { user: 254950, nice: 0, sys: 115040, idle: 15583670, irq: 0 } } ]
```

lo0 is the “virtual loopback interface”.
Packets sent to this are routed **internally** to the network loopback.
This bypasses local network interface hardware, reducing load on network resources.

Readline ...

```
var rl = require('readline');
```

- Reads streams one line at a time
- Supports many control keys when reading from stdin
 - see `_twrite` method in `readline.js`
- Functions
 - `createInterface(input, output, completer)`
 - `input` and `output` are streams
 - typically `process.stdin` and `process.stdout` (should probably be the defaults)
 - can be a file streams
 - `completer` is a function that provides tab completion; can omit; see example ahead
 - `setPrompt(prompt, [length])`
 - sets prompt displayed when `prompt()` is called
 - if `length` is specified, the prompt will be right-padded with spaces to meet it
 - `prompt()`
 - outputs prompt specified in `setPrompt()`
 - user can enter a line without a prompt

Key	Action
ctrl-c	emits SIGINT
ctrl-z	emits SIGSTP
ctrl-b	back one char
ctrl-f	forward one char
ctrl-left	go to left word boundary
ctrl-right	got to right word boundar
ctrl-a	go to start
ctrl-e	go to end
ctrl-h	delete char to left of cursor
ctrl-d	delete char under cursor
ctrl-u	delete entire line
ctrl-k	delete to end
ctrl-w	delete back to word boundary
ctrl-backspace	same as ctrl-w
ctrl-delete	delete forward to word boundary
ctrl-p	previous in history
ctrl-n	next in history

... Readline

More Functions

- **question(text, callback)**
 - outputs `text` instead of the specified prompt
 - passes line entered by user to callback
 - does not emit '`line`' event, but if additional lines are entered, '`line`' events are generated for those
- **write(text)** - writes string to input as if user typed it
- **pause()** - used internally for tab completion
- **resume()** - used internally for tab completion
- **close()** - marks interface as closed and emits '`close`' event, but doesn't close input stream

Events

- '`line`' - when user presses enter key or there is a newline in the stream
- '`close`' - when `close()` is called or ctrl-c or ctrl-d are pressed

```
var rl = require('readline');

var intf = rl.createInterface(
  process.stdin, process.stdout);

intf.on('line', function (line) {
  // Only invoked if more than one line is entered.
  // The question method doesn't emit this event.
  console.log('line event: got ' + line);
});

intf.question('What is your name? ',
  function (name) {
    console.log('Hello, ' + name + '!');
});
```

Readline Guess Example

```
var rl = require('readline');

// Generate a random number between 1 and 10.
var answer = Math.floor(Math.random() * 10) + 1;
console.log('A number between 1 and 10 has been selected.');

var intf = rl.createInterface(process.stdin, process.stdout);
function prompt(msg) {
    intf.setPrompt(msg);
    intf.prompt();
}

intf.on('line', function (line) {
    var number = parseInt(line, 10);
    if (isNaN(number)) {
        prompt('Enter a number: ');
    } else if (number < answer) {
        prompt('Too low: ');
    } else if (number > answer) {
        prompt('Too high: ');
    } else {
        console.log('CORRECT!');
    }
    // Allow the program to terminate.
    intf.close();
    process.stdin.destroy();
}
);

prompt('Guess the number: '');
```

Readline Completion Example

```
var rl = require('readline');
var fruits = ('apple banana blackberry blueberry cherry grape grapefruit ' +
    'lemon lime orange peach pear plum strawberry').split(' ');

function completer(partial) {
    var options = fruits.filter(function (word) {
        return word.indexOf(partial) === 0;
    });
    return [options, partial];
}

console.log('Enter names of fruits.');
console.log('Press tab for completion.');
console.log('To exit, enter "exit" or press ctrl-c or ctrl-d.');
var intf = rl.createInterface(process.stdin, process.stdout, completer);
intf.setPrompt('fruit: ');
intf.prompt();
intf.on('line', function (line) {
    if (line === 'exit') {
        intf.close();
        process.stdin.destroy(); // allows program to terminate
    } else {
        console.log('got ' + line);
        intf.prompt();
    }
});
```

gets fruits that start with entered text

```
// Asynchronous version
function completer(partial, cb) {
    var options = fruits.filter(function (word) {
        return word.indexOf(partial) === 0;
    });
    cb(null, [options, partial]);
}
```

TTY

```
var tty = require('tty');
```

- Intercepts terminal keystrokes
 - including whether shift, ctrl and meta keys were down
- Important for intercepting certain keystrokes before the operating system acts on them
 - for example, ctrl-c normally sends an interrupt signal (SIGINT) that causes a Node program to stop



TTY Example

```
var tty = require('tty');

process.stdin.resume(); // must do before entering raw mode
tty.setRawMode(true); // can't intercept key presses without this

// char string is only set for normal characters.
// key object is set for all key presses.
// Properties of key include name, ctrl, meta and shift.
process.stdin.on('keypress', function (char, key) {
    console.log('char =', char);
    if (key) {
        console.log('key =', key);
        var name = '';
        if (key.shift) key.name = key.name.toUpperCase();
        if (key.meta) name += 'meta ';
        if (key.ctrl) name += 'ctrl ';
        name += key.name;
        console.log('You pressed ' + name);
        if (key.ctrl && key.name == 'c') {
            console.log('exiting');
            process.exit()
        }
    }
});
```

On Mac OS X, neither the option nor the command key cause the `meta` property to be set!

Events

- Many Node classes inherit from `EventEmitter`
- Custom classes can also
- Objects that are event emitters
 - always emit '`newListener`' when listeners are added
 - often emit '`error`' when an error occurs in one of their methods
- Event listeners
 - functions that are invoked when events are emitted
 - passed any data emitted with the event
 - not passed the event name unless `EventEmitter` subclasses are specifically written to do so




EventEmitter Methods ...

```
var EventEmitter = require('events').EventEmitter;
```

only property exported

- **setMaxListeners(n)**
 - sets the maximum number of listeners that can be registered for a given event
 - default is 10; set to zero for unlimited
 - useful for finding bugs where an excessive number listeners are being registered
 - outputs warning using `console.error` and calls `console.trace`, but does not throw
- **on(event, listener) Or addListener(event, listener)**
 - registers a listener function for a given event
- **once(event, listener)**
 - registers a listener function for a given event and removes it after its first invocation
- **emit(event, args)**
 - invokes listener functions for the event in the order they registered; passes all args to them
 - synchronous! - listener functions are run immediately, not added to event loop queue
 - workaround
 - listener functions can add a function to event loop queue by passing it to `process.nextTick`



light emitting diodes

... EventEmitter Methods

- **removeListener(event, listener)**
 - unregisters a listener function for a given event
- **removeAllListeners([event])**
 - unregisters all listener functions for a given event or all events
- **listeners(event)**
 - returns a live array of all listener functions for a given event
 - can delete function elements to unregister
 - can push function elements to register

Event Example

```
var EventEmitter = require('events').EventEmitter;
var util = require('util');

function Thermostat() {}
util.inherits(Thermostat, EventEmitter);

Thermostat.prototype.set = function (temperature) {
  this.temperature = temperature;
  if (temperature < 32) {
    this.emit('cold', temperature);
  } else if (temperature > 100) {
    this.emit('hot', temperature);
  }
};

var t = new Thermostat();
t.on('cold', function (temp) {
  console.log(temp + ' is too cold!');
});
t.on('hot', function (temp) {
  console.log(temp + ' is too hot!');
});

t.set(50);
t.set(0); // outputs "0 is too cold!"
t.set(110); // outputs "110 is too hot!"
```

Path

```
var path = require('path');
```

- Methods

p parameters are string file paths

- **normalize(*p*)** - returns a new path after resolving .. and ., and replacing consecutive slashes with one, in path *p*
- **join(*path-parts*)** - returns a path created by joining any number of path parts and normalizing
- **resolve([*from...*], *to*)** - resolves relative path *to* to an absolute path also normalizes by prefixing with the *from* values from right to left, why right to left? using the first combination found to exist or the current directory if none exist
- **relative(*from*, *to*)** - returns a relative path that can be used to cd from *from* to *to*
- **dirname(*p*)** - returns the directory portion of path *p*
- **basename(*p*, [*ext*])** - returns the filename portion of path *p*, removing the extension *ext* if supplied
- **extname(*p*)** - returns the extension of path *p*, including a leading dot
- **exists(*p*, *callback*)** - passes a boolean to callback indicating whether the path *p* exists
- **existsSync(*p*)** - returns a boolean indicating whether the path *p* exists



moved to the "File System" module in Node version 7

Path Examples

```
var path = require('path');

console.log(path.normalize('../fs/../console///demo.js'));
// ../console/demo.js

var dirs = ['../url', '../vm', '../zlib'];
var args = dirs.concat('demo.js');
console.log(path.resolve.apply(null, args));
// /Users/Mark/Documents/OCI/SVN/training/Node.js/labs/zlib/demo.js

var absPath = path.resolve('../foo.txt');
// Recall that __dirname holds the absolute path to the current directory.
// var absPath = __dirname + '/foo.txt'; // same as above


console.log(path.dirname(absPath)); // parent of current directory
console.log(path.basename(absPath, '.txt')); // foo
console.log(path.extname(absPath)); // .txt

path.exists(absPath, function (existsP) {
  console.log(absPath + ' exists? ' + existsP); // false
});
```

File System

```
var fs = require('fs');
```

- Wraps access to POSIX file I/O functions
- Provides asynchronous (preferred) and synchronous versions of most functions
 - asynchronous functions take a callback function as their last argument
 - callback functions take an error description as their first argument
 - synchronous functions can throw errors
- Contains many more functions than any other core module
 - buckle up, six slides worth coming next!
 - for parameter details, see <http://nodejs.org/docs/latest/api/fs.html>



File System Functions ...

- Open/Close
 - `open/openSync` - takes a path and returns a file descriptor
 - `close/closeSync` - takes a file descriptor
- Reading
 - `read/readSync` - takes a file descriptor and a `Buffer`; reads specified range of bytes from file into `Buffer`
 - `readFile/readFileSync` - takes a file path; reads entire file; returns data in a `Buffer`
- Writing
 - `write/writeSync` - takes a file descriptor and a `Buffer`; writes specified range of bytes from `Buffer` into file starting at a given position
 - `writeFile/writeFileSync` - takes a file path and a string or `Buffer`; writes bytes in string or `Buffer` to file, replacing existing content



Use of async functions instead of sync functions is strongly encouraged to avoid blocking the event loop with long-running I/O operations.

- `'r'` - Open file for reading. An exception occurs if the file does not exist.
- `'r+'` - Open file for reading and writing. An exception occurs if the file does not exist.
- `'w'` - Open file for writing. The file is created (if it does not exist) or truncated (if it exists).
- `'w+'` - Open file for reading and writing. The file is created (if it does not exist) or truncated (if it exists).
- `'a'` - Open file for appending. The file is created if it does not exist.
- `'a+'` - Open file for reading and appending. The file is created if it does not exist.

modes are used by these functions and their "Sync" counterparts:
`chmod`, `fchmod`, `lchmod`,
`mkdir` and `open`

... File System Functions ...

- Streams

- `createReadStream` - returns an `fs.ReadStream` object
- `createWriteStream` - returns an `fs.WriteStream` object
- see detail on stream objects later

- Directories

- `mkdir/mkdirSync` - takes file path and optional access permissions mode (ex. '755') and creates a directory
- `readdir/readdirSync` - takes file path and gets array of directory contents
- `rmdir/rmdirSync` - takes directory path and deletes directory only if empty

- Links

- `link/linkSync` - creates a file that is a link (a.k.a. hard link) to another
- `symlink/symlinkSync` - creates a file that is a symbolic link (a.k.a. soft link) to another
- `readLink/readLinkSync` - gets info. about the file referred to by a link
- `unlink/unlinkSync` - deletes a link or file; note there is no `rm` function



For more on **hard links**, see
http://en.wikipedia.org/wiki/Hard_link

For more on **symbolic links**, see
http://en.wikipedia.org/wiki/Symbolic_link


... File System Functions ...

● Statistics

- **stat/statSync** - takes a file path; returns an **fs.Stats** object that provides details about the file
- **fstat/fstatSync** - same as **stat** versions, but takes a file descriptor object instead of a file path
- **lstat/lstatSync** - same as **stat** versions, but if file path is to a link, describes the link instead of the target file
- all return an **fs.Stats** object
 - methods: **isFile**, **isDirectory**, **isBlockDevice**, **isCharacterDevice**, **isSymbolicLink**, **isFIFO**, **isSocket**
 - properties: **dev**, **ino**, **mode**, **nlink**, **uid**, **gid**, **rdev**, **size**, **blksize**, **blocks**, **atime**, **mtime**, **ctime**
 - **atime**, **mtime** and **ctime** are **Date** objects

● Timestamps

- **utimes/utimesSync** - takes a file path, atime and mtime; changes atime and mtime values for the file
- **futimes/futimesSync** - same as **utimes** versions, but takes a file descriptor object instead of a file path



atime is time of last access

mtime is time of last content modification

ctime is time of last content, owner or permission change

... File System Functions ...

- Change owner
 - **chown/chownSync** - takes a file path, user id and group id
 - **fchown/fchownSync** - same as **chown** versions, but takes a file descriptor instead of a file path
 - **lchown/lchownSync** - same as **chown** versions, but if file path is to a link, changes the link instead of the target file
- Change mode (access permissions)
 - **chmod/chmodSync** - takes a file path and a mode (an octal number or a string) ex. 0755 or '755'
 - octal literals are not allowed in ES5 strict mode
 - **fchmod/fchmodSync** - same as **chmod** versions, but takes a file descriptor instead of a file path
 - **lchmod/lchmodSync** - same as **chmod** versions, but if file path is to a link, changes the link instead of the target file



... File System Functions ...



• Watching

- **watchFile** - takes a file path, optional options and a callback
 - file path cannot be to a directory
 - calls callback every time the file is accessed (not under Mac OS X!) or modified
 - default options are typically good; see doc for detail
 - callback is passed current and previous `fs.Stats` objects
 - to detect file modification, compare current mtime value to previous one
- **unwatchFile** - takes a file path; stops watching for file access
- **watch** - takes a file path, optional options and a callback
 - file path can be to a directory (typical case)
 - does not detect changes in nested directories
 - default options are typically good; see doc for detail
 - callback is invoked when the file or directory being watched has a change
 - passed an event string (always 'change') and the associated file path (useful when watching a directory and a file in it changes)
 - returns an `fs.FSWatcher` object that emits 'change' and 'error' events and has a `close` method

Currently the file path isn't consistently passed to the callback. It never is under Mac OS X because that OS doesn't provide the information.

When `watch` indicates that the content of a directory has changed, `fs.readdir` can be used to determine which files have changed. See the `watch` example coming up.

From Ben Noordhuis ... " `fs.watch` on OS X and the BSDs is backed by the `kqueue` event mechanism. It has a couple of known shortcomings, lack of *time updates being one of them.

Linux and Windows use the `inotify` and `ReadDirectoryChangesW` interfaces respectively, which are more robust.

Use `fs.watchFile` if you want consistent behavior across Unices. It's not nearly as efficient as `fs.watch` though, and it isn't supported on Windows."

... File System Functions

- Other
 - **fsync/fsyncSync** - synchronizes in-memory data with data on disk
 - **realpath.realpathSync** - resolves relative file paths to absolute paths
 - **rename/renameSync** - renames and/or moves a file
 - takes "from path" and "to path"; "to path" must include file name, not just directory
 - **truncate/truncateSync** - truncates or extends a file to a given byte length

Reading Files - Three Ways

```
var async = require('async');
var fs = require('fs');
var filePath = 'foo.txt';

fs.readFile(filePath, function (err, buf) {
  if (err) throw err;
  console.log(buf.toString());
});

var rs = fs.createReadStream(filePath);
rs.on('data', function (buf) {
  console.log(buf.toString());  
    can be called multiple times  
    for large files
});  
    can also listen for 'error' events

var maxSize = 100;
var buf = new Buffer(maxSize);
var openFile = fs.open.bind(null, filePath, 'r');
var readFile = function (fd, cb) {
  fs.read(fd, buf, 0, buf.length, 0, function (err, bytesRead) {
    console.log(buf.toString());
    cb(err, fd);
  });
};
async.waterfall([openFile, readFile, fs.close], function (err) {
  if (err) throw err;
});  
    see more on async module  
    in "Userland Modules" section
```

This approach has the following advantages:
• can read from a specified chunk of the file
• can read into a specified chunk of the **Buffer**
Disadvantages include:
• all the things above MUST be specified
• the code is longer and more complicated

Writing Files - Three Ways

```
var async = require('async');
var fs = require('fs');

var filePath = 'foo.txt';
var data = 'red\ngreen\nblue\n';

// If file already exists, content is replaced.
fs.writeFile(filePath, data, function (err) {
  if (err) throw err;
});

var ws = fs.createWriteStream(filePath);
ws.write(data);      can listen for 'error' events
ws.end();

var buf = new Buffer(data);
var open = fs.open.bind(null, filePath, 'w');
var write = function (fd, cb) {
  fs.write(fd, buf, 0, buf.length, 0, function (err) {
    cb(err, fd);
  });
}
async.waterfall([open, write, fs.close], function (err) {
  if (err) throw err;
});
```

This approach has the following advantages:
• can write into a specified chunk of the file
• can write from a specified chunk of the **Buffer**
Disadvantages include:
• all the things above MUST be specified
• the code is longer and more complicated

Watch Example ...

```
var async = require('async');
var fs = require('fs');
var dir = '.';

function getStats(dir, cb) {
  fs.readdir(dir, function (err, files) {
    if (err) return cb(err);

    var stats = {};

    // This function is passed to async.every below.
    var iterator = function (file, cb) {
      // Skip hidden files (start with a period).
      if (/^\./.test(file)) return cb(true);
      // Skip Vim backup files (end with a tilde).
      if (/~$/.test(file)) return cb(true);

      fs.stat(file, function (statErr, stat) {
        if (statErr) {
          err = statErr;
        } else {
          stats[file] = stat;
        }
        cb(!err); // stops async.every when there is an error
      });
    };
    async.every(files, iterator, function (result) {
      cb(err, stats);
    });
  });
}

// Gets an fs.Stats object for every file in a given directory.
// dir is a directory path.
// cb is a callback that is passed err and an array of fs.Stats objects.
```

... Watch Example ...

```
function report(name, oldStat, newStat) {
  if (!oldStat && newStat) {
    console.log(name, 'was created');
    return;
  }

  var modified = newStat.mtime > oldStat.mtime;
  if (modified) {
    var diff = newStat.size - oldStat.size;
    var suffix = Math.abs(diff) === 1 ? 'byte' : 'bytes';
    var desc =
      diff > 0 ? 'increased by ' + diff + ' ' + suffix :
      diff < 0 ? 'decreased by ' + -diff + ' ' + suffix :
      'did not change';
    console.log(name, 'content modified, size', desc);
  }
}

var oldStats;
getStats(dir, function (err, stats) {
  oldStats = stats;
});
```

Reports activity for a single file.
name is a file name.
oldStat and newStat are fs.Stats objects.

... Watch Example

```
fs.watch(dir, function () {
  getStats(dir, function (err, newStats) {
    if (err) {
      return console.error(err);
    }

    Object.keys(oldStats).forEach(function (name) {
      if (!newStats[name]) {
        console.log(name, 'was deleted');
      }
    });

    Object.keys(newStats).forEach(function (name) {
      report(name, oldStats[name], newStats[name]);
    });


    oldStats = newStats;
  });
});
```

Under Mac OS X, `null` is always passed to the callback for `filePath`.
The callback is invoked when any file in the directory is created, deleted, or has its contents modified.
It is not invoked when

- a file is merely accessed
- the owner of a file is changed
- the permissions on a file are changed

Streams


- Two types
 - readable streams - created by `fs.createReadStream(file-path, [options])`
 - writable streams - created by `fs.createWriteStream(file-path, [options])`
 - options include `flags` (a mode at bottom of slide 11) and `encoding ('ascii', 'base64', 'binary', 'hex', 'ucs2' or 'utf8')`
- A stream can be one or both (duplex) types
- Classes
 - `Stream` inherits from `EventEmitter` defined in lib/stream.js
 - `ReadStream` and `WriteStream` inherit from `Stream` defined in lib/fs.js



There are several ways to create duplex streams in the core modules including:
the `Stream` pipe method,
the `net.createServer` and `net.connect` functions
(both return a `net.Socket` object which is a duplex stream)
and the `tls.connect` function.

Examples of non-duplex streams include:
`http.ServerRequest`, `http.ServerResponse`,
`fs.ReadStream` and `fs.WriteStream`.

Custom streams of both types can also be created.
For an example, see <https://github.com/dominictarr/event-stream>.



Readable Streams ...

- Events
 - **open** - when stream is ready; callback is passed a file descriptor object
 - **data** - when data has been read
 - callback is passed a `Buffer` object or a string if `setEncoding` was called on the stream
 - **end** - when end of stream is reached
 - no more '`data`' events will be emitted
 - **error** - when a read error occurs
 - **close** - when underlying file descriptor is closed
- Properties
 - **readable** - boolean indication of whether the stream can be read
 - changes to false if an `error` or `end` event is delivered or the `destroy` method is called on the stream

... Readable Streams

- Methods

- `setEncoding(encoding)` - sets character encoding used
 - valid values are `'ascii'`, `'base64'`, `'binary'`, `'hex'`, `'ucs2'` and `'utf8'`
- `pause()` - temporarily stops `'data'` events
- `resume()` - resumes `'data'` events
- `destroy()` - closes underlying file descriptor
 - no more events will be emitted after `close`
- `destroySoon()` - closes underlying file descriptor
 - only after writes complete if the stream is also writable
- `pipe(destination, [options])` - connects this stream to a writable stream

- See example on slide 17

```
var fs = require('fs');
var rs = fs.createReadStream('TaleOfTwoCities.txt');
rs.pipe(process.stdout);
```

everything read from the file
is written to stdout

Reading a File By Lines

```
var fs = require('fs');

function readLines(filePath, cb) {
  var rs = fs.createReadStream(filePath, {bufferSize: 80});
  var leftover = '';

  rs.on('data', function (buf) {
    var lines = buf.toString().split('\n');
    lines[0] = leftover + lines[0];
    leftover = lines.pop(); // chunk at end
    lines.forEach(function (line) {
      cb(line);
    });
  });

  rs.on('end', function () {
    if (leftover.length > 0) {
      cb(leftover);
    }
  });
}

readLines('./story.txt', console.log);
```

callback is invoked once for each line

See slightly better implementation in node-liner userland module.
npm install liner

Writable Streams ...

- Events

- **open** - when stream is ready; callback is passed a file descriptor object
- **drain** - when “kernel buffer” is empty meaning it is safe to write again
- **error** - when write error occurs
- **close** - when underlying file descriptor has been closed
- **pipe** - when stream is passed to **pipe** method of a **ReadStream**

Kernel buffers are used internally by Node to buffer output in case destination streams cannot keep up.

- Properties

- **bytesWritten** - number of bytes written so far
- **writable** - boolean indication of whether stream can be written
 - changes to **false** if an **error** event is delivered or the **end** or **destroy** method is called on the stream (see next slide)

... Writable Streams

- Methods

- `write(string, encoding='utf8', [fd])` - writes a string to stream
 - `fd` parameter is a UNIX-only, rarely used option
- `write(buffer)` - writes contents of `Buffer` to stream

- `end()` - terminates stream
- `end(string, [encoding])` - writes a string to stream and then terminates it
- `end(buffer)` - writes contents of `Buffer` to stream and then terminates it


- `destroy()` - closes underlying file descriptor
 - no more events will be emitted after `close`
- `destroySoon()` - closes underlying file descriptor
 - only after writes complete if stream is also writable

- See example on slide 18

Zlib

```
var zlib = require('zlib');
```

- Supports three kinds of compression and decompression
 - Deflate - from Wikipedia, "Deflate is a lossless data compression algorithm that uses a combination of the LZ77 algorithm and Huffman coding."
 - Deflate Raw - same as Deflate, but doesn't append a zlib header
 - GZIP - based on the Deflate algorithm
- Highly configurable
- Seven classes - create instances with `zlib.createName([options])` ;
 - `Gzip`, `Deflate` and `DeflateRaw` are writeable streams that compress
 - `Gunzip`, `Inflate` and `InflateRaw` are readable streams that decompress
 - `Unzip` is a readable stream that detects the compression type and decompresses
- Convenience functions
 - perform seven operations corresponding to the seven classes without streams
 - `deflate`, `deflateRaw`, `gzip`, `gunzip`, `inflate`, `inflateRaw` and `unzip`
 - each takes a string or `Buffer` object and a callback function that is passed an `Error`, if any, and the result as a `Buffer`



Zlib Example

```
var fs = require('fs');
var zlib = require('zlib');

function zipToFile(data, filePath, cb) {
    zlib.gzip(data, function (err, buffer) {
        if (err) return cb(err);
        fs.writeFile(filePath, buffer, cb);
    });
}

function unzipFromFile(filePath, cb) {
    fs.readFile(filePath, function (err, buffer) {
        if (err) return cb(err);
        zlib.gunzip(buffer, function (err, buffer) {
            cb(err, buffer.toString());
        });
    });
}

var filePath = 'message.gz';
var data = 'This is a message';
zipToFile(data, filePath, function (err) {
    if (err) throw err;
    unzipFromFile(filePath, function (err, result) {
        if (err) throw err;
        console.log('result =', result);
    });
});
});
```

String Decoder

```
var StringDecoder = require('string_decoder').StringDecoder;
```

- Not documented yet
- Handles writing data from buffers that do not end in a complete multi-byte character
- Used by
 - core modules `fs`, `http`, `net`, `repl` and `tls`
 - npm's `read` module
 - which it uses for "`npm init`" to prompt for `package.json` information



Net ...

```
var net = require('net');
```

- Provides methods for implementing TCP servers and clients
- Methods
 - **createServer([options], [callback])**
 - typically used server-side
 - returns a `net.Server` object
 - callback is passed a `net.Socket` object
 - register listeners for events on socket in callback
 - **connect(port, [host], [callback])** - for TCP
connect(path, [callback]) - for Unix socket
 - asynchronously creates a new connection
 - typically used client-side
 - returns a `net.Socket` object and passes nothing to callback
 - `host` defaults to localhost
 - **createConnection(args)**
 - alias for `connect` method



for communicating between processes on same host

... Net

- Methods

- **isIP(s)** - returns 0 if **s** is not an IP address string, 4 if IPv4, and 6 if IPv6
- **isIPv4(s)** - returns boolean indicating whether **s** is a version 4 IP address string
 - pattern is **d.d.d.d** where **d** is an integer between 0 and 255
 - can be represented in 32 bits
- **isIPv6(s)** - returns boolean indicating whether **s** is a version 6 IP address string
 - pattern is **h:h:h:h:h:h:h:h** where each **h** is a 4 character hex value
 - can be represented in 128 bits
 - leading zeros in an **h** value may be omitted
 - **h** values that are all zeros can be replaced by a single zero or omitted
 - all colons must be retained, except more than two consecutive colons can be replaced by only two colons once within an address
 - ex. **1:2:0:0:0:0:7:8** is equivalent to **1:2:::::7:8** and **1:2::7:8**

net.Server Class ...

- Kind of object returned by `net.createServer` function
- Methods
 - `listen(port, [host], callback)` - for TCP
`listen(path, callback)` - for Unix socket
 - listens for new connections
 - if `host` is omitted, will listen for connections from any host
 - returns nothing and passes nothing to callback
 - `pause(ms)`
 - stop accepting new connections for `ms` milliseconds, perhaps for throttling
 - `close()`
 - asynchronously stop accepting new connections permanently
 - a 'close' event is emitted when complete
 - `address()`
 - returns an object containing `port` and `address` (IP) properties



... net.Server Class

- Events

```
register for these with  
server.on(event-name, callback);
```

- **listening** - emitted when server is ready to accept connections
- **connection** - emitted when a connection is made
 - `net.Socket` object is passed to callback
- **close** - emitted when server is no longer accepting connections
- **error** - emitted when an error occurs
 - `Error` object is passed to callback

- Properties

- **maxConnections** - set to limit number of connections
- **connections** - will be set to current number of connections



net.Socket Class ...

- Represents a TCP or Unix socket
- Kind of object returned by `net.connect` function
- Properties
 - `remoteAddress` - remote IP address
 - `remotePort` - remote port number
 - `bufferSize` - size of `Buffer` that holds data to be written before it is sent
 - `bytesRead` - number of bytes read
 - `bytesWritten` - number of bytes written



... net.Socket Class ...

- Methods

- `connect(port, [host], callback)` - for TCP
`connect(path, callback)` - for Unix socket
 - usually `net.connect` is used instead of this
 - might use this to implement a custom socket (by writing a new class that inherits `net.Socket`) or to reuse a closed `Socket` to connect to a different server
 - asynchronously opens a new connection
 - `host` defaults to localhost
 - returns nothing and passes nothing to callback
- `setEncoding(encoding)` - options are 'ascii', 'base64' and 'utf8' (default)
- `write(data, [encoding], [callback])`
 - encoding defaults to 'utf8', callback is invoked after all data has been written
- `end(data, [encoding])`
 - optionally writes more data; closes socket; server will receive 'end' event

... net.Socket Class ...

- Methods

- `pause()` - pauses reading of data; for throttling an upload
- `resume()` - resumes reading of data after a call to `pause()`
- `setTimeout(ms, [callback])`
 - invokes callback once if no reads or writes within ms
 - set to zero (default) for no timeout to wait forever and never invoke a callback
- `address()`
 - returns IP address and port of socket in a object with `address` and `port` properties
- `destroy()` - advanced
- `setNoDelay(bool)` - advanced
- `setKeepAlive(enable, [initialDelay])` - advanced

... net.Socket Class

- Events
 - **connect** - when connection is established
 - **data** - when data is received
 - callback is passed a **Buffer** or string containing the data
 - **end** - when **end()** has been called on socket on other end
 - **timeout** - when timeout occurs (see **setTimeout** method)
 - **drain** - when write **Buffer** becomes empty
 - **error** - when any socket-related error occurs
 - callback is passed an **Error** object
 - **close** - when fully closed
 - callback is passed boolean indicating whether it was closed due to an error

net Example

```
var net = require('net');  
var PORT = 8019;  
  
var server = net.createServer(function (socket) {  
    console.log('client connected'); 2  
  
    socket.on('data', function (data) {  
        console.log('received "' + data + '"'); 5  
    });  
  
    socket.on('end', function () {  
        console.log('client disconnected'); 7  
        server.close();  
    });  
  
    socket.write('hello');  
});  
  
server.on('error', function (err) {  
    console.error(err.code === 'EADDRINUSE' ?  
        'port ' + PORT + ' is already in use' :  
        err);  
});  
  
server.listen(PORT, function () {  
    console.log('listening on ' + PORT); 1  
});
```

Server

```
var net = require('net');  
  
var socket = net.connect(8019, function () {  
    console.log('connected to server'); 3  
});  
  
socket.on('data', function (data) {  
    console.log('received "' + data + '"'); 4  
    socket.write('goodbye');  
    socket.end();  
});  
  
socket.on('end', function (data) {  
    console.log('disconnected from server'); 6  
});
```

Client

	Output from server	Output from client
1)	listening on 8019	3) connected to server
2)	client connected	4) received "hello"
5)	received "goodbye"	6) disconnected from server
7)	client disconnected	

Datagram

```
var dgram = require('dgram');
```

- User Datagram Protocol (UDP)
 - supports datagram sockets
- Datagram overview
 - messages are broken into packets
 - packets are separately addressed and routed
 - faster because it foregoes the handshaking overhead of TCP
 - doesn't guarantee reliability, packet ordering or data integrity
 - suitable when error checking and correction isn't needed or is provided by the application
 - suitable when dropping packets is better than waiting for them
- To create a datagram socket
 - `var dgs = createSocket(type, [callback])`
 - creates a datagram socket of a given type ('`udp4`' or '`udp6`')
 - optional callback gets '`message`' events (more in two slides)

"I have a UDP joke to tell you, but you might not get it" ... unknown

Packet size varies based on the Maximum Transmission Unit (MTU) of the transmission technology used.

For **IPv4** the minimum size is 68 bytes and the recommended size is 576 bytes.

For **IPv6** the minimum size is 1280 bytes.

Typically the actual packet size is at least 1500 bytes.

Datagram Socket Methods

- **`dgs.send(buffer, offset, length, port, address, [callback])`**
 - sends a message that is in a specified chunk of a `Buffer` object
 - callback is passed `err` and number of bytes sent
- **`dgs.bind(port, [address])`**
 - starts listening on a given `port`
 - if `address` is specified, only listens on specified network interface instead of all
- **`dgs.close()`**
 - closes the datagram socket
- **`dgs.address()`**
 - gets address of socket in an object with `address` and `port` properties
- and more

see output from
`os.networkInterfaces()`
later

Datagram Events

- **'message'**
 - when a message is received
 - callback is passed a `Buffer` and `rinfo` object with `address` and `port` properties
- **'listening'**
 - when socket begins listening
- **'close'**
 - when call to `close` method completes
- **'error'**
 - when an error occurs
 - callback is passed an `Error` object

Datagram Server Example

```
var dgram = require('dgram');

var type = 'udp4'; // or 'udp6'
var server = dgram.createSocket(type);

server.on('message', function (msg, rinfo) {
  console.log('got "' + msg + '" from ' +
    rinfo.address + ':' + rinfo.port);

  msg = new Buffer('pong');
  server.send(msg, 0, msg.length, rinfo.port, rinfo.address, function (err, bytes) {
    console.log('bytes sent: ', bytes);
    server.close();
  });
});

server.on('error', function (err) {
  console.error(err);
});

server.on('listening', function () {
  var addr = server.address();
  console.log('listening on ' + addr.address + ':' + addr.port);
});

var PORT = 1234;
server.bind(PORT);
```

Datagram Client Example

```
var dgram = require('dgram');

var type = 'udp4'; // or 'udp6'
var client = dgram.createSocket('udp4');

client.on('message', function (msg, rinfo) {
  console.log('got "' + msg + '" from ' +
    rinfo.address + ':' + rinfo.port);
  client.close(); // only expecting one message
});

client.on('error', function (err) {
  console.error(err);
});

client.on('listening', function () {
  var addr = client.address();
  console.log('listening on ' + addr.address + ':' + addr.port);
});

var msg = new Buffer('ping');
var HOST = 'localhost';
var PORT = 1234;
client.send(msg, 0, msg.length, PORT, HOST, function (err, bytes) {
  console.log('bytes sent: ', bytes);
});
```

Server Output

```
listening on 0.0.0.0:1234
got "ping" from 127.0.0.1:49617
bytes sent: 4
```

Client Output

```
listening on 0.0.0.0:49617
bytes sent: 4
got "pong" from 127.0.0.1:1234
```

automatically selected port

Domain Name System (DNS)

```
var dns = require('dns');
```

- Resolves IP address from a domain name
 - `lookup` function
- Resolves domain name from an IP address
 - `reverse` function
- Retrieves many types of DNS records from a domain name
 - supported DNS record types are A (IPv4), AAAA (IPv6), CNAME (canonical name), MX (mail exchange), NS (name server), PTR (reverse IP lookup), TXT (text), SRV (service locator)
 - `resolve` function takes an array of DNS record types to retrieve
 - these functions return a specific type of DNS record:
`resolve4, resolve6, resolveCname, resolveMx, resolveNs, resolveTxt, resolveSrv`
- For information on DNS record types,
see http://en.wikipedia.org/wiki/List_of_DNS_record_types

DNS Example

```
var dns = require('dns');

var domain = 'www.google.com';

dns.lookup(domain, function (err, address, family) {
  if (err) {
    throw err;
  }
  console.log(domain, address, 'IPv' + family);

  dns.reverse(address, function (err, domains) {
    if (err) {
      console.error('reverse lookup failed');
    } else {
      console.log(domains);
    }
  });
});
```

Output

```
www.google.com 74.125.65.106 IPv4
[ 'gx-in-f106.1e100.net' ]
```

HTTP

```
var http = require('http');
```

- Low-level API
- Typically the **express** module is used which builds on the **connect** module which builds on this
 - so we'll just cover the basics
- Supports streaming of requests and responses
 - rather than buffering until all the data is ready
- Use **querystring** core module to parse query parameters
 - covered in more detail later
- Can send HTTP requests with **http.request** function
 - userland module **request** is often used instead



HTTP Example ...

```
var http = require('http');
var qs = require('querystring');

var PORT = 3002;

// Create an HTTP server and give it a 'request' listener.
var srv = http.createServer(function (req, res) {
  var url = req.url;

  // Many browsers, including Chrome, ask for this first.
  if (url === '/favicon.ico') {
    res.statusCode = 404;
    res.end(); // could also return an icon file and 200 status
    return;
  }

  console.log('method =', req.method);
  console.log('url =', url);
  console.log('headers =', req.headers);
  console.log('HTTP version =', req.httpVersion);

  var index = url.indexOf('?');
  var path = url.substring(0, index);
  console.log('path =', path);
  var queryString = url.substring(index + 1);
  var params = qs.parse(queryString); // can't pass entire URL
  console.log('query parameters =', params);
```

Sample output is based on browsing
http://localhost:3002/foo/bar?
month=April&color=yellow

see output
two slides
ahead

... HTTP Example ...

```
// Decide what to write in response based on path and query parameters.  
// Express supports defining "routes" which makes this easier.  
  
// If there is data in the request body, it can be received in chunks.  
var data = '';  
req.on('data', function (chunk) {  
    data += chunk;  
});  
req.on('end', function () {  
    // All the data has been received now.  
    console.log('data =', data);  
});  
  
var status = 200;  
var responseHeaders = {  
    'Content-Type': 'text/plain'  
};  
// Can set response status and other headers in one call.  
//res.writeHead(status, responseHeaders);  
  
// Can set response status and each header separately.  
res.statusCode = status;  
res.setHeader('Content-Type', 'text/plain');
```

chunk size is limited
by TCP packet size

... HTTP Example

```
// Write the response body after all headers have been written.  
  
// Can write response body in one call.  
//res.end('Hello, World!');  
  
// Can write response body in chunks.  
res.write('Hello');  
res.write(', ');  
res.write('Chunks!');  
res.end();  
});  
  
srv.listen(PORT, function () {  
  console.log('ready');  
});
```

Output

```
ready  
connection created  
method = GET  
url = /foo/bar?month=April&color=yellow  
headers = { host: 'localhost:3002',  
  'user-agent': 'Mozilla/5.0 (Macintosh; Intel Mac OS X 10.7; rv:9.0.1) Gecko/20100101 Firefox/9.0.1',  
  accept: 'text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8',  
  'accept-language': 'en-us,en;q=0.5',  
  'accept-encoding': 'gzip, deflate',  
  'accept-charset': 'ISO-8859-1,utf-8;q=0.7,*;q=0.7',  
  connection: 'keep-alive',  
  'cache-control': 'max-age=0' }  
HTTP version = 1.1  
path = /foo/bar  
query parameters = { month: 'April', color: 'yellow' }  
data = nothing in request body
```

Uniform Resource Locator (URL) ...

```
var url = require('url');
```

- Contains methods to resolve and parse URL strings
- **URL** object properties

- **protocol** - ex. 'http:'
- **auth** - ex. 'username:password'
- **host** - includes port; ex. 'company.com:3000'
- **port** - port number; ex. '3000'
- **hostname** - ex. 'company.com'
- **query** - query object; ex. {month: 'April', color: 'yellow'}
- **search** - query prepended with ?; ex. '?month=April&color=yellow'
- **pathname** - portion after **host** and before **search**; starts with a slash; ex. '/foo/bar'
- **path** - **pathname** plus **search**; ex. '/foo/bar?month=April&color=yellow'
- **hash** - fragment identifier; ex. '#baz'
- **href** - entire URL string;
ex. 'http://username:password@company.com:3000/foo/bar?month=April&color=yellow#baz'

All examples below assume the full URL is 'http://
username:password@company.com:3000/foo/bar?
month=April&color=yellow#baz'

... URL

- Functions

- **parse(urlString, parseQueryString=false, slashesDenoteHost=false)**
 - creates and returns a `URL` object from a URL string
 - if `parseQueryString` is true, `query` property will be an object where keys are query parameter names and values are query parameter values;
ex. { month: 'April', color: 'yellow' }
 - otherwise `query` property value is just the query string portion as a string without leading ?
 - if `slashesDenoteHost` is true, host will be obtained from first slashed part after //;
`ex. url.parse('http://foo/bar/baz', false, true)` returns a URL object where `host = 'foo'`, `path = '/bar/baz'`, and `pathname` is the same
- **format(urlObject)**
 - takes a URL object and returns a URL string
- **resolve(from, to)**
 - returns a URL string created by treating `from` as the base URL and `to` as a relative URL
 - see example on next slide

URL Example

```
'use strict';
var url = require('url');

var urlString =
  'http://username:password@company.com:3000/' +
  'foo/bar?month=April&color=yellow#baz';
var urlObj = url.parse(urlString, true, true);
console.log('urlObj =', urlObj);

urlObj.auth = 'fred:wilma';
urlObj.query.month = 'September';
urlObj.query.color = 'blue';
urlObj.hash = '#barney';
urlString = url.format(urlObj);
console.log('urlString =', urlString);

urlObj = { protocol: 'http:',
  slashes: true,
  auth: 'username:password',
  host: 'company.com:3000',
  port: '3000',
  hostname: 'company.com',
  href: 'http://username:password@company.com:3000/foo/bar?month=April&color=yellow#baz',
  hash: '#baz',
  search: '?month=April&color=yellow',
  query: { month: 'April', color: 'yellow' },
  pathname: '/foo/bar',
  path: '/foo/bar?month=April&color=yellow' }
urlString = http://fred:wilma@company.com:3000/foo/bar?month=April&color=yellow#barney
resolvedUrl = http://www.ociweb.com/knowledge-sharing/tech-com/sett'
```

Output

Query Strings

```
var qs = require('querystring');
```


- Contains methods to parse and create query strings
 - part of URLs from ? to end
- Used by "url" module
- Functions
 - **stringify(obj, sep='&', eq='=')**
 - creates a query string from key/value pairs in obj
 - why would different delimiter characters ever be desired?
 - ex. `qs.stringify({month: 'April', color: 'yellow'})`
returns 'month=April&color=yellow'
 - **parse(str, sep='&', eq='=')**
 - creates an object containing key/value pairs from a query string
 - ex. `qs.parse('month=April&color=yellow')`
returns { month: 'April', color: 'yellow' }
 - other functions are mainly for internal use



Crypto

```
var crypto = require('crypto');
```


- Provides functions for working with security credentials that are used with HTTP and HTTPS
- Works with concepts such as
 - Privacy Enhanced Email (PEM) credential
 - cryptographic hash
 - digest
 - Hash-based Message Authentication Code (HMAC)
 - cipher / decipher
 - signer object
 - verification object
 - Diffie-Hellman key exchange
 - asynchronous PBKDF2
- Relies on OS `openssl` command
- Beyond my knowledge to say anything further



TLS/SSL

```
var tls = require('tls');
```

- Secure Socket Layer (SSL)
- Transport Layer Security (TLS) - an upgrade to SSL 3.0
- Both are cryptographic protocols for secure internet communication
 - public/private key infrastructure
 - prevents eavesdropping and tampering with message content
- Functions
 - `tls.createServer(options, [connectionListener])`
 - called by server code
 - options include key, certificate and certificate authority (CA) file contents
 - can also set `rejectUnauthorized` option to `true` to reject connections not authorized by a CA in list of authorized CAs
 - returns a `tls.Server` object (see next slide)
 - `tls.connect(port, [host], [options], [connectionListener])`
 - called by client code
 - options include key, certificate and certificate authority (CA) file contents
 - returns a `tls.CleartextStream` object (see next slide)



TLS Classes

- **tls.Server**
 - type of object returned by `tls.createServer()`
 - “a subclass of `net.Server` and has the same methods”
 - including `listen(port)`
 - “Instead of accepting just raw TCP connections, this accepts encrypted connections using TLS or SSL.”
- **tls.CleartextStream**
 - type of object returned by `tls.connect()`
 - has same methods and events as readable and writable streams
 - “a stream on top of the encrypted stream that makes it possible to read/write an encrypted data as a cleartext data”

TLS Example Server

```
var fs = require('fs');
var tls = require('tls');

var opts = {
  key: fs.readFileSync('mykey.pem'),
  cert: fs.readFileSync('mycert.pem'),
};

var server = tls.createServer(opts, function (cts) {
  console.log('server connected',
    cts.authorized ? 'authorized' : 'unauthorized');

  cts.setEncoding('utf8');           if encoding isn't specified then data
                                    will be a Buffer instead of a string
                                    (no default encoding)
  cts.write('ping');
  cts.on('data', function (data) {
    console.log('got', data, 'from client');
    server.close();
    process.exit(0);
  });
  cts.on('end', function () {
    console.log('got end event from client');
  });
});

server.listen(8000, function() {
  console.log('ready');
});
```

see HTTPS section (slide 51) for command
to generate key and certificate .pem files

will only get 'authorized'
if `rejectUnauthorized`
option is `true`

TLS Example Client

```
var fs = require('fs');
var tls = require('tls');

var opts = {
  key: fs.readFileSync('mykey.pem'),
  cert: fs.readFileSync('mycert.pem'),
};

var cts = tls.connect(8000, opts);
cts.setEncoding('utf8');
cts.on('secureConnect', function () {
  console.log('client connected',
    cts.authorized ? 'authorized' : 'unauthorized');
});
cts.on('data', function (data) {
  console.log('got', data, 'from server');
  cts.write('pong');
});
cts.on('end', function () {
  console.log('got end event from server; process will exit');
});
cts.on('error', function (e) {
  var msg = e.code === 'ECONNREFUSED' ?
    'failed to connect; is server running?' : e.toString();
  console.error(msg);
});
```

can also pass this callback to `tls.connect()`; this will be set to the `cts` object inside it

if encoding isn't specified then `data` will be a `Buffer` instead of a string (no default encoding)

TLS/SSL Advanced Functionality

- Start a TLS session on an existing TCP connection
- Next Protocol Negotiation (NPN)
 - TLS handshake extension to use one TLS server for multiple protocols (HTTP and SPDY)
- Server Name Indication (SNI)
 - TLS handshake extensions to use one TLS server for multiple hostnames with different SSL certificates

HTTPS

```
var https = require('https');
```

- HTTP over SSL/TLS
 - Secure Socket Layer (SSL) preceded Transport Layer Security (TLS)
 - these are cryptographic protocols
 - from Wikipedia, “encrypt the segments of network connections above the Transport Layer, using asymmetric cryptography for key exchange, symmetric encryption for privacy, and message authentication codes for message integrity”
- Need .pem files for key and certificate
 - “Privacy Enhanced Mail”
 - one way to create is to run following command and answer prompts

```
openssl req -x509 -nodes -days 365 -newkey rsa:2048 -keyout mykey.pem -out mycert.pem
```

type of certificate	makes key	valid for this	type of key	management and
signing request	unencrypted	many days	size in bits	
no password required to use				
- **https.Server** is a subclass of **tls.Server**
- **https.request** function sends a request to a secure web server
 - supports same options as **http.request** function
 - userland module **request** is often used instead



HTTPS Example

```
var fs = require('fs');
var https = require('https');

var PORT = 3002;
var opts = {
  key: fs.readFileSync('mykey.pem'),
  cert: fs.readFileSync('mycert.pem')
};

var srv = https.createServer(opts, function (req, res) {
  // Many browsers, including Chrome, ask for this first.
  if (req.url === '/favicon.ico') {
    res.statusCode = 404;
    res.end(); // could also return an icon file and 200 status
    return;
  }

  res.statusCode = 200;
  res.end('Hello, World!');
});

srv.listen(PORT, function () {
  console.log('ready');
});
```

browse with `https://localhost:3002` or
run "`curl -k https://localhost:3002`"
(-k allows SSL connections without certificates)

Virtual Machine (VM) ...

```
var vm = require('vm');
```

- Compiles a string of JavaScript code and runs it or saves it so it can be run later without recompiling
- The code does not have access to variables in local scope, regardless of the `vm` function used
 - to run code that can access variables in local scope, use the JavaScript `eval` function
- Syntax errors in the code string passed to these functions are reported to stderr and an `Error` is thrown
- Functions in this module that run code return the value of the last expression evaluated
 - `return` statements cannot be used in the top-level of a code string, only inside function definitions within a code string

... VM

• Functions

- `runInThisContext(code, [filename])`
 - global object for code is current global object; assign properties to `global` to make them accessible
- `runInNewContext(code, [sandbox], [filename])`
 - global object for code is sandbox object
 - creates a new context which has overhead
- `createContext([sandbox])`
 - creates a `Context` object that can be passed to `vm.runInContext()`
- `runInContext(code, context, [filename])`
 - context must be created by calling `vm.createContext(sandbox)` (see next slide)
 - global object for code is sandbox object passed to `createContext`
 - context object also holds built-in objects and functions
 - **more efficient than `runInNewContext` when the same context will be used multiple times**
- `createScript(code, [filename])`
 - compiles code and returns a `Script` object that can be used execute the code later
 - see `Script` methods on next slide

The optional filename arguments appear in stack traces to help with debugging.

VM Script Class

- Objects created by calling `vm.createScript(code)`
- Methods
 - `runInThisContext()`
 - global object for code is current global object
 - assign properties to `global` to make them accessible
 - `runInNewContext([sandbox])`
 - global object for code is sandbox object

VM Example

```
var assert = require('assert');
var vm = require('vm');

var code = "Math.pow(x, y)";
var code2 = "console.log('z =', z); " + code;

console.log('global =', global); // has lots of variables and functions
global.x = 3;
global.y = 2;
global.z = 19;
// Note how global functions (in this case just console)
// are explicitly being made available in the sandbox and context.
var sandbox = {x: 2, y: 3, z: 19, console: console};
var context = vm.createContext({x: 2, y: 4, z: 19, console: console});

assert.equal(vm.runInThisContext(code), 9); // 3 ^ 2
assert.equal(vm.runInNewContext(code2, sandbox), 8); // 2 ^ 3
assert.equal(vm.runInContext(code2, context), 16); // 2 ^ 4

var script = vm.createScript(code);
assert.equal(script.runInThisContext(), 9); // 3 ^ 2
assert.equal(script.runInNewContext(sandbox), 8); // 2 ^ 3
assert.equal(script.runInContext(context), 16); // 2 ^ 4
```

Child Processes

```
var cp = require('child_process');
```

- **Functions** [all of these return a `ChildProcess` object]

- **`spawn(command, args=[], [options])`**
 - starts a new process that runs a given command and returns a `ChildProcess` object
 - `args` holds command-line flags and arguments
 - `cwd` option specifies directory in which command runs (defaults to current)
 - `env` option specifies environment variables available in child process (defaults to `process.env`)
 - to obtain output, listen for `data` events on `stdout` and `stderr` properties
- **`exec(command, args=[], options, callback)`**
 - runs a command in a shell, buffers output to `stdout` and `stderr`, and passes it to a callback function of the form `function (err, stdout, stderr)`
 - supports a `timeout` option
 - callback is passed status code, `stdout Buffer` and `stderr Buffer`
- **`execFile(file-path, args=[], [options], callback)`**
 - executes commands in specified file in current process
 - callback is passed status code, `stdout Buffer` and `stderr Buffer`
- **`fork(script-path, args=[], options)`**
 - similar to `spawn`, but returned object has a `send` method that emits '`message`' events

process doesn't end when end of script is reached; must call `process.exit()` in script

ChildProcess Class

- Inherits from **EventEmitter**
- Events
 - **exit** - emitted after child process ends
 - callback function takes a status code and a signal
 - a code is passed on normal termination
 - a signal is passed if terminated by a signal
- Properties
 - **stdin** - standard input stream
 - **stdout** - standard output stream
 - **stderr** - standard error stream
 - **pid** - process id
- Methods
 - **send(*message*)**
 - sends message to child process
 - **kill(*signal='SIGTERM'*)**
 - sends a given signal to the child process

Child Process Example #1

```
var child_process = require('child_process');

var cp = child_process.spawn(
  'ls', ['-l', '..']);
console.log('pid =', cp.pid); runs the "ls -l" command
                                in the parent directory

cp.stdout.on('data', function (data) {
  console.log('data =', data.toString());
});

cp.on('exit', function (code, signal) {
  console.log('exit code =', code);
  console.log('exit signal =', signal);
});
```

pid = 16511
total 0

```
drwxr-xr-x  7 Mark  staff  238 Jan 28 18:36 addons
drwxr-xr-x  4 Mark  staff  136 Dec  7 20:52 async
drwxr-xr-x  3 Mark  staff  102 Nov 21 08:50 buffers
drwxr-xr-x  5 Mark  staff  170 Nov 15 15:03 callbacks
drwxr-xr-x  8 Mark  staff  272 Feb 18 14:04 child_process
...
drwxr-xr-x  5 Mark  staff  170 Jan   8 13:19 vm
drwxr-xr-x  4 Mark  staff  136 Feb 15 18:38 zlib

exit code = 0
exit signal = null
```

Output

Child Process Example #2

finds every required module in the .js files in a below the parent directory

```
var child_process =
  require('child_process');

var args = ['js', 'require'];
var opts = {cwd: '..'};
var file = 'child_process/myFind.sh';
var cp = child_process.execFile(file, args, opts, function (err, data) {
  if (err) {
    return console.error(err);
  }

  var re = /require\(['"](.*)['"]\)\//;
  var requires = {};
  data.split('\n').forEach(function (line) {
    var matches = re.exec(line);
    if (matches) {
      requires[matches[1]] = true;
    }
  });
  Object.keys(requires).sort().forEach(function (req) {
    console.log(req);
  });
});
```

```
#!/bin/bash
# Finds all files with a given file extension
# in and below the current directory
# that contain a given string.
# For example, myFind java "implements Foo"

if [ $# -ne 2 ]; then
  echo usage: myFind {file-extension} {search-string}
  exit 1
fi

find . -name "*.$1" | xargs grep "$2"
```

Output

```
./lib/math
./build/Release/demo
./build/Release/hello
./demo1
./helper
assert
async
child_process
...
util
vm
zlib
```

Cluster

```
var cluster = require('cluster');
```

- “easily create a network of processes that all share server ports”
 - works with any TCP-based server, including HTTP and HTTPS
- Builds on “Child Processes” module
- Initial process is called “master”
 - only process that listens on selected port
 - uses inter-process communication (IPC) pipes to communicate with workers
- Forked processes are called “workers”
 - typically want to fork a number of workers not greater than number of processors
 - get number of processors with `os.cpus().length`
 - no guarantees about order of selection of workers to handle requests
 - distributes connections across workers, but doesn’t distribute requests
 - once a client gets a connection, all their requests will go to the same worker

“The difference between `cluster.fork()` and `child_process.fork()` is simply that cluster allows TCP servers to be shared between workers.

`cluster.fork` is implemented on top of `child_process.fork`.

The message passing API that is available with `child_process.fork` is available with cluster as well.”

“The Jewel Box (also known as NGC 4755, the Kappa Crucis **Cluster** and Caldwell 94) is an open cluster in the constellation of Crux.” ... Wikipedia



Cluster Masters

- Can fork workers
- Can send messages to workers
 - `worker.send('message');`
- Can listen for messages from workers
 - `worker.on('message', function (obj) {...});`
- Can listen for death of workers
 - `cluster.on('death', function (worker) {...});`
 - anything that kills the process triggers this, including OS `kill` command and throwing an uncaught exception
 - can optionally fork replacement workers
 - typically the only job of master after it forks workers

Cluster Workers

- Have a unique id
 - in `process.env.NODE_WORKER_ID` within their process
- Typically start a server such as an HTTP server
- Can send messages to their master
 - `process.send(obj);`
- Can listen for messages from master
 - `process.on('message', function (msg) { ... });`
- Cannot send messages to other workers
- Cannot fork more workers
- Are killed if their master dies

Cluster Example ...

```
var cluster = require('cluster');

if (cluster.isMaster) { // cluster.isWorker is also set
    var requestCount = 0;
    var handleMsg = function (msg) {
        if (msg.cmd === 'gotRequest') {
            requestCount++;
            console.log('requestCount =', requestCount);
        }
    };

    cluster.on('death', function (worker) {
        console.log('worker with pid', worker.pid, 'died - starting new worker');
        worker = cluster.fork();
        worker.on('message', handleMsg);
    });

    // Fork worker processes.
    var cpuCount = require('os').cpus().length;
    for (var i = 1; i < cpuCount; i++) {
        var worker = cluster.fork();
        worker.on('message', handleMsg);
    }
}
```

same code is run
for the master
and each worker

... Cluster Example

```
} else { // for workers
  // Start an HTTP server in worker processes.
  var http = require('http');
  var PORT = 8000;
  http.Server(function (req, res) { // not a constructor function
    if (req.url === '/favicon.ico') {
      res.writeHead(404);
      res.end(); // could also return an icon file and 200 status
      return;
    }

    // Simulate taking a while to process request.
    setTimeout(function () {
      res.statusCode = 200;
      res.end('Hello from process ' + process.pid + '!\n');

      console.log('worker with pid', process.pid, 'handled a request');

      // Send message to master process.
      process.send({cmd: 'gotRequest'});
    }, 1000); // one second
  }).listen(PORT);

  var workerId = process.env.NODE_WORKER_ID; // numbered starting from 1
  console.log('worker server', workerId, 'ready, pid', process.pid);
}
```

1. browse `http://localhost:8000`
2. kill the process that handled the request
3. refresh the page and note that a different process handles the request

Node.js Resources

- Main site - <http://nodejs.org/>
- API doc - <http://nodejs.org/docs/latest/api/>
- NPM Registry Search - <http://search.npmjs.org/>
- How To Node - <http://howtonode.org/>
- node-toolbox - <http://toolbox.no.de/>
- NodeUp podcast - <http://nodeup.com/>
- Felix Geisendoerfer's guide - <http://nodeguide.com>