



SAINT LOUIS JAVA USER GROUP
MAY 2014

STEVEN BORRELLI

steve@borrelli.org

@stevendborrelli

ABOUT ME

FIRST COMPUTER:



SYSTEMS ENGINEERING
MANAGEMENT

FOUNDER, ASTERIS (JAN 2014)

@



ORGANIZER OF STL MACHINE
LEARNING AND DOCKER STL

WHY DOCKER?

Docker makes it easy to:

Package

Deploy

Share



Server Applications

Think:

`java -jar`

vs.

`./configure; make install`

DOCKER FACTS

Written by Docker, Inc. (Formerly Dotcloud)

Automates the management and control of
Linux containers

Rewrite of their proprietary PAAS container
engine (written in Python)

Written in Go / Apache 2 License

11,700+ Github stars

DOCKER TIMELINE

JANUARY 2013: PROJECT START

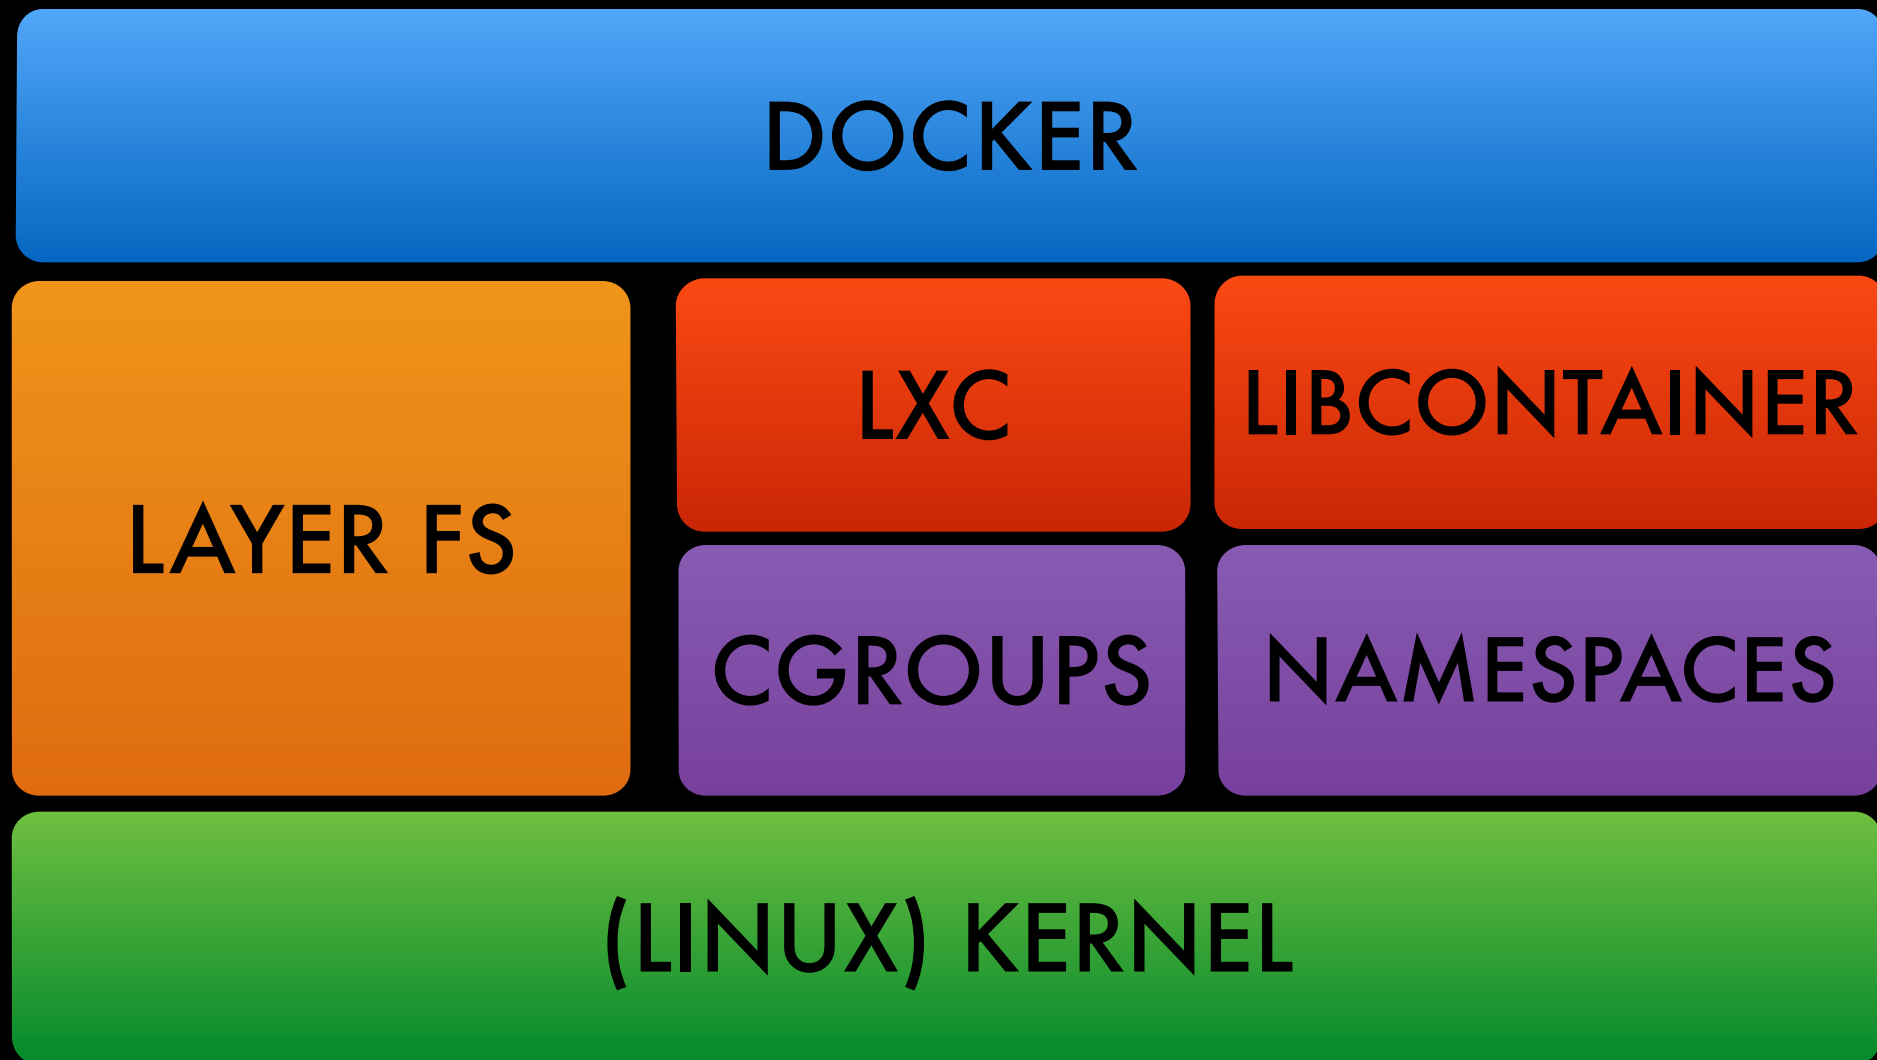
MARCH 2013: INITIAL GITHUB RELEASE

MONTHLY RELEASE CADENCE

MAY 7, 2014: 0.11 RELEASE

MAY 8, 2014: 0.11.1 RELEASE

DOCKER ARCHITECTURE



NAMESPACES VS. CGROUPS

Namespaces provide isolation:

- pid (processes)
- net (network interfaces, routing...)
- ipc (System V IPC)
- mnt (mount points, filesystems)
- uts (hostname)
- user (UIDs)

Control groups control resources:


- cpu (CPU shares)
- cpusets (limit processes to a CPU)
- memory (swap, dirty pages,
- blockio (throttle reads/writes)
- devices
- net_cls, net_prio: control packet class and priority

What's the difference between containers
and virtual machines (VMs)?

VIRTUALIZATION

Virtualization Defined

For those more visually inclined...




HARDWARE VIRTUALIZATION


1966-1972	IBM CP/CMS
1989	Insignia SoftPC
1997	Connectix VirtualPC
1999	VMWare Workstation
2001	IBM AIX LPAR
2002	Xen
2006	<i>Amazon EC2</i>
2007	Sun Logical Domains
2007	Linux KVM
2007	InnoTek VirtualBox
2008	Microsoft Hyper-V

CONTAINERS

Containers vs. VMs



Containers are isolated, but share OS and, where appropriate, bins/libraries



PROCESS VIRTUALIZATION

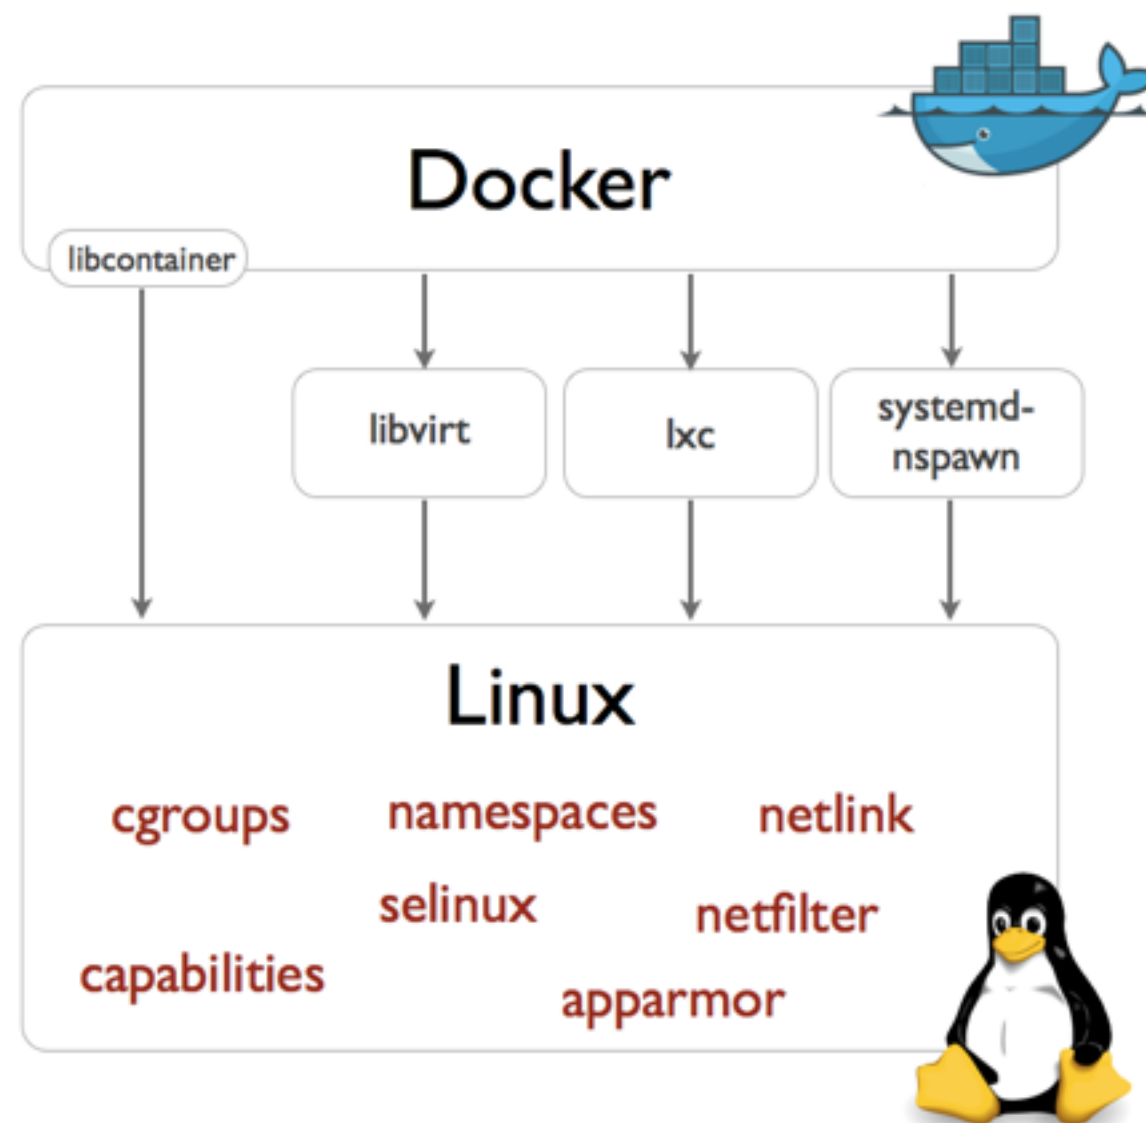
1979-1982	UNIX Chroot
1998	FreeBSD Jail
2001	Parallels Virtuozzo
2001	Linux-VServer
2005	Solaris Zones
2005	OpenVZ
2008	<i>Linux LXC</i>
2007+	<i>PAAS: Heroku, Joyent, CloudFoundry</i>
2013	Docker

Differences between containers and virtual machines

- Weaker isolation in containers
- Containers run near-native speed CPU/IO
- Containers launch in around 0.1 second (libcontainer)
- Less memory overhead

NOTABLE CHANGES

0.9: LIBCONTAINER



NOTABLE CHANGES

0.10:

- TLS support on docker API
- Systemd integration via API instead of /proc
- Lots of cleanups

NOTABLE CHANGES

0.11:

- Release Candidate for 1.0
- Multiple registries
- Direct host network access
- SELinux support

EXAMPLES

RUNNING A CONTAINER

Start a container:

```
$ sudo docker run -i -t ubuntu:12.04 /bin/bash
root@09aa796197bc:/#
```

Mount host filesystems:

```
$ sudo docker run -i -t -v /var/log:/var/host/logs ubuntu /bin/bash
root@bb52ddbdb91c:/# ls /var/host/logs | head -3
apt
auth.log
boot.log
```

MAPPING PORTS

Example: run Zookeeper + Exhibitor

```
sudo docker run --name zookeeper -d -e EXHIBITOR_HOSTNAME=$(hostname) -p 2181:2181 -p 2888:2888  
-p 3888:3888 -p 8080 -p 8443 asteris/zookeeper
```

Port 2181 on host will be mapped to 2181 on container

```
sudo docker run --name zookeeper -d -e EXHIBITOR_HOSTNAME=$(hostname) -p 2181:2181 -p 2888:2888  
-p 3888:3888 -p 8080 -p 8443 asteris/zookeeper
```

Host ports will be dynamically allocated by docker

CONTAINER ID	IMAGE	COMMAND	CREATED	STAT
US	PORTS	NAMES		
485bdc8a5312	asteris/zookeeper:3.4.5	/bin/sh -c /opt/exhi	5 minutes ago	Up 5
minutes	0.0.0.0:2181->2181/tcp, 0.0.0.0:2888->2888/tcp, 0.0.0.0:3888->3888/tcp, 0.0.0.0:49156->8080/tcp, 0.0.0.0:49157->8443/tcp	zookeeper		

DIRECT HOST NETWORK

New in 0.11, allows a container to access host adapters:

```
# docker run -d --net=host tomcat
```

Port 8080 on the container is 8080 on the host:

tcp	0	0 0.0.0.0:22	0.0.0.0:*	LISTEN
tcp6	0	0 :::22	:::*	LISTEN
tcp6	0	0 :::8080	:::*	LISTEN

IMMUTABLE SERVERS

Physical server lifetime is measured in years.

A container's lifetime can be as short as a few seconds.

Treat containers like a build artifact.

If you need to make changes, build a new container.

JAVA DOCKERFILE

```
FROM dockerfile/ubuntu
```

```
RUN add-apt-repository -y ppa:webupd8team/java
```

```
RUN apt-get update
```

```
RUN echo debconf shared/accepted-oracle-license-v1-1 select true | debconf-set-selections
```

```
RUN echo debconf shared/accepted-oracle-license-v1-1 seen true | debconf-set-selections
```

```
RUN apt-get install -y oracle-java7-installer
```

TOMCAT DOCKERFILE

```
FROM dockerfile/java
```

```
RUN apt-get -y update && apt-get -y install tomcat7 tomcat7-admin tomcat7-examples
```

```
RUN echo "JAVA_HOME=/usr/lib/jvm/java-7-oracle" >> /etc/default/tomcat7
```

```
EXPOSE 8080
```

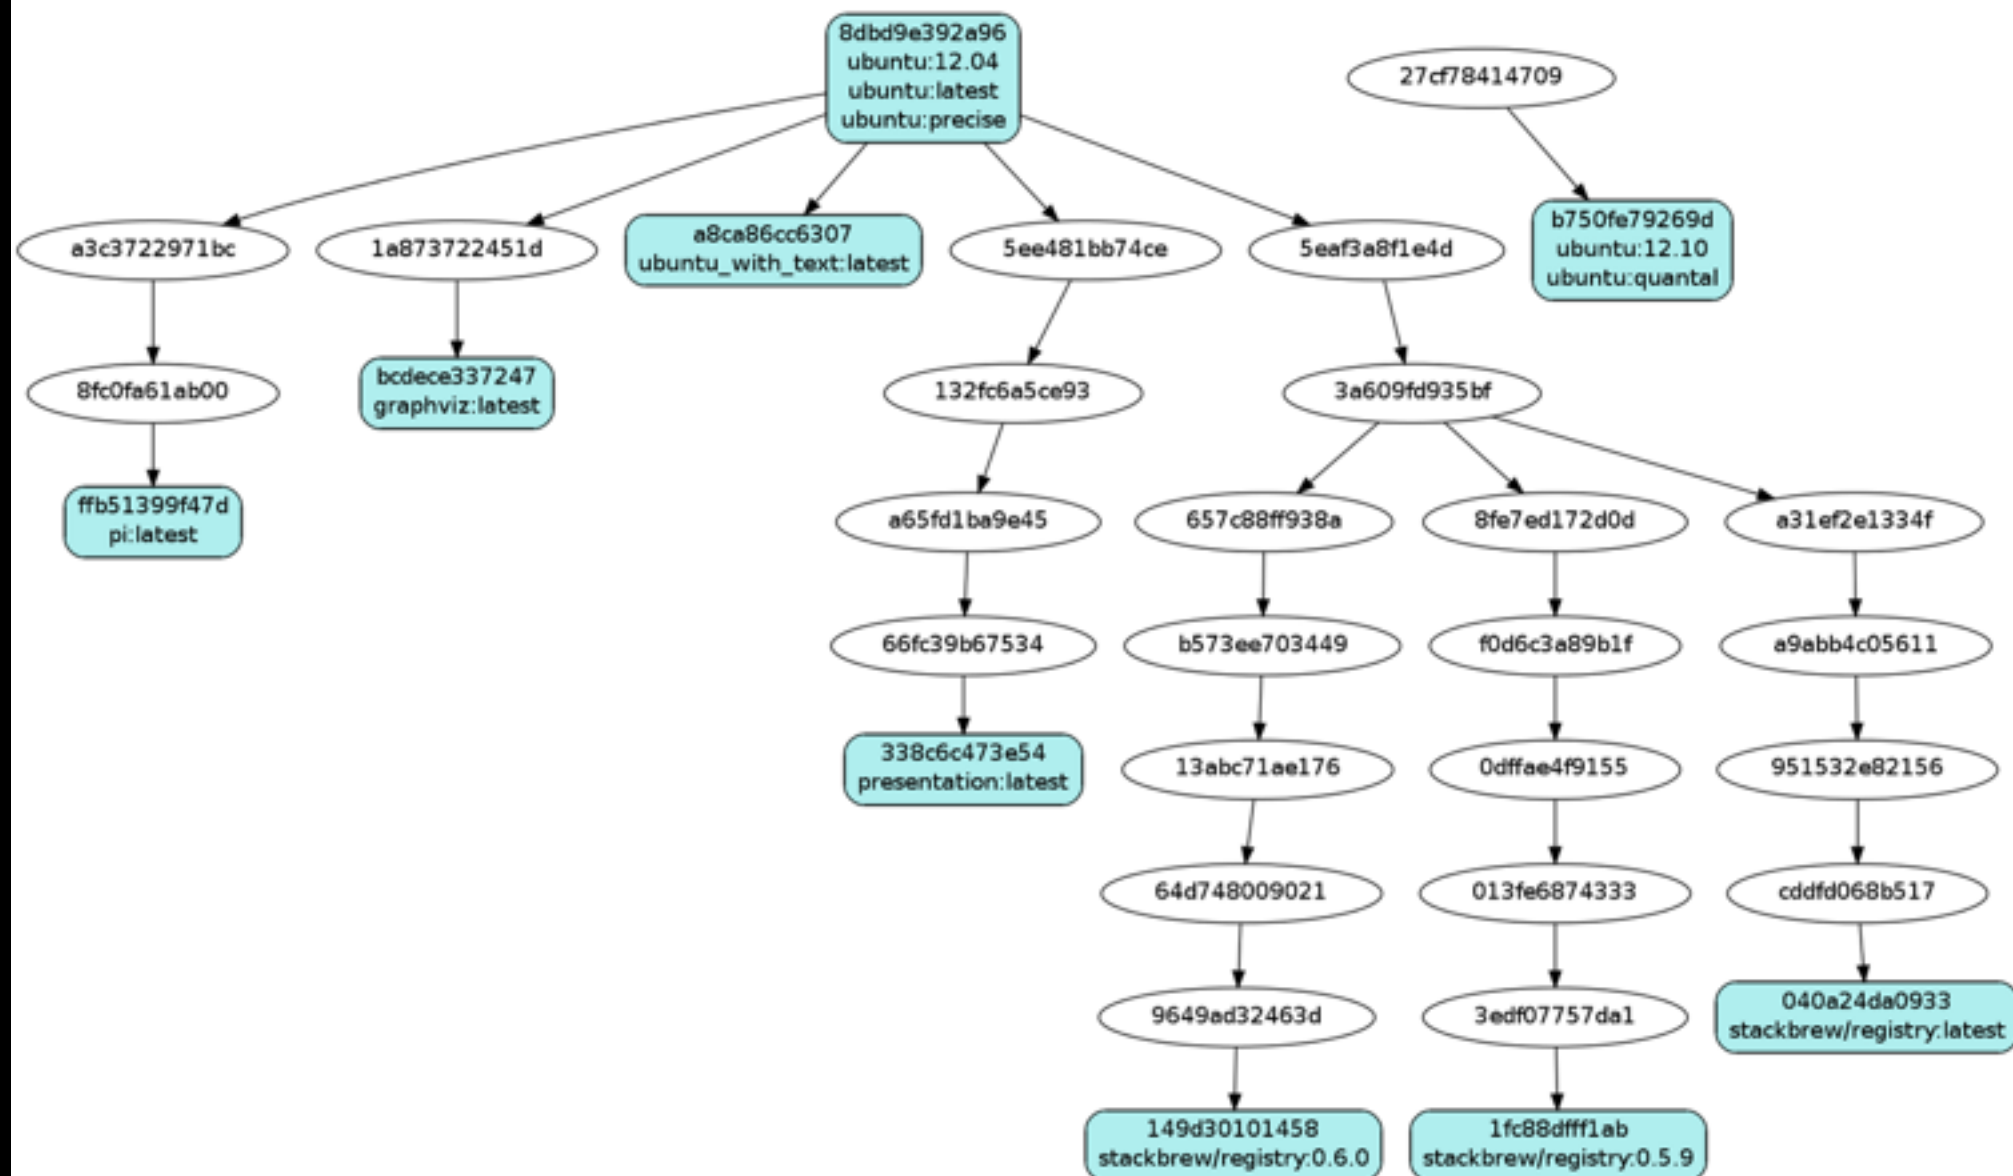
```
ENV CATALINA_HOME /usr/share/tomcat7
```

```
ENV CATALINA_BASE /var/lib/tomcat7
```

```
CMD /usr/share/tomcat7/bin/catalina.sh start && tail -f /var/lib/tomcat7/logs/catalina.out
```


LAYERED FS IS A GRAPH

```
docker$ sudo docker images --viz=true | dot -Tpng -o graph.png
docker$
```



RUNNING JAVA IN DOCKER

Problem: keep configuration out of containers

- Pass in environment variables to Java vars (i.e. in `start.sh`):

```
java -Dkeystore.password=${KEY_PASS}
```

- When you run the container, set the vars:

```
docker run -e SSL_PASS=password tomcat
```

RUNNING JAVA IN DOCKER

Problem: keep configuration out of containers

- Link from a volume container

```
docker run -v /opt/properties -v /opt/ssl \
-name TOMCAT-CFG busybox true
```

```
docker run -t -i -rm -volumes-from TOMCAT-
CFG -name appsrv1 tomcat
```

- Mount host filesystem:

```
docker run -v /opt/ssl:/opt/ssl tomcat
```


RUNNING JAVA IN DOCKER

Running stateless containers

- Use a redis/DB to store session data
- Use a shared FS (hdfs/nfs, etc.)
or Object storage (Swift, S3) for data
- Send logs to a centralized location
- Docker future: storage plugins

WHAT'S NEXT?

Stripped-down operating systems



WHAT'S NEXT?

New operational models:

Continuous Delivery

Automated routing

Distributed consensus (Paxos, Raft)

Service Discovery (Zookeeper, etcd, serf, skydns, consul)

Distributed scheduling (Fleet, Mesos, YARN)

SUMMARY

Easy to build, run & share containers

Rapidly expanding ecosystem

Better performance vs. VMs

Layered filesystem gives us git-like control of images.

Reduces complexity of system builds

Q & A