



Apache Cassandra

STL Java Users Group

Cliff Gilmore

DataStax Solutions Architect / Engineer

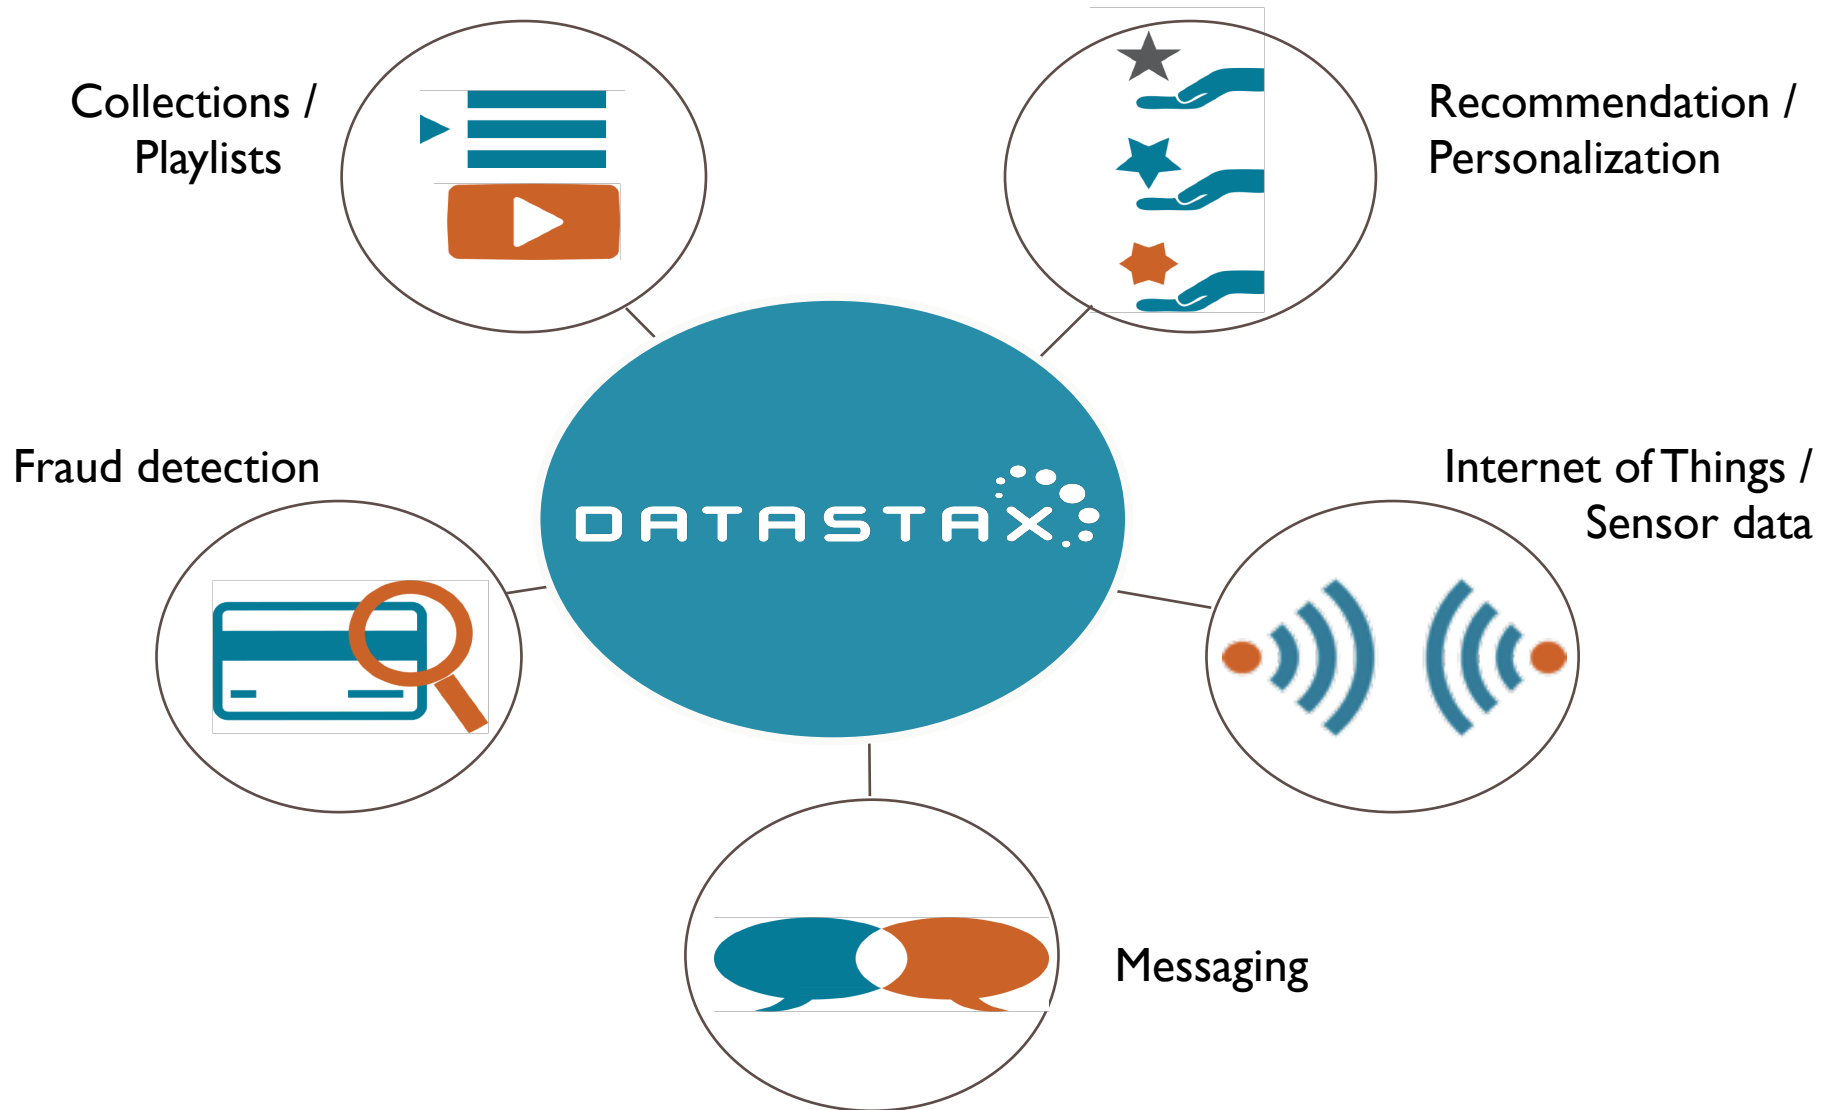
Aug 14, 2014

Agenda

- Cassandra Overview
- Cassandra Architecture
- Cassandra Query Language
- **Interacting with Cassandra using Java**
- About DataStax

CASSANDRA OVERVIEW

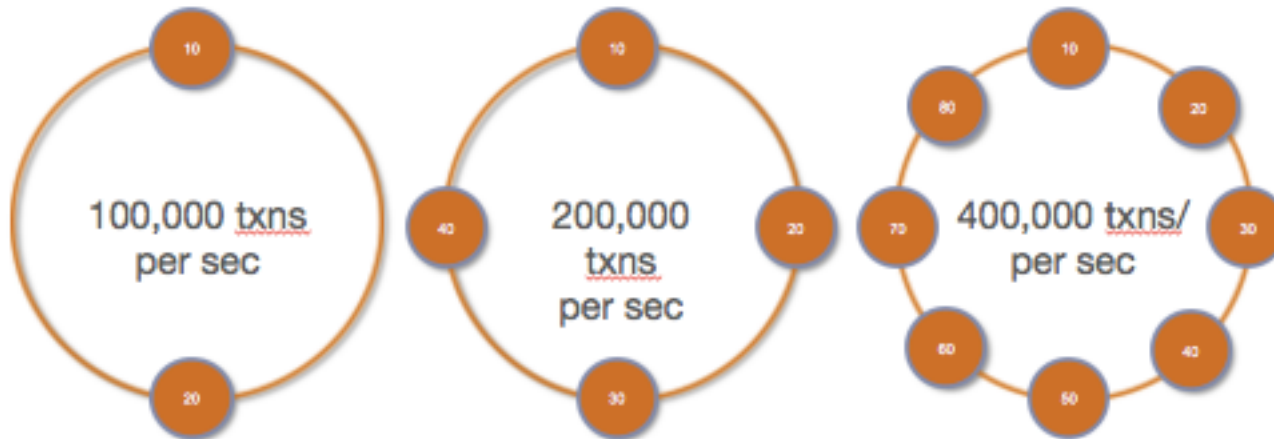
Who is using DataStax?



What is Apache Cassandra?

Apache Cassandra™ is a massively scalable NoSQL database.

- Continuous availability
- High performing writes and reads
- Linear scalability
- Multi-data center support



The NoSQL Performance Leader

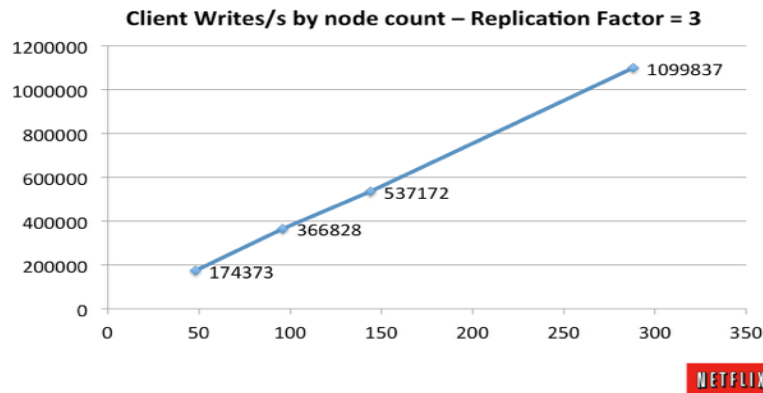


“In terms of scalability, there is a clear winner throughout our experiments. Cassandra achieves the highest throughput for the maximum number of nodes in all experiments with a linear increasing throughput.”

Source: Solving Big Data Challenges for Enterprise Application Performance Management benchmark paper presented at the Very Large Database Conference, 2013.

Netflix Cloud Benchmark...

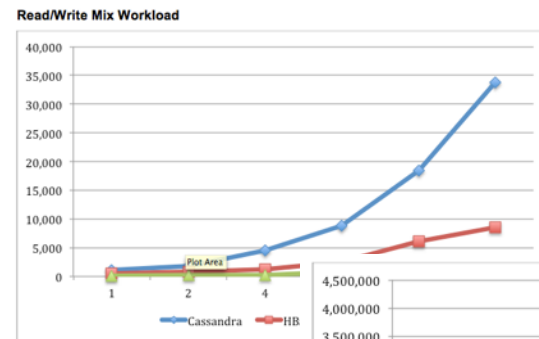
Scale-Up Linearity



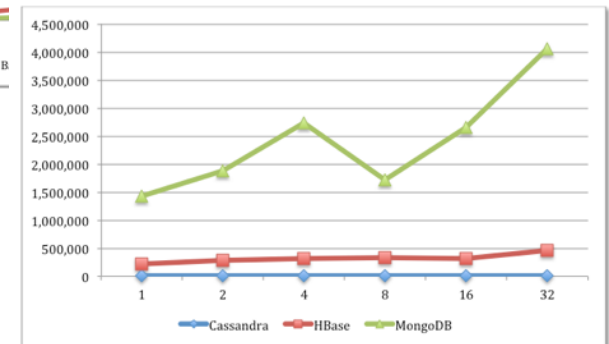
Source: Netflix Tech Blog

End Point Independent NoSQL Benchmark

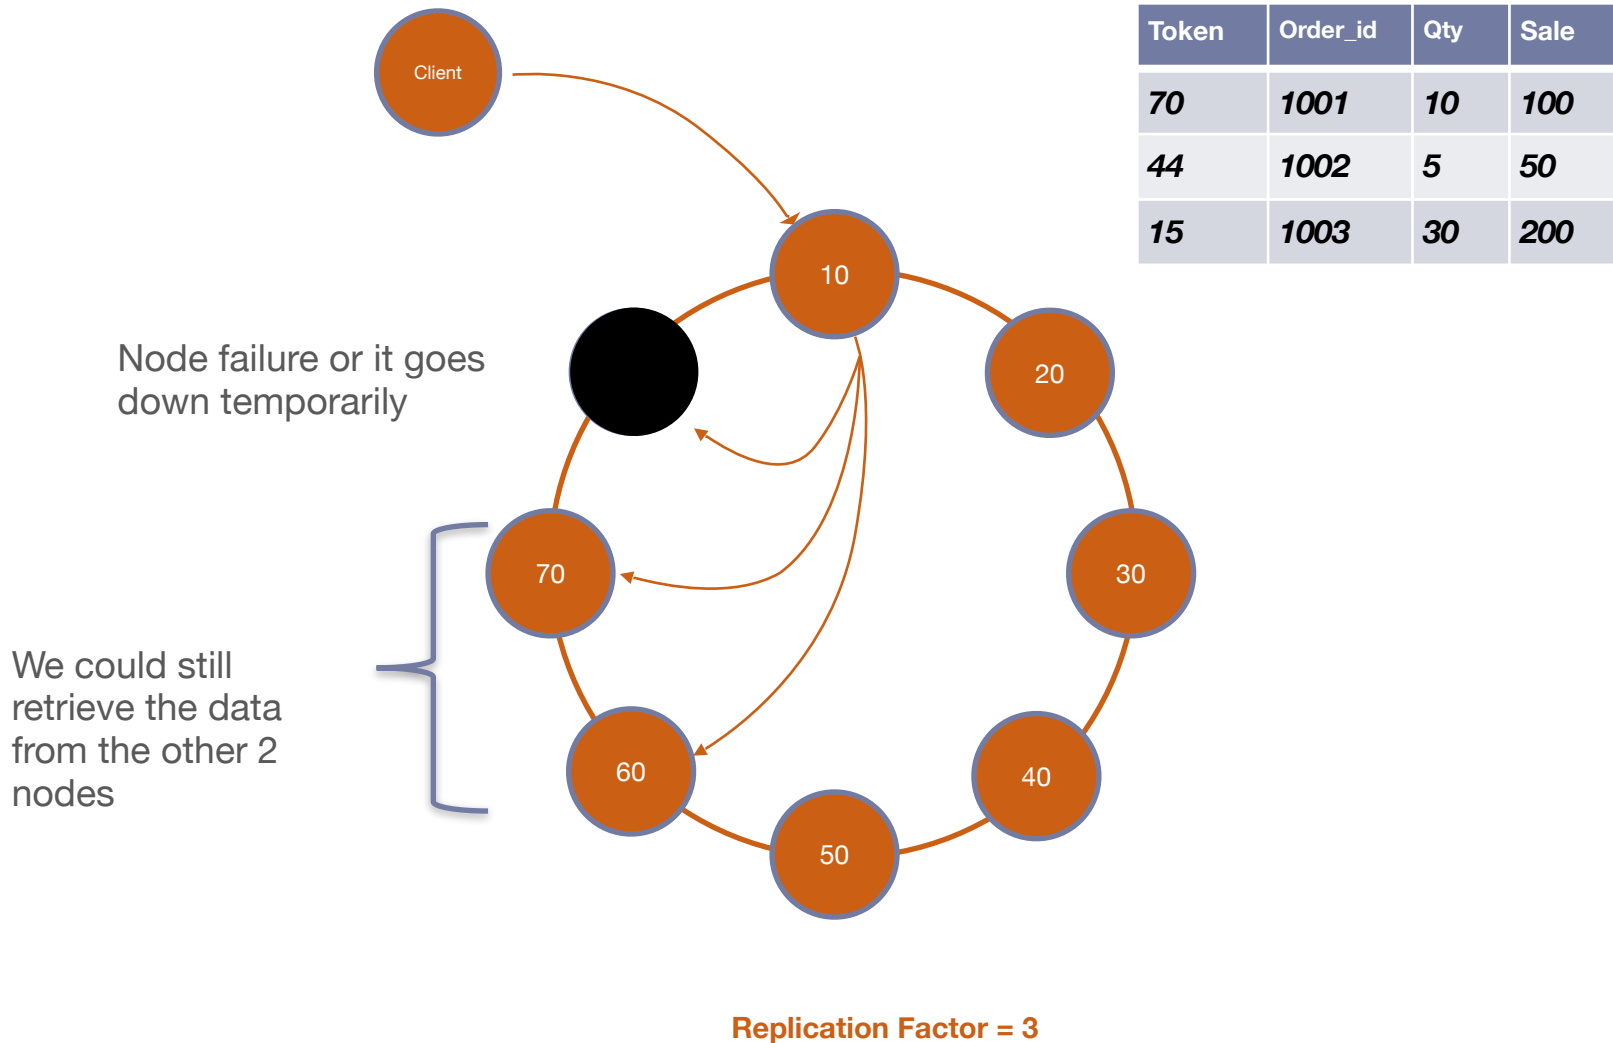
Highest in throughput...



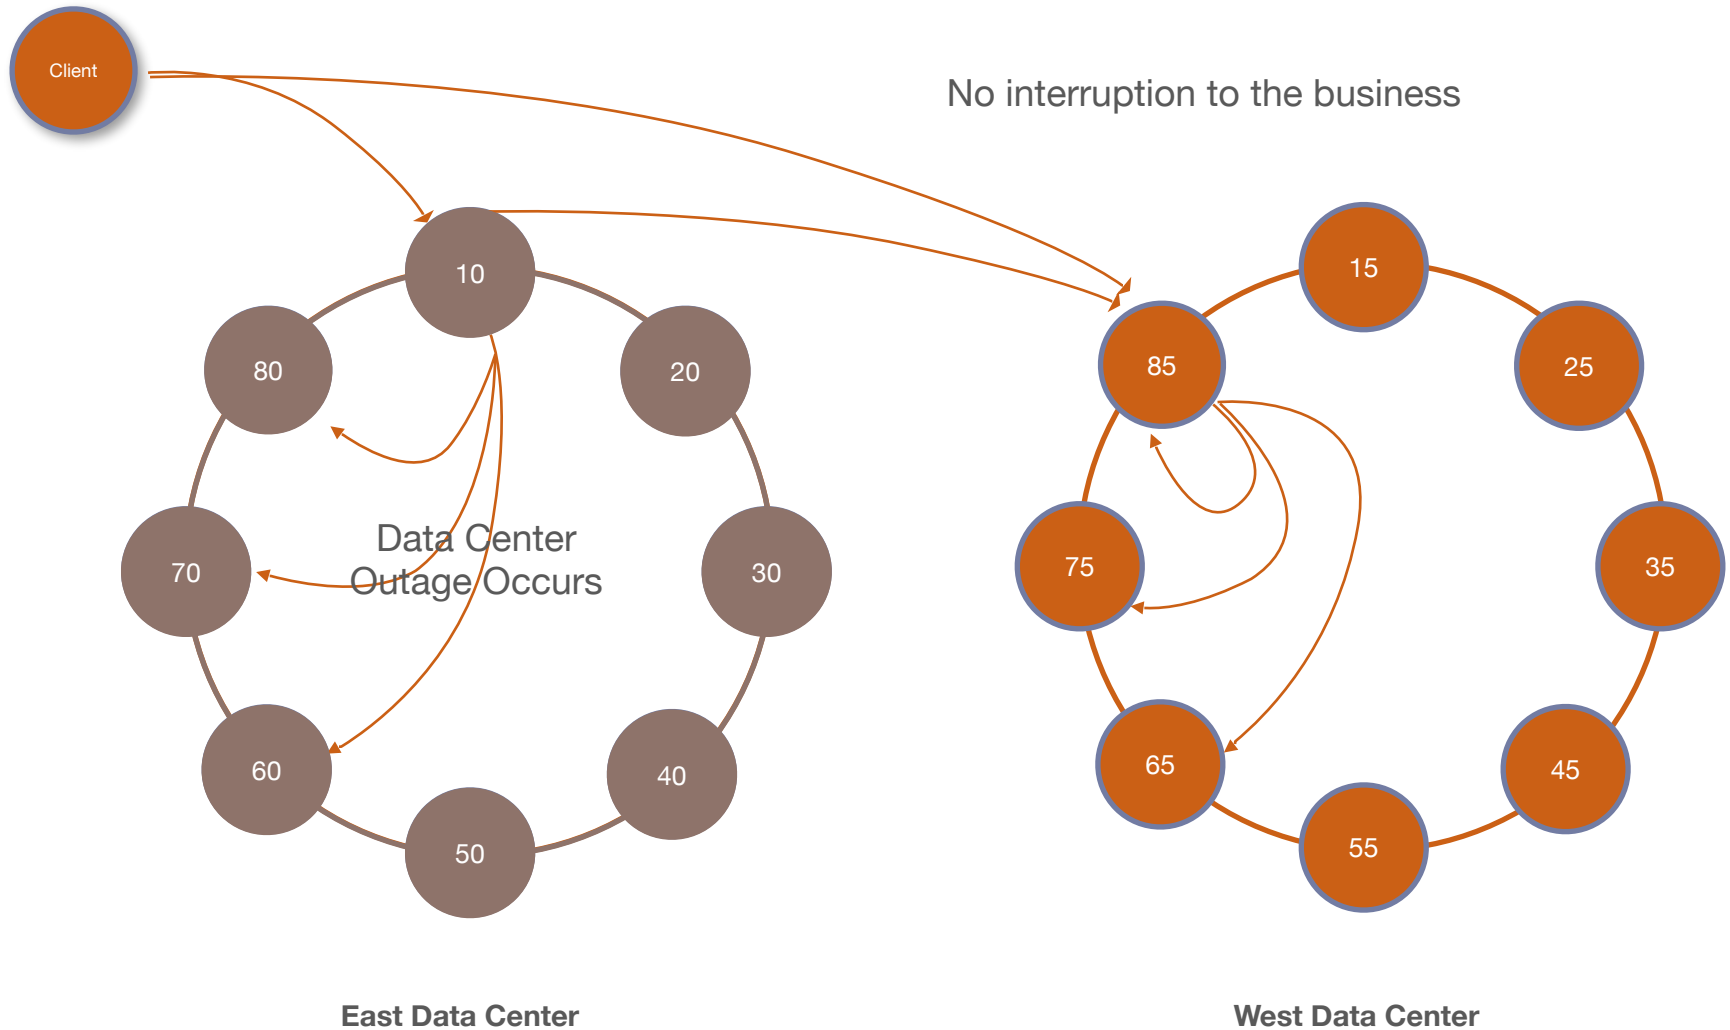
Lowest in latency...



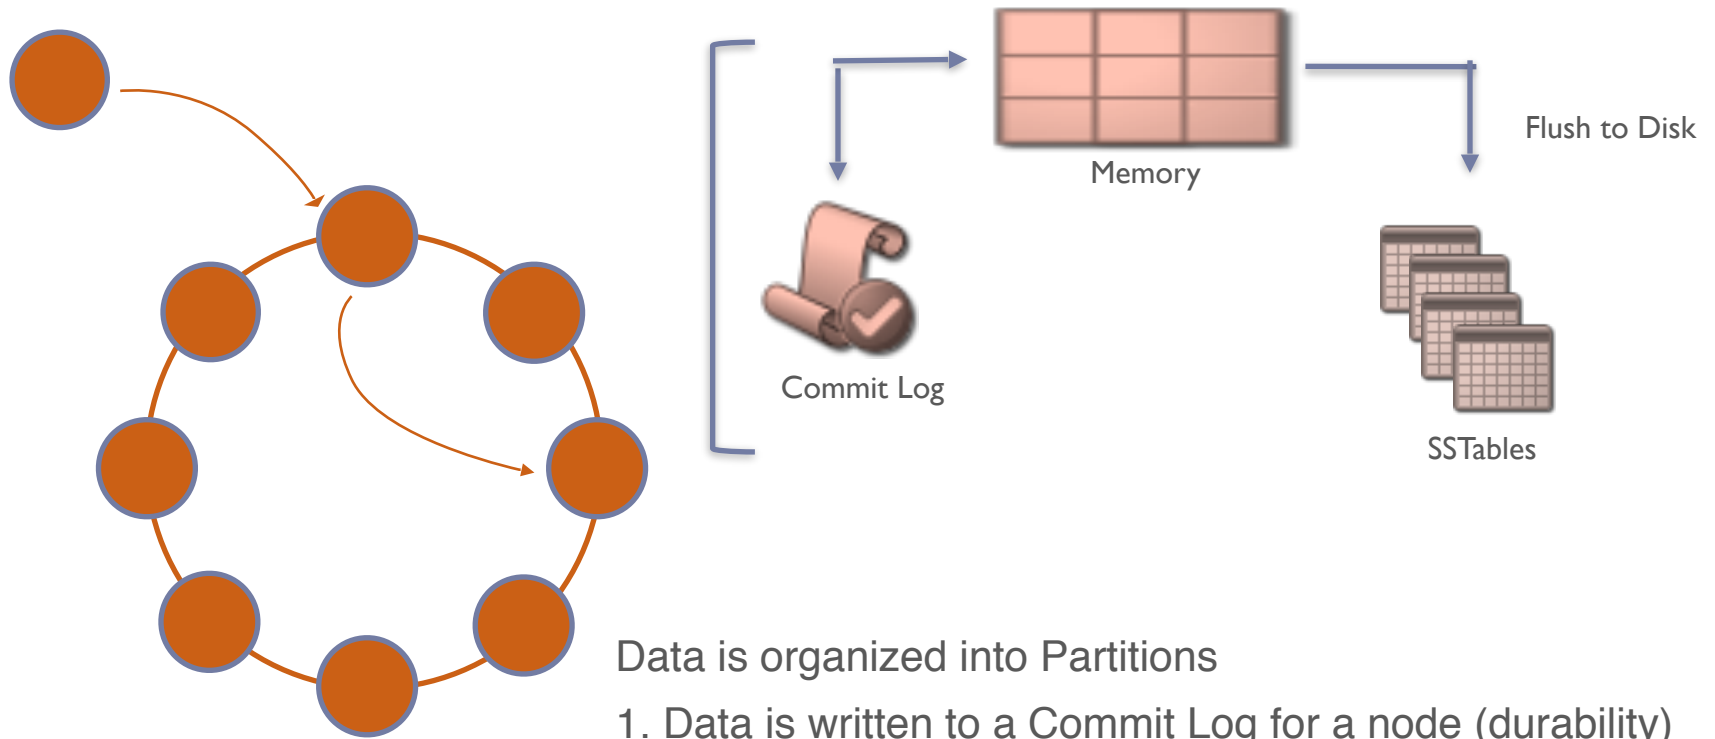
Cassandra is Fault Tolerant



Multi Data Center Support



Writes in Cassandra

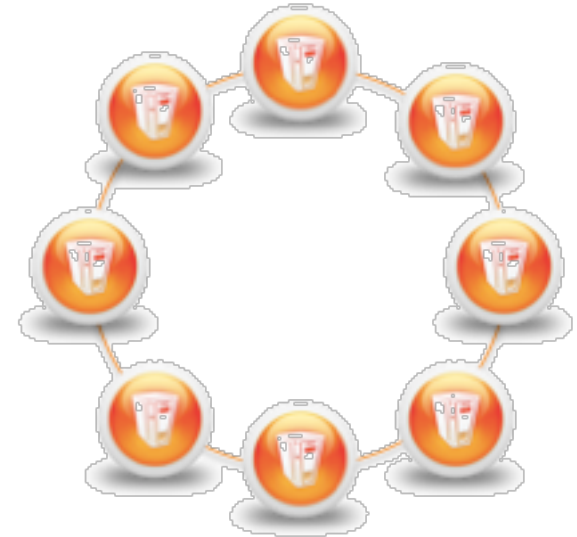


Data is organized into Partitions

1. Data is written to a Commit Log for a node (durability)
2. Data is written to MemTable (in memory)
3. MemTables are flushed to disk in an SSTable based on size.

SSTables are immutable

Tunable Data Consistency



Writes

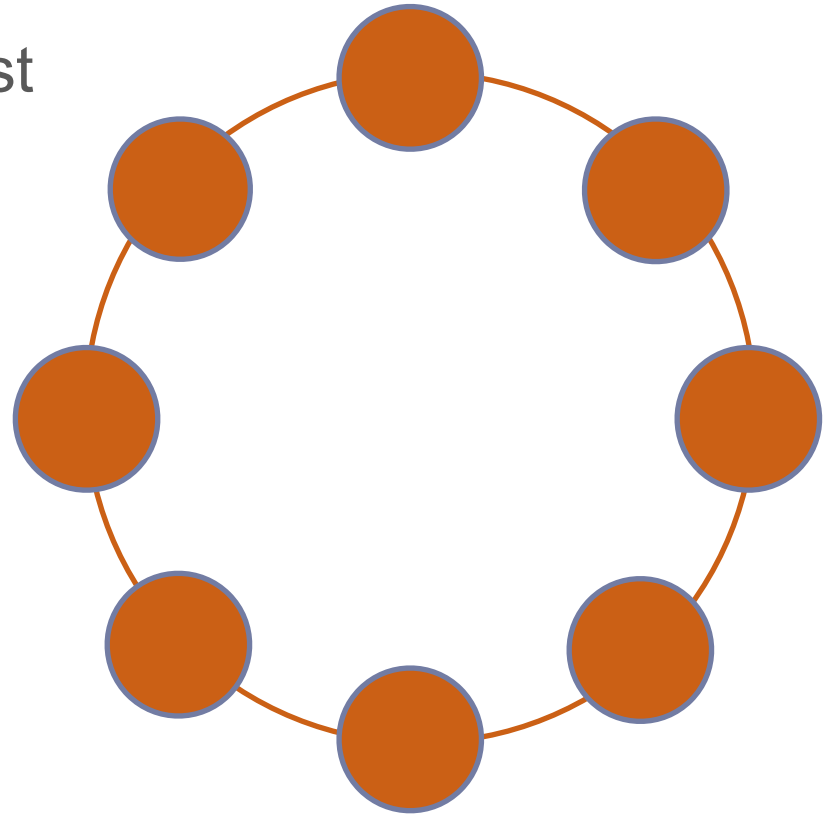
- Any
- One
- Quorum
- Local_Quorum
- Each_Quorum
- All

Reads

- One
- Quorum
- Local_Quorum
- Each_Quorum
- All

Built for Modern Online Applications DATASTAX

- Architected for today's needs
- Linear scalability at lowest cost
- 100% uptime
- Operationally simple



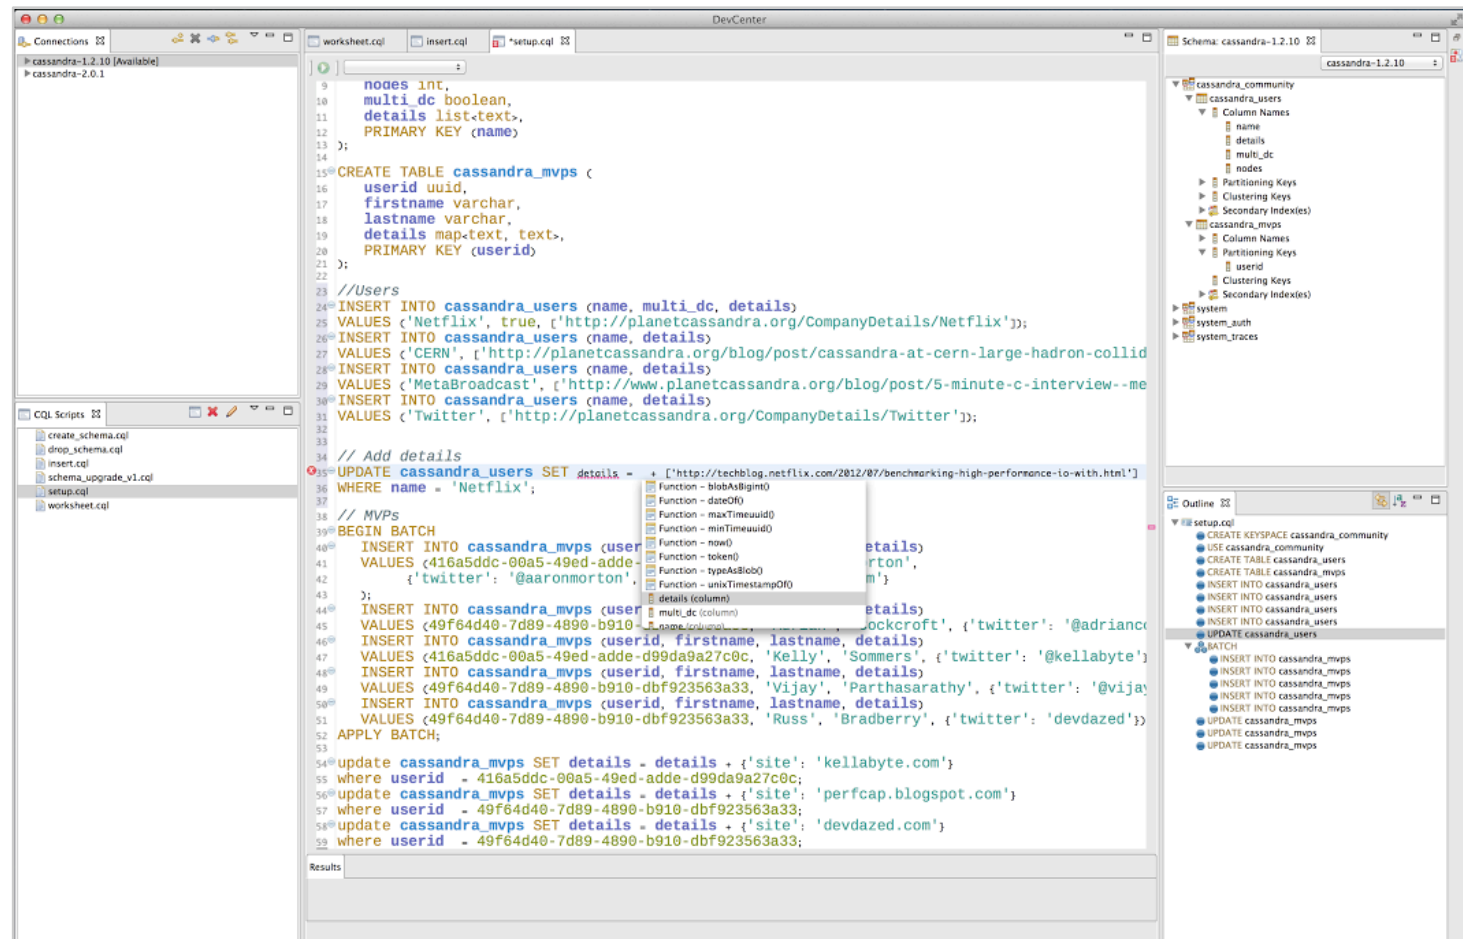
Cassandra Query Language

CQL - DevCenter



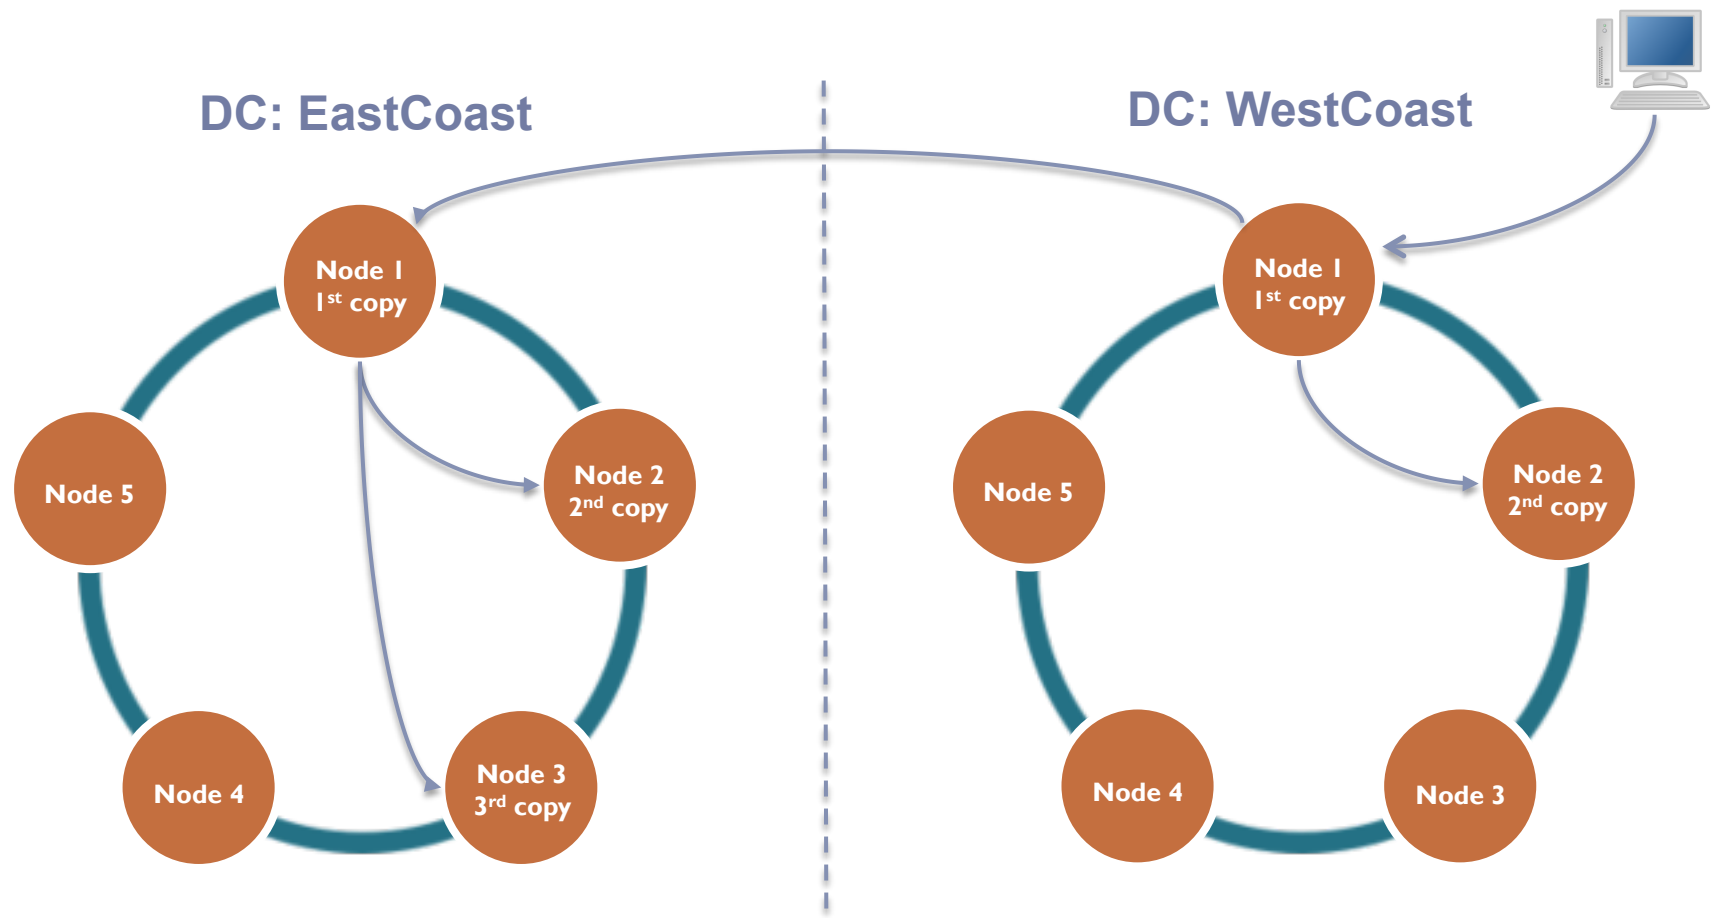
A SQL-like query language for communicating with Cassandra

DataStax DevCenter – a free, visual query tool for creating and running CQL statements against Cassandra and DataStax Enterprise.



CQL - Create Keyspace

```
CREATE KEYSPACE demo WITH REPLICATION =  
{'class' : 'NetworkTopologyStrategy', 'EastCoast': 3,  
'WestCoast': 2};
```



```
CREATE TABLE users (  
    username text,  
    password text,  
    create_date timestamp,  
    PRIMARY KEY (username, create_date desc);
```

```
INSERT INTO users (username, password, create_date)  
VALUES ('caroline', 'password1234', '2014-06-01 07:01:00');
```

```
SELECT * FROM users WHERE username = 'caroline' AND  
create_date = '2014-06-01 07:01:00';
```

Predicates

On the **partition key**: = and IN

On the **cluster columns**: <, <=, =, >=, >, IN



Collection Data Types

CQL supports having columns that contain collections of data.

The collection types include:

Set, List and Map.

```
CREATE TABLE users (  
    username text,  
    set_example set<text>,  
    list_example list<text>,  
    map_example map<int,text>,  
    PRIMARY KEY (username)  
);
```

Favor sets over list – better performance

Plus much more...

Light Weight Transactions

```
INSERT INTO customer_account (customerID, customer_email)  
VALUES ('LauraS', 'lauras@gmail.com') IF NOT EXISTS;
```

```
UPDATE customer_account SET customer_email='laurass@gmail.com'  
IF customer_email='lauras@gmail.com';
```

Counters

```
UPDATE UserActions SET total = total + 2  
WHERE user = 123 AND action = 'xyz';
```

Time to live (TTL)

```
INSERT INTO users (id, first, last) VALUES ('abc123', 'abe',  
'lincoln') USING TTL 3600;
```

Batch Statements

```
BEGIN BATCH  
  INSERT INTO users (userID, password, name) VALUES ('user2',  
'ch@ngem3b', 'second user')  
  UPDATE users SET password = 'ps22dhds' WHERE userID =  
'user2'  
  INSERT INTO users (userID, password) VALUES ('user3',  
'ch@ngem3c')  
  DELETE name FROM users WHERE userID = 'user2'  
APPLY BATCH;
```

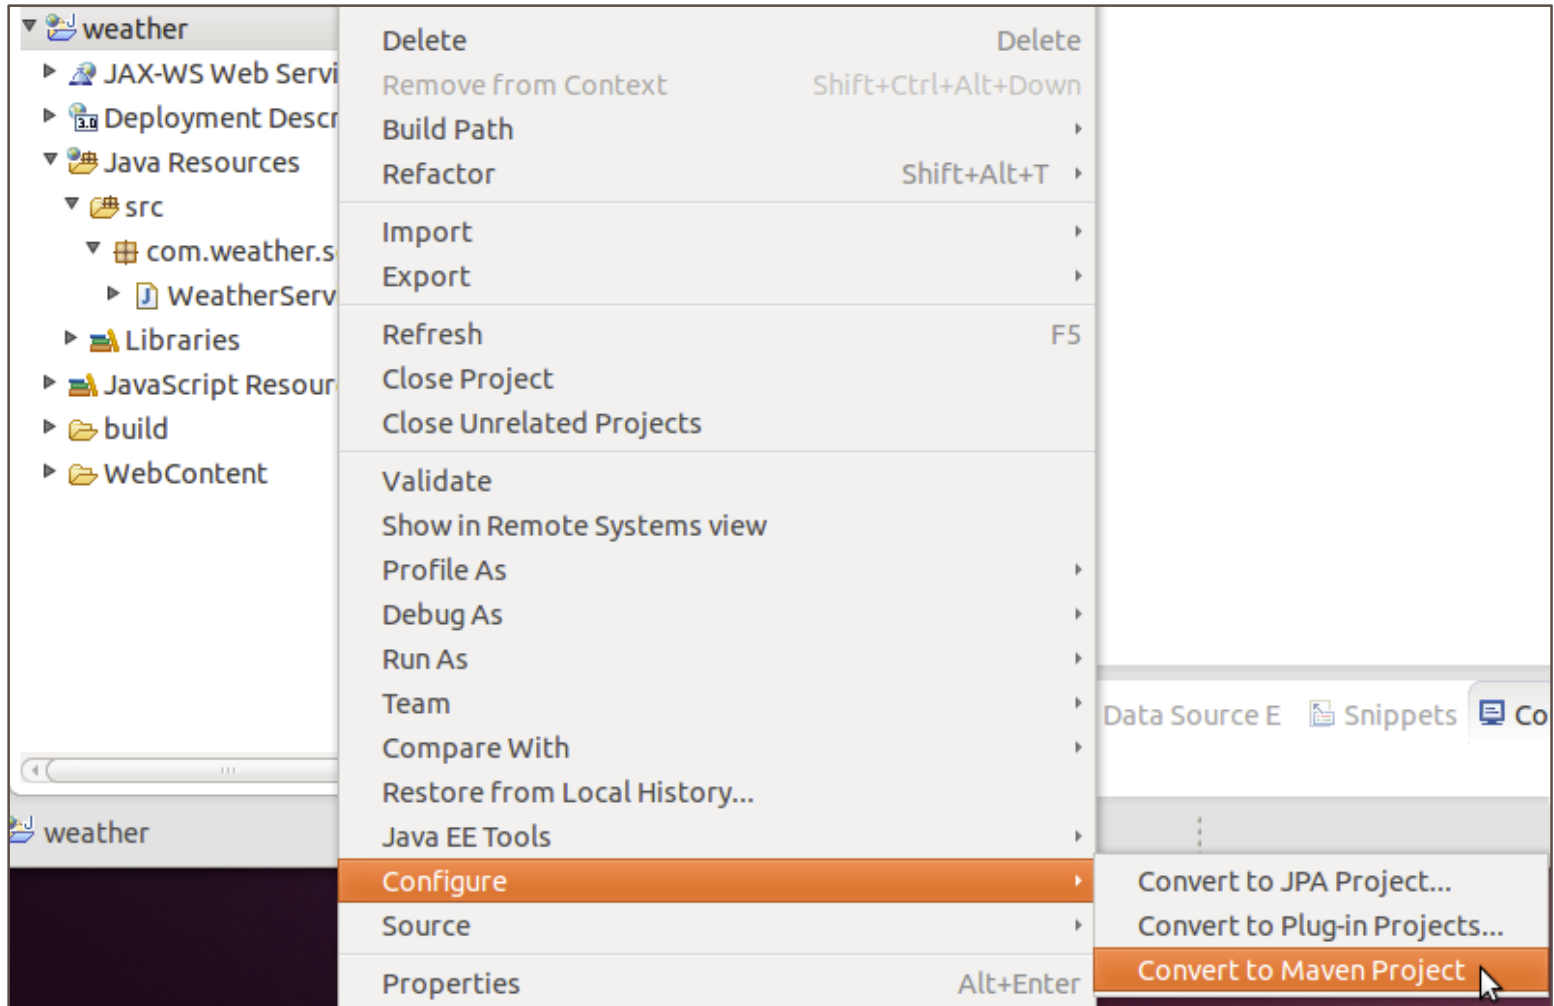
JAVA CODE EXAMPLES

- Written for CQL 3.0
- Uses the binary protocol introduced in Cassandra 1.2
- Uses Netty to provide an asynchronous architecture
- Can do asynchronous or synchronous queries
- Has connection pooling
- Has node discovery and load balancing

<http://www.datastax.com/download>

Add .JAR Files to Project

Easiest way is to do this with Maven, which is a software project management tool



Add .JAR Files to Project

In the pom.xml file, select the **Dependencies** tab

Click the **Add...** button in the left column

Enter the DataStax Java driver info

Group Id:	<input type="text" value="com.datastax.cassandra"/>
Artifact Id:	<input type="text" value="cassandra-driver-core"/>
Version:	<input type="text" value="1.0.3-dse"/> ▼

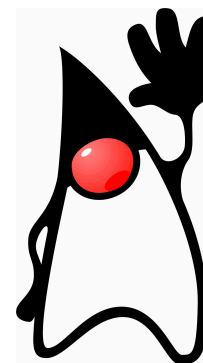
```
<dependency>
  <groupId>com.datastax.cassandra</groupId>
  <artifactId>cassandra-driver-core</artifactId>
  <version>1.0.3-dse</version>
</dependency>
```

Connect & Write

```
Cluster cluster = Cluster.builder()  
    .addContactPoints("10.158.02.40", "10.158.02.44")  
    .build();
```

```
Session session = cluster.connect("demo");
```

```
session.execute(  
    "INSERT INTO users (username, password) "  
    + "VALUES('caroline', 'password1234')"  
);
```



Note: Cluster and Session objects should be long-lived and re-used

Read from Table

```
ResultSet rs = session.execute("SELECT * FROM users");
```

```
List<Row> rows = rs.all();
```

```
for (Row row : rows) {  
    String userName = row.getString("username");  
    String password = row.getString("password");  
}
```

Asynchronous Read

```
ResultSetFuture future = session.executeAsync(  
    "SELECT * FROM users");  
  
for (Row row : future.get()) {  
    String userName = row.getString("username");  
    String password = row.getString("password");  
}
```

Note: The future returned implements Guava's `ListenableFuture` interface. This means you can use all Guava's `Futures`¹ methods!

¹<http://docs.guava-libraries.googlecode.com/git/javadoc/com/google/common/util/concurrent/Futures.html>

Read with Callbacks

```
final ResultSetFuture future =  
    session.executeAsync("SELECT * FROM users");  
  
future.addListener(new Runnable() {  
  
    public void run() {  
        for (Row row : future.get()) {  
            String userName = row.getString("username");  
            String password = row.getString("password");  
        }  
    }  
}, executor);
```

Parallelize Calls

```
int queryCount = 99;

List<ResultSetFuture> futures = new
ArrayList<ResultSetFuture>();

for (int i=0; i<queryCount; i++) {
    futures.add(
        session.executeAsync("SELECT * FROM users "
            +"WHERE username = '"+i+"'"));
}

for(ResultSetFuture future : futures) {
    for (Row row : future.getUninterruptibly()) {
        //do something
    }
}
```

Prepared Statements

```
PreparedStatement statement = session.prepare(  
    "INSERT INTO users (username, password) "  
    + "VALUES (?, ?)");
```

```
BoundStatement bs = statement.bind();
```

```
bs.setString("username", "caroline");  
bs.setString("password", "password1234");
```

```
session.execute(bs);
```

Query Builder

```
Query query = QueryBuilder
    .select()
    .all()
    .from("demo", "users")
    .where(eq("username", "caroline"));
```

```
ResultSet rs = session.execute(query);
```

Determine which node will next be contacted once a connection to a cluster has been established

```
Cluster cluster = Cluster.builder()  
    .addContactPoints("10.158.02.40", "10.158.02.44")  
    .withLoadBalancingPolicy(  
        new DCAwareRoundRobinPolicy("DC1"))  
    .build();
```

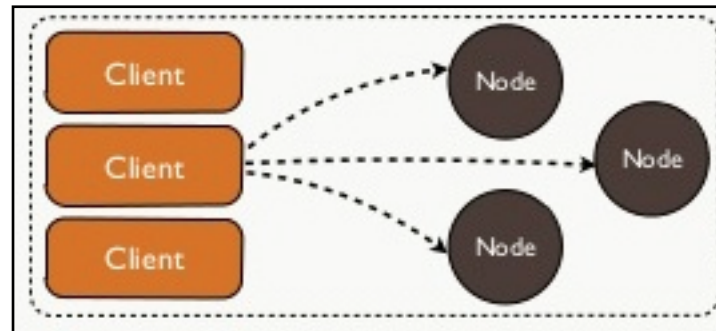
Name of the local DC



Policies are:

- **RoundRobinPolicy**
- **DCAwareRoundRobinPolicy** (default)
- **TokenAwarePolicy**

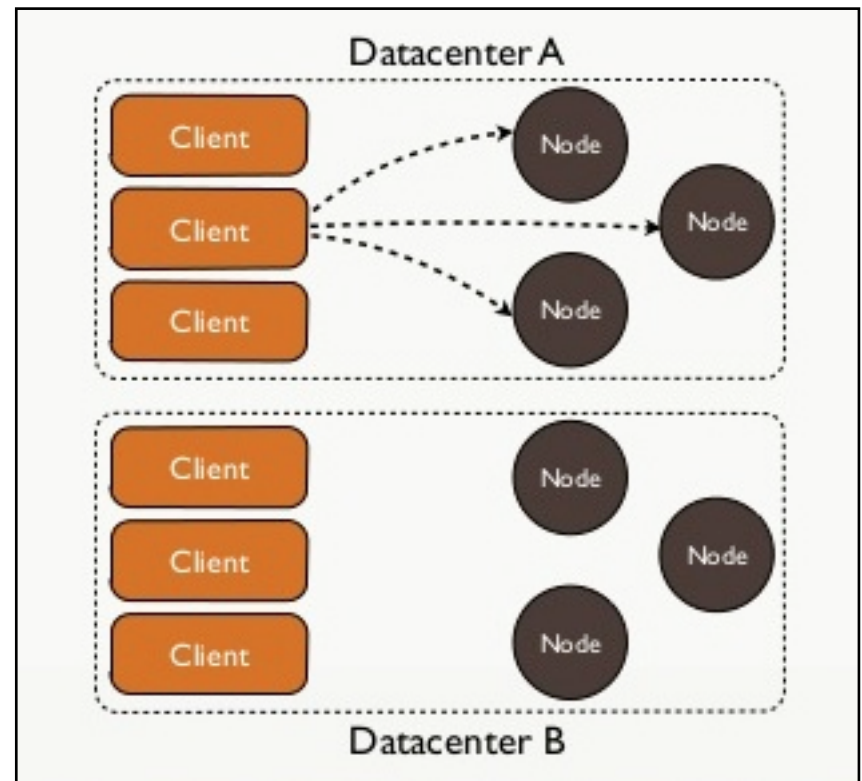
- Not data-center aware
- Each subsequent request after initial connection to the cluster goes to the next node in the cluster



- If the node that is serving as the coordinator fails during a request, the next node is used

DCAwareRoundRobinPolicy

- Is data center aware
- Does a round robin within the local data center
- Only goes to another data center if there is not a node available to be coordinator in the local data center



- Is aware of where the replicas for a given token live
- Instead of round robin, the client chooses the node that contains the primary replica to be the chosen coordinator
- Avoids unnecessary time taken to go to any node to have it serve as coordinator to then contact the nodes with the replicas

- Community Site
(<http://planetcassandra.org>)
- Documentation
(<http://www.datastax.com/docs>)
- Downloads
(<http://www.datastax.com/download>)
- Getting Started
(<http://www.datastax.com/documentation/gettingstarted/index.html>)
- DataStax
(<http://www.datastax.com>)



ABOUT DATASTAX

About DataStax



300+

Employees



Founded in April 2010

Santa Clara, Austin, New York, London

30

Percent

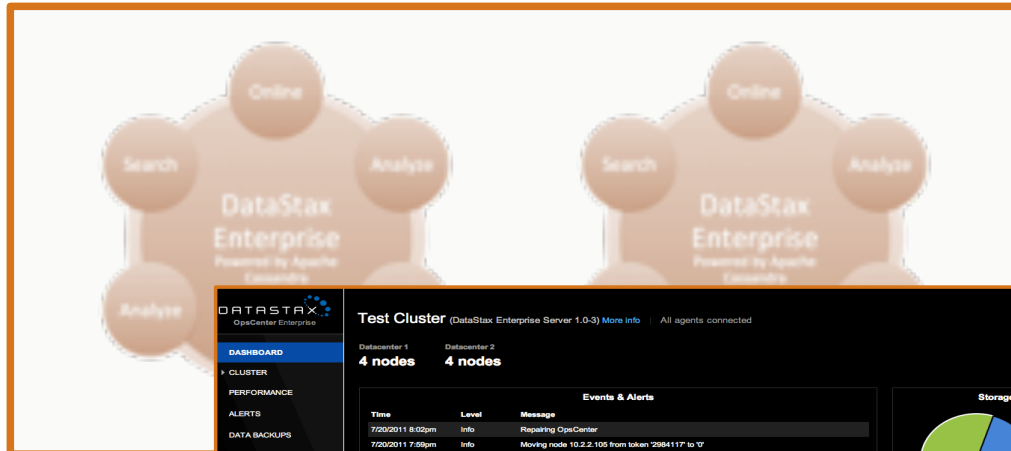


500+

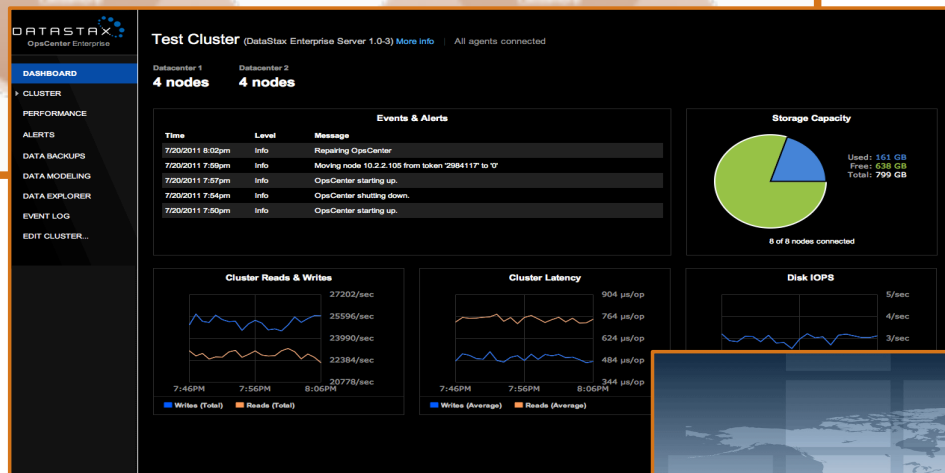
Customers



DataStax delivers Apache Cassandra to the Enterprise



Certified /
Enterprise-ready Cassandra



Visual Management &
Monitoring Tools

24x7 Support & Training



CASSANDRASUMMIT2014

September 10 - 11 | #CassandraSummit



SAN FRANCISCO

WORLD'S LARGEST GATHERING
OF CASSANDRA DEVELOPERS.

FREE ADMISSION.
SERIOUSLY.

INTRODUCING **DATASTAX ENTERPRISE 4.5**

The World's Fastest, Most Scalable
Distributed Database Technology

[DOWNLOAD NOW »](#)

[PRESS RELEASE »](#)



Thank You!

cgilmore@datastax.com
