

OverOps

Continuous Reliability:

- **An Agile Process to Deliver
Higher Quality Applications**


- Rich Borucki

Agenda

- What is Continuous Reliability?
- Current Challenges Delivering Quality Applications
- The Path to Continuous Reliability
- Q&A

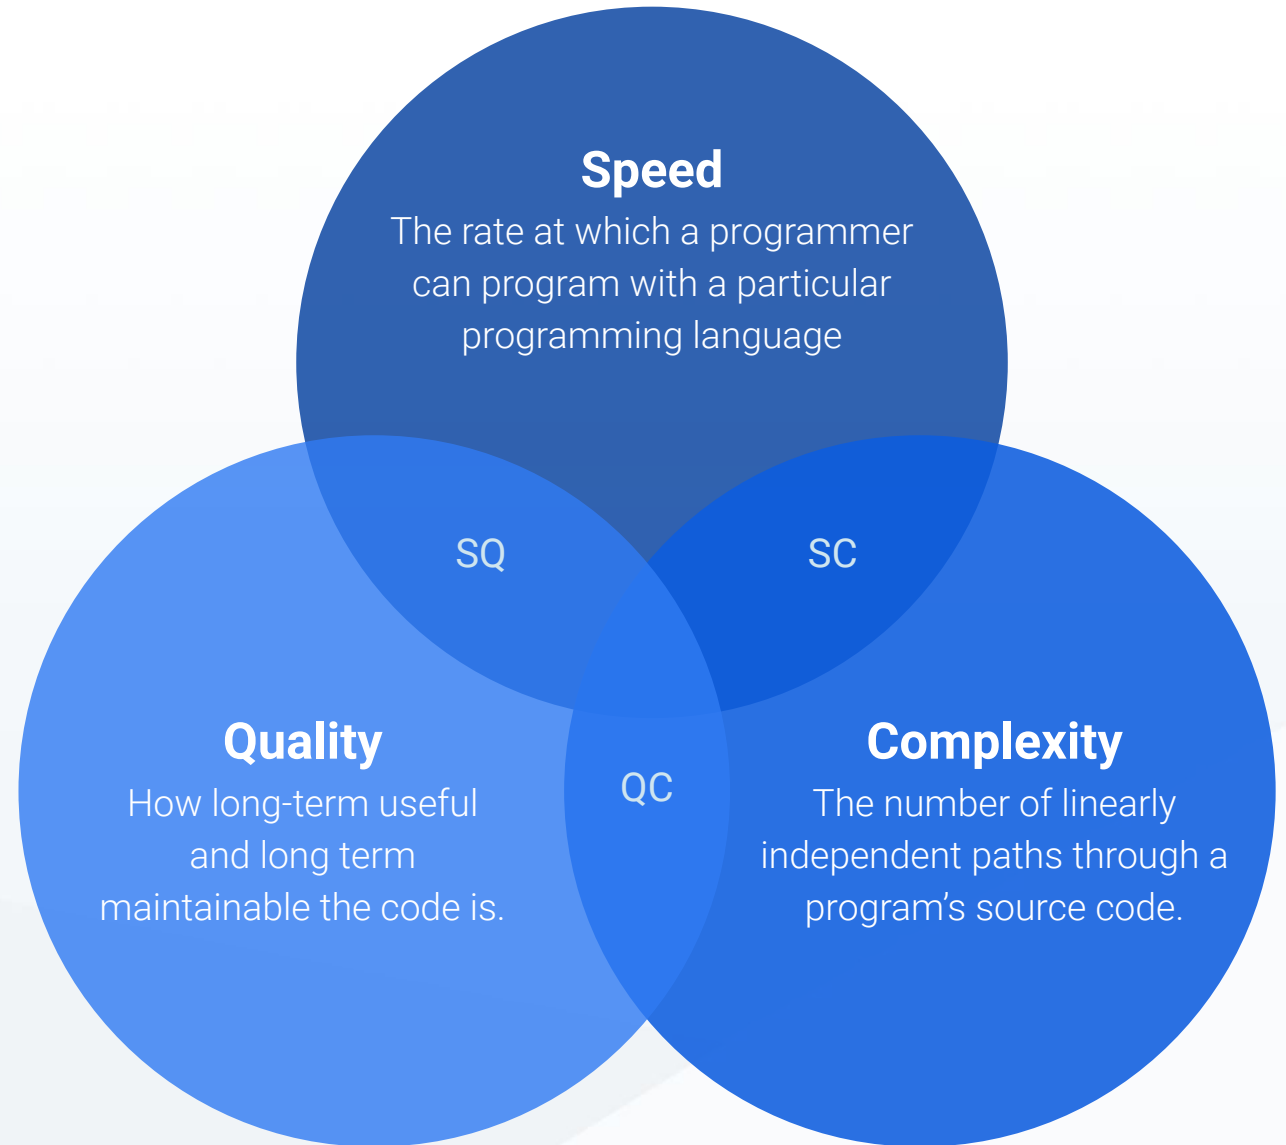
What is Continuous Reliability?

Reliability is the **probability** of a piece of software operating without failure while in a specified environment over a **set duration of time**. In a perfect world, a reliable piece of software is completely defect free, does not create downtime and performs correctly in every scenario.



What is Continuous Reliability?

Continuous Reliability is the notion of balancing speed, complexity and quality by taking a continuous, proactive approach to reliability across the software delivery life cycle. It is achieved through the creation of data-driven quality gates and feedback loops.



Current Challenges

Applications Aren't Truly Tested Until Production



There is a lot of dev humor about this...

Reliability Challenges



Technical debt



3rd-party code



Finger-pointing



No way to prioritize



Unknown error number



Lack of context

Legacy Tools & Processes

Tools/Processes that put application quality at risk:

- Manual code reviews – peer reviews
- Manual testing – QA team writes and runs test cases
- Basic unit testing of methods – Devs write tests against methods
- Manual deployments
 - Wait for the QA region to become available
 - Manual build and deploy of applications

Questions to Ask Ourselves

How do we know when to promote code?

- “Head of QA gives the green light”
- “Query APM or Logs and see if there are less errors than last build”
- “Don’t have a choice because management wants features released”

How do we know if we have introduced new or resurfaced errors?

- “Don’t really know...all tests passed”
- “We do a lot of regex queries against the logs and we think we know”

How do we know release-over-release if our code is improving?

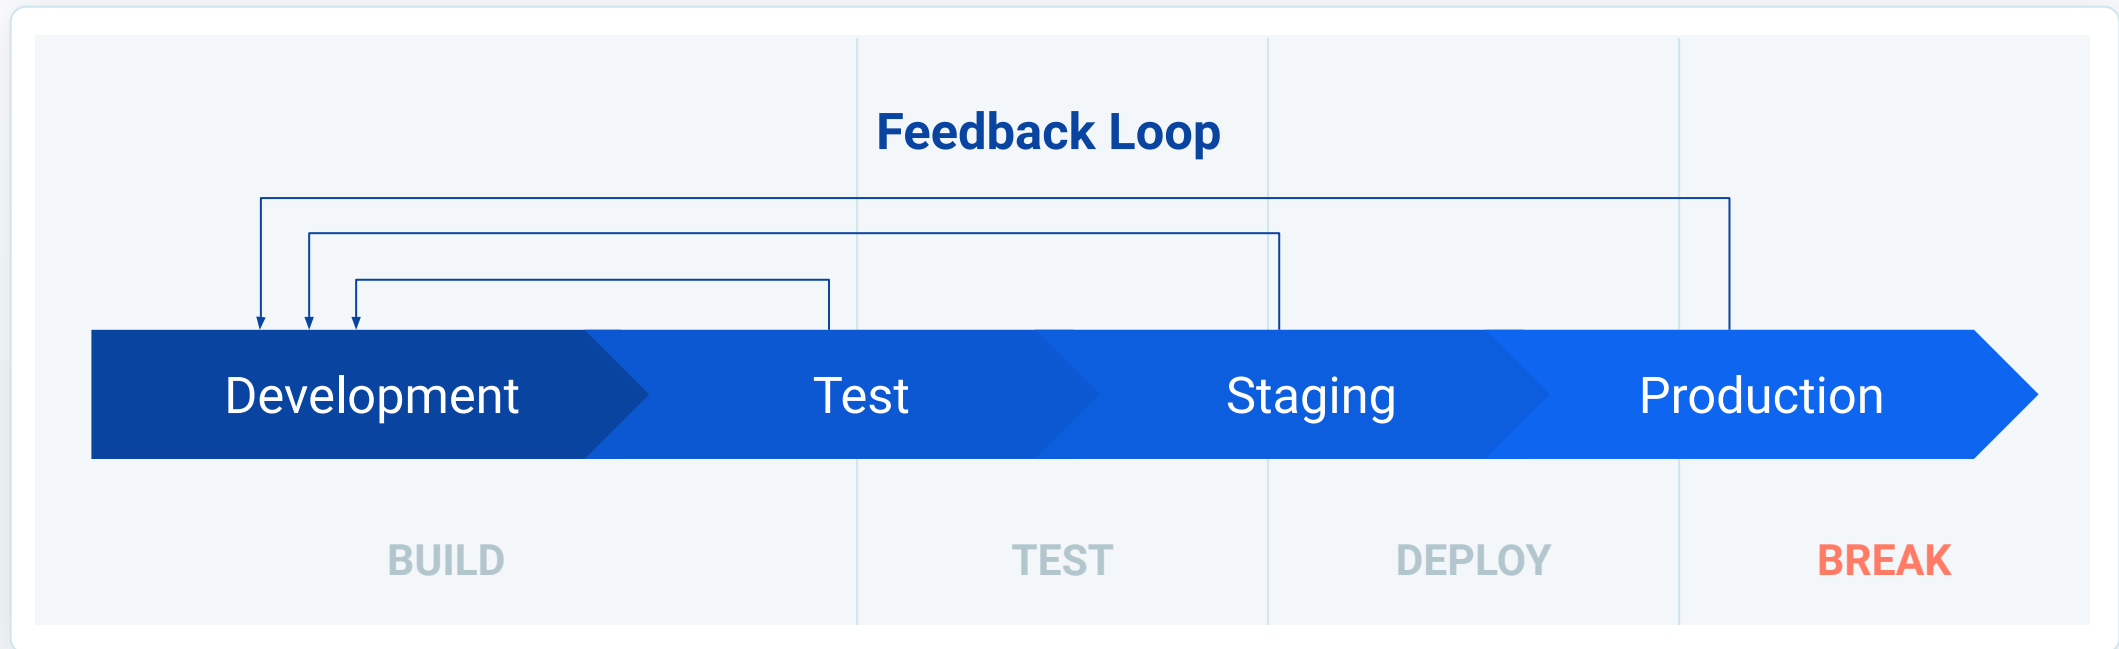
- “Less customer complaints”
- “Less Sev1 issues”
- “Don’t really know”

The Path to Continuous Reliability

We Need to Build Smarter CI/CD Pipelines

Developer feedback, alerting, routing, context, metrics

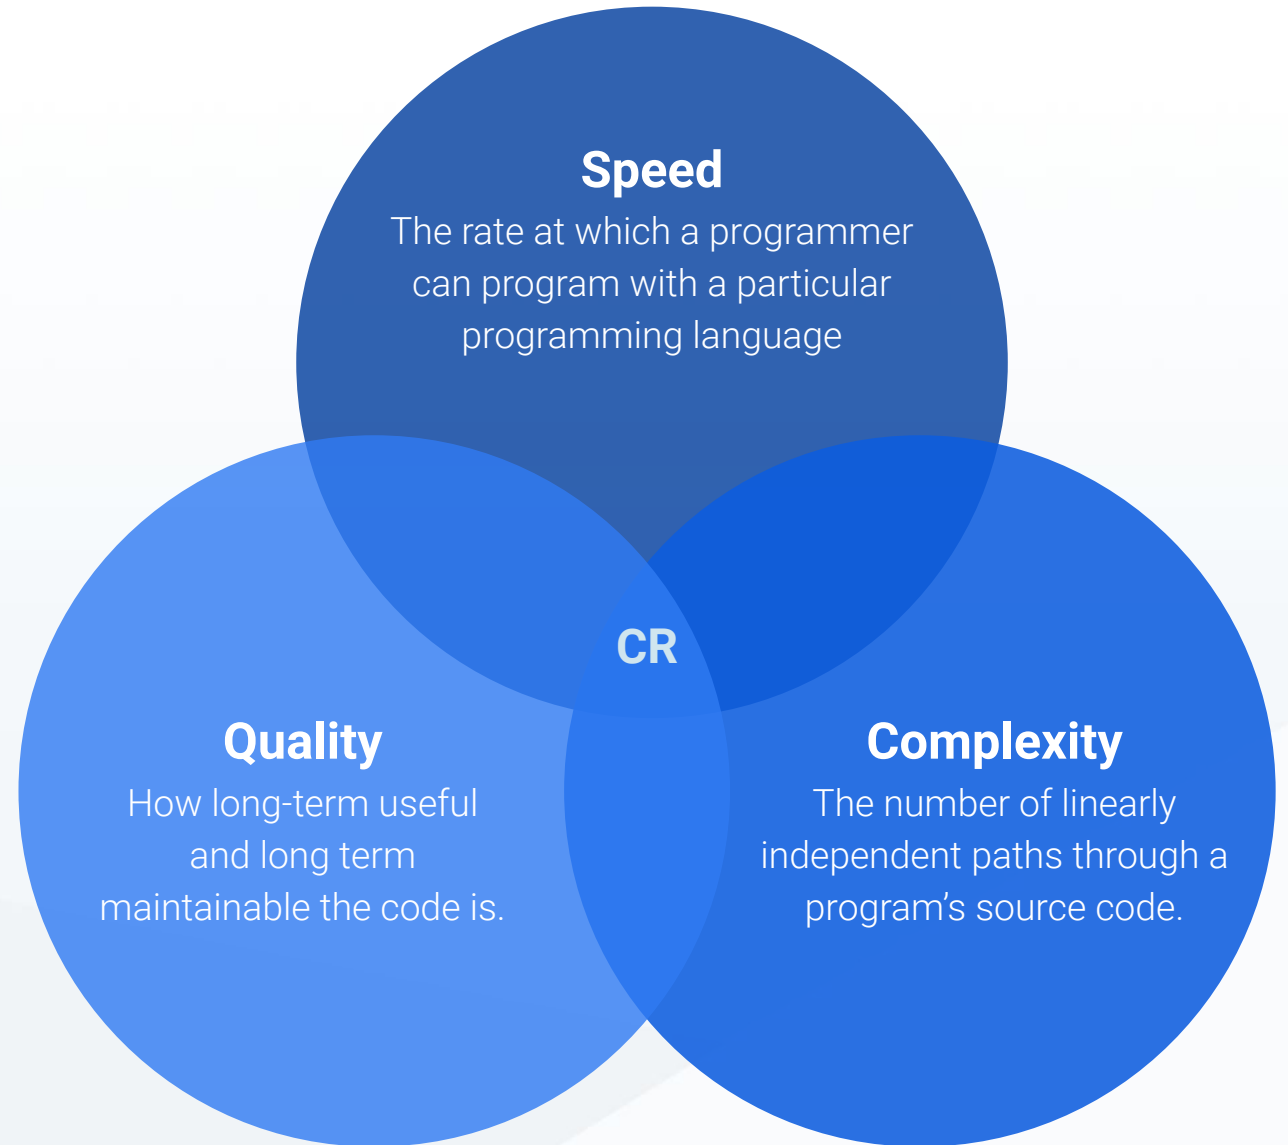
Automate as much as possible to eliminate human intervention



CI/CD Workflow

Continuous Reliability (CR) Roadmap

- **Automate** as much as possible to eliminate human intervention
- **Know** your unknowns
- **Identify** errors sooner – shift left
- **Capture** more context to provide feedback to developers
- **Create** more/better code coverage
- **Collect** better data – metrics matter
- **Build** a culture of accountability



Automate...Automate...Automate

- Static code analysis tools – automated approach
 - Goal is to identify weaknesses in code that might lead to vulnerabilities before testing
- Mix of JUnit, JMeter, Selenium, Load Runner, etc.
 - New tools emerging to identify gaps in code coverage
- Limited to no manual testing
- Automated builds and deployments (CI/CD)
 - QA environment auto build and teardown
- Dynamic code analysis tool – automated way to detect, deduplicate, alert and measure errors while code is being exercised
- Feedback loops as code moves through the pipeline
 - Quality gates identifying issues that stops a build from getting promoted

Know Your Unknowns



Know Your Unknowns

- Critical exceptions – Nullpointers, IndexOutOfBoundsException, etc.
- Slowdowns – which processes are running slower
- Errors that are happening more frequently
- Unhandled and swallowed exceptions – don't show up in the logs
 - Little known fact – it takes 50-100 microseconds to generate the exception

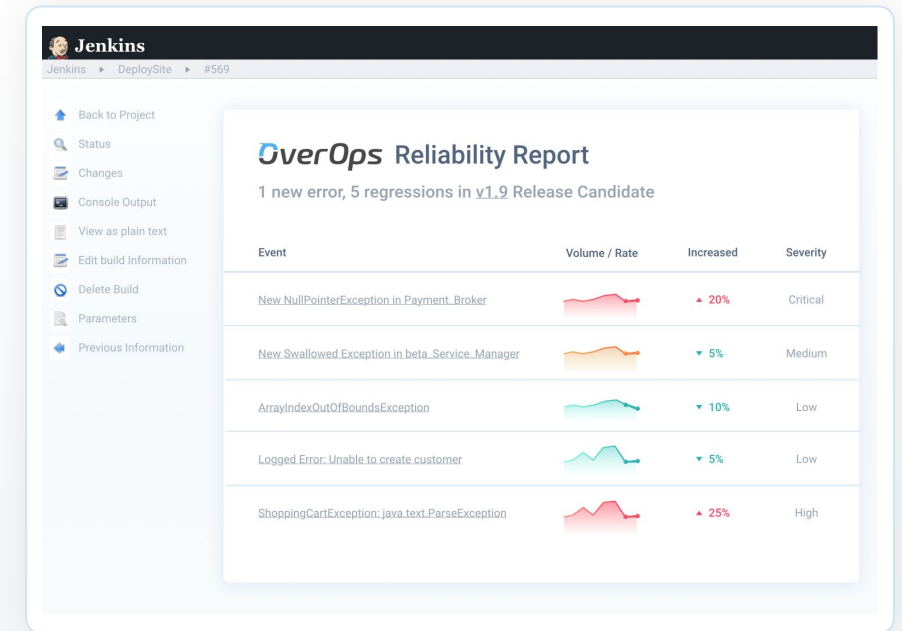
```
STRING VALUE == NULL;
TRY {
    IF (VALUE.EQUALSIGNORECASE ("HELLO") )
    {
        // DO SOMETHING IF THEY MATCH
    }
} CATCH {EXCEPTION E} {
    // TODO: HANDLE EXCEPTION
    // I FORGOT TO HANDLE THIS....OOPS
}
```

```
BOOLEAN PASS == TRUE;

TRY {
    INTEGER.PARSEINT (VALUE);
} CATCH {EXCEPTION E} {
    PASS == FALSE;
}
```

Shift Left

- Test early and often – daily, hourly, at code check-in, etc.
- Implement quality gates to ensure code quality
 - Static and dynamic gates
 - Instant feedback to Dev team



Shift Left

=



Early Detection



Cost Effectiveness



Time Saving



Release Smooth

Context Beyond the Stack Trace

Shameless OverOps plug


- Dynamic code analysis
- Detects 100% of all errors
 - Even errors not in the logs
- Provides metrics
 - Error counts and rates
- Detects slowdowns
- Detects increasing errors

The screenshot displays the OverOps interface. On the left, a 'Logged Error' sidebar lists several errors, with 'DBManager.removeObjectsInList' selected. The main area shows the Java code for this method. A blue line highlights the error location, and a callout box shows the exception details: 'IllegalStateException' with the message 'Connection pool shut d...'. On the right, a 'Record Variables' panel lists variables and their values, including 'Thread-local state', 'Log message', 'currentEX', '(Object ID)', 'cause', 'detailMessage', 'stackTrace', 'suppressedExcep...', 'failedBatch', 'failedBatches', and 'failedItemsCount'.

Record Variables	
Thread-local state	4 items
Log message	Error deleting batch, 1 item...
currentEX	IllegalStateException
(Object ID)	9
cause	IllegalStateException
detailMessage	"Connection pool shut d...
stackTrace	Object[]
suppressedExcep...	Collection[Unmodifiabl...
failedBatch	DynamoDBMapper\$Fail...
failedBatches	Object[]
failedItemsCount	1

Create More/Better Code Coverage

- Implement new tools to help detect code coverage gaps
- Leverage context from tools like OverOps to help create tests
 - Context of code path, data, environment, etc.



Metrics Matter

- Build a metrics hub
 - Error data from applications, databases, middleware, queues, etc.
 - System metrics, CPU, memory, GC, database connections, file handles, threads, etc.
- Correlate data to identify anomalies, build self-healing patterns, identify potential issues earlier
- Use machine learning to look at past issues to predict future issues
- Know release-over-release code quality
- Application health scorecard – weighted average of errors & slowdowns
- Know who your best developers are

Env

OverOps Prod: S6295 ▾

Apps

All ▾

Server Groups

All ▾

Deployments

All ▾

Type

All ▾

Transactions

All ▾

Top

10 ▾

🗪

☰

Deployments

.v4.0.2190*	🟡🟡🟡🟡🟡🟡🟡🟡🟡🟡🟡🟡
.v4.0.2189	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.v4.0.2186	🟡🟡🟡🟡🟡🟡🟡🟡🟡🟡🟡🟡

Applications

.api	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.app	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.service	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.sqs-aa	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.sqs-sql	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.sqs-stats	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢

Code Tiers

.AWS*	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.AWS-DynamoDB*	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.JDBC*	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.Application	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.Apache	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.Google	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.Apache-http	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.Guava-cache	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.Java-IO	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢
.Java-Net	🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢🟢

Drill-down

Deploy	Diff	New	Inc	Slow	Avg Score
<u>v4.0.2190*</u>	↔		<u>2 p1</u>	<u>2 (1 p1)</u>	82.5
<u>v4.0.2189</u>	↔	<u>7</u>			92.4
<u>v4.0.2186</u>	↔	<u>21</u>	<u>5 p1</u>	<u>8 (5 p1)</u>	82.9

Drill-down

App	New	Inc	Slow	Avg Score
<u>api</u>	<u>10 (1 p1)</u>		<u>2 p1</u>	94.6
<u>app</u>	<u>7</u>	<u>3 p1</u>	<u>2 p1</u>	93.9
<u>service</u>	<u>4</u>	<u>5 (4 p1)</u>	<u>1</u>	95.0
<u>sqs-aa</u>	<u>6</u>	<u>1 p1</u>	<u>2 (1 p1)</u>	96.1
<u>sqs-sql</u>	<u>3</u>		<u>1 p1</u>	98.2
<u>sqs-stats</u>			<u>1 p1</u>	99.3

Drill-down

Tier	New	Inc	Avg Score
<u>AWS*</u>	<u>3</u>		97.9
<u>AWS-DynamoDB*</u>			100.0
<u>JDBC*</u>			100.0
<u>Application</u>	<u>25 (1 p1)</u>	<u>6 (5 p1)</u>	86.8
<u>Apache</u>		<u>1 p1</u>	99.3
<u>Google</u>	<u>1</u>		99.6
<u>Apache-http</u>			100.0
<u>Guava-cache</u>			100.0
<u>Java-IO</u>			100.0
<u>Java-Net</u>			100.0

Env OverOps Prod: S6295 Apps All Server Groups All Deployments All Type All Transactions All

Reliability Score

83.6

Failed Transactions

0.01%

Slow Transactions

5

New Errors

24

Error Increases

5


Unique Errors

665

Transaction Throughput



Errors By App



Transaction Response (ms)



Exceptions By Tier



Transaction Errors



Critical Exceptions



Culture of Accountability

It has become acceptable to have lots of errors in applications. When issues happen, we play the blame game, resulting in outages, Sev1 issues, customer incidents, etc.

What to do?

Change the status quo

- Bug bounties, bonuses for applications that meet SLAs
- Error burn down time for each sprint
- Set goals for code quality in your team

Hold everyone accountable

- Correlate all errors to commits and build a leaderboard
- Require the dev who wrote the code to fix it

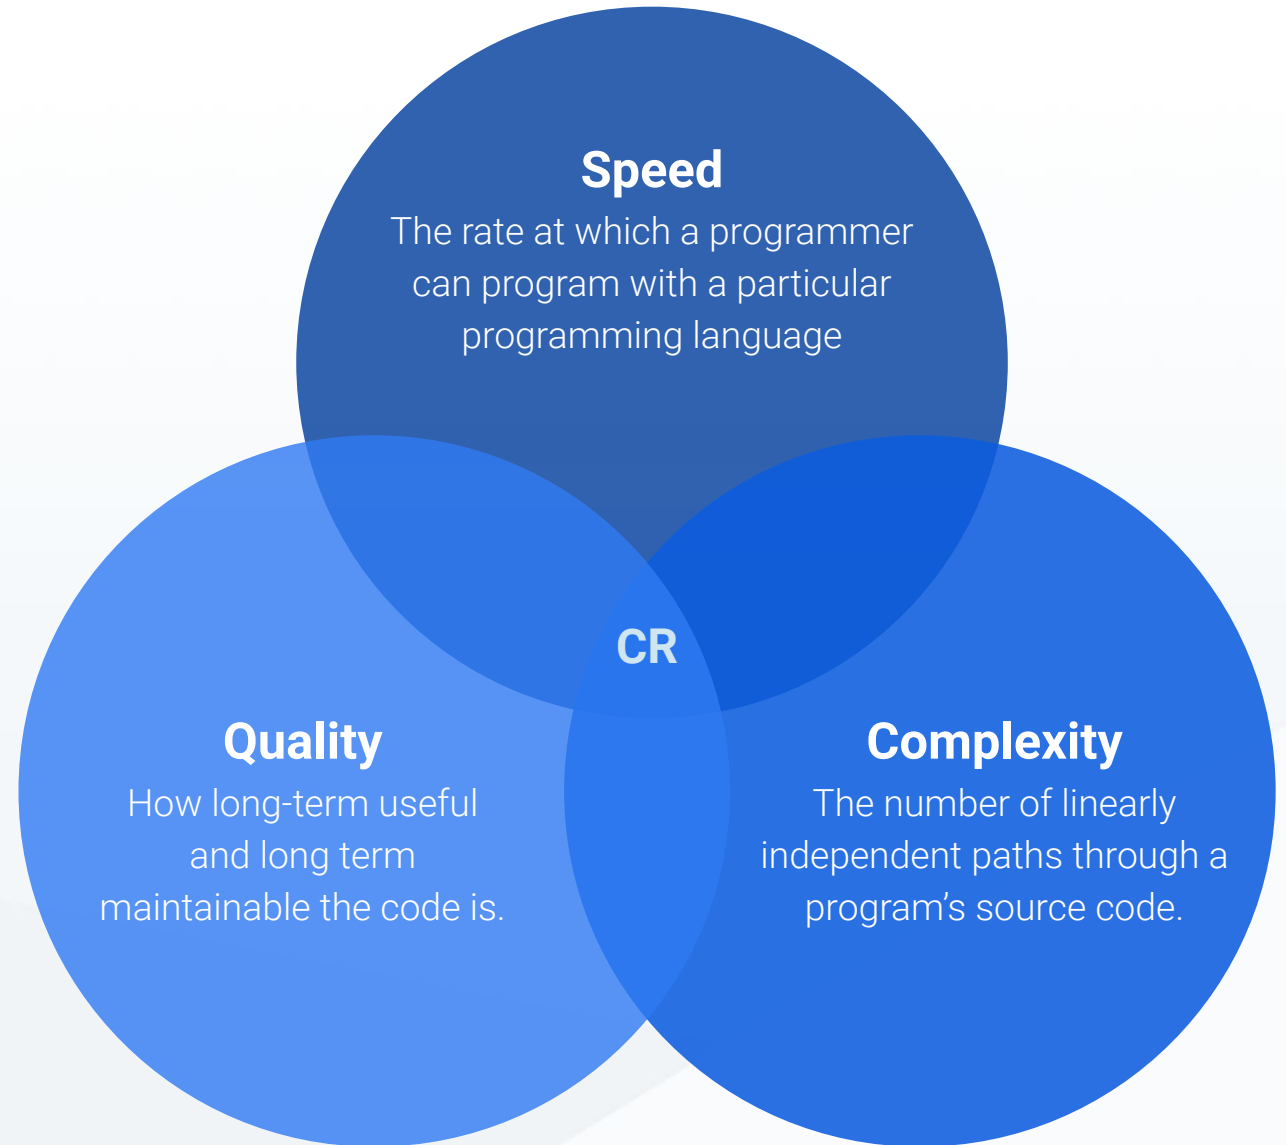
Leverage tools like OverOps that detect 100% of errors and work to fix them all

- Don't allow code to be promoted that has errors



Continuous Reliability (CR) Roadmap

- **Automate** as much as possible to eliminate human intervention
- **Know** your unknowns
- **Identify** errors sooner – shift left
- **Capture** more context to provide feedback to developers
- **Create** more/better code coverage
- **Collect** better data – metrics matter
- **Build** a culture of accountability



OverOps



Questions?