

RMI and CORBA

Why both are valuable tools

by

Brian Gilstrap

with thanks to Mark Volkmann

Introduction

- This talk is about RMI and CORBA, not “RMI vs CORBA”
- It’s about using the right tool for the job
- RMI and CORBA have different “design centers”; as a result, they are good for solving different (but overlapping) sets of problems
- This talk centers around Java

RMI & CORBA Characteristics

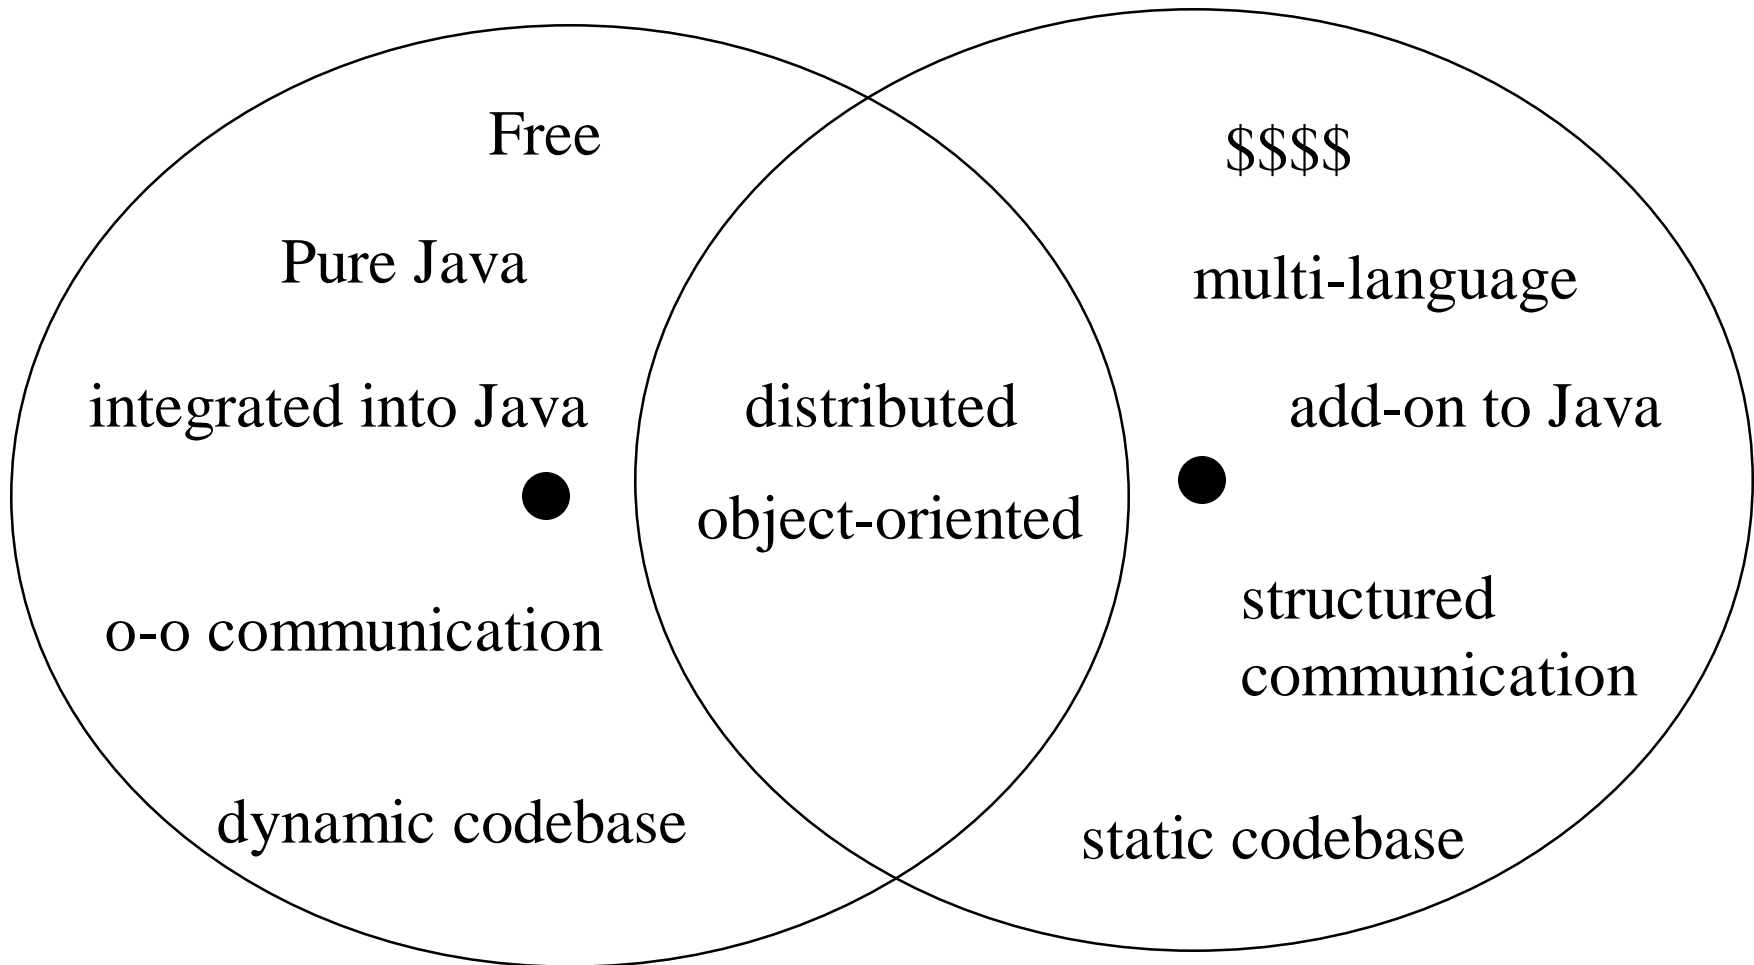
RMI

- Pure Java
- Free
- Simple to use
- Integrated into platform
- Object-oriented communication
- Dynamic codebase
- Performance is not as good (depending on application)
- Less mature (but learned from CORBA issues/mistakes)
- Maturing quickly

CORBA

- Multi-language
- \$\$\$
- Harder to use
- Add-on
- Structured communication
- Static codebase
- Performance is better (depending on application)
- More mature (but 1st of its kind in the industry)
- Maturing slowly

These characteristics result in Design Centers



RMI features/limitations

- Pure Java
 - RMI is pure Java (no other languages), but it supports distributed OO computing in Java very well
 - Can still mix in socket communication and native code (JNI) to access non-Java (e.g. legacy) systems
 - this may be easier or harder than using CORBA, depending on the problem

RMI features/limitations (cont)

- Free
 - comes as part of the Java Development Kit (JDK) and is part of the Java Runtime Environment (JRE)
 - no runtime costs, developer costs, etc.

RMI Features/Limitations (Cont)

- Simple to use - providing remote access
 - extend “Remote” interface & handle exceptions
 - use the “rmic” utility to generate stubs & skeletons (produces ‘.class’ files)
 - registry provides bootstrap for finding remote objects
 - must know host

RMI Features/Limitations (Cont)

- Simple to use - sending objects
 - implement “Serializable” interface
 - optionally specify codebase
 - automatic version mismatch detection

RMI Features/Limitations (Cont)

- Integrated into platform
 - standard way of using RMI
 - source code is the same for any platform
 - bytecode is cross-platform just like other Java
 - remote method calls are nearly identical to normal method calls
 - distributed garbage collection preserves automatic memory management

RMI Features/Limitations (Cont)

- Object-oriented communication
 - can easily send references to remotely accessible objects to other VMs
 - using the “Remote” interface
 - can easily send copies of objects to other VMs
 - using “Serializable” interface
 - can get more sophisticated
 - overriding “read/writeObject”
 - using “Externalizable”

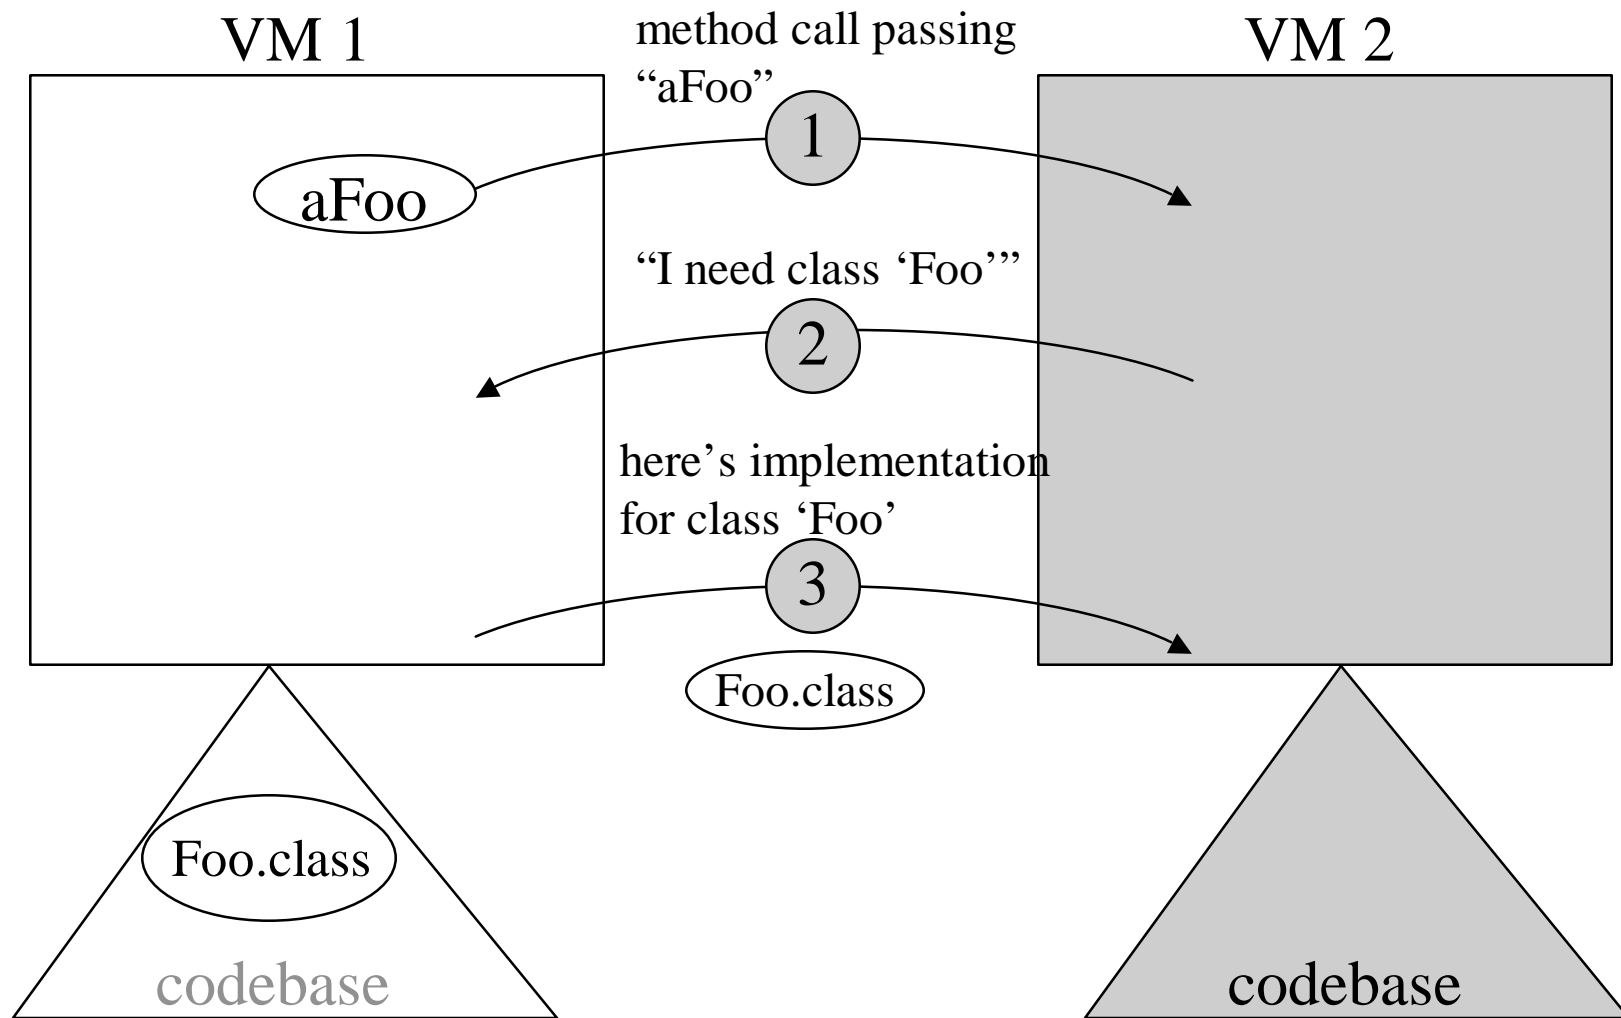
RMI Features/Limitations (Cont)

- Object-oriented communication (continued)
 - distributed garbage collection
 - allows server VMs to garbage collect unused objects/resources when those objects/resources are no longer needed

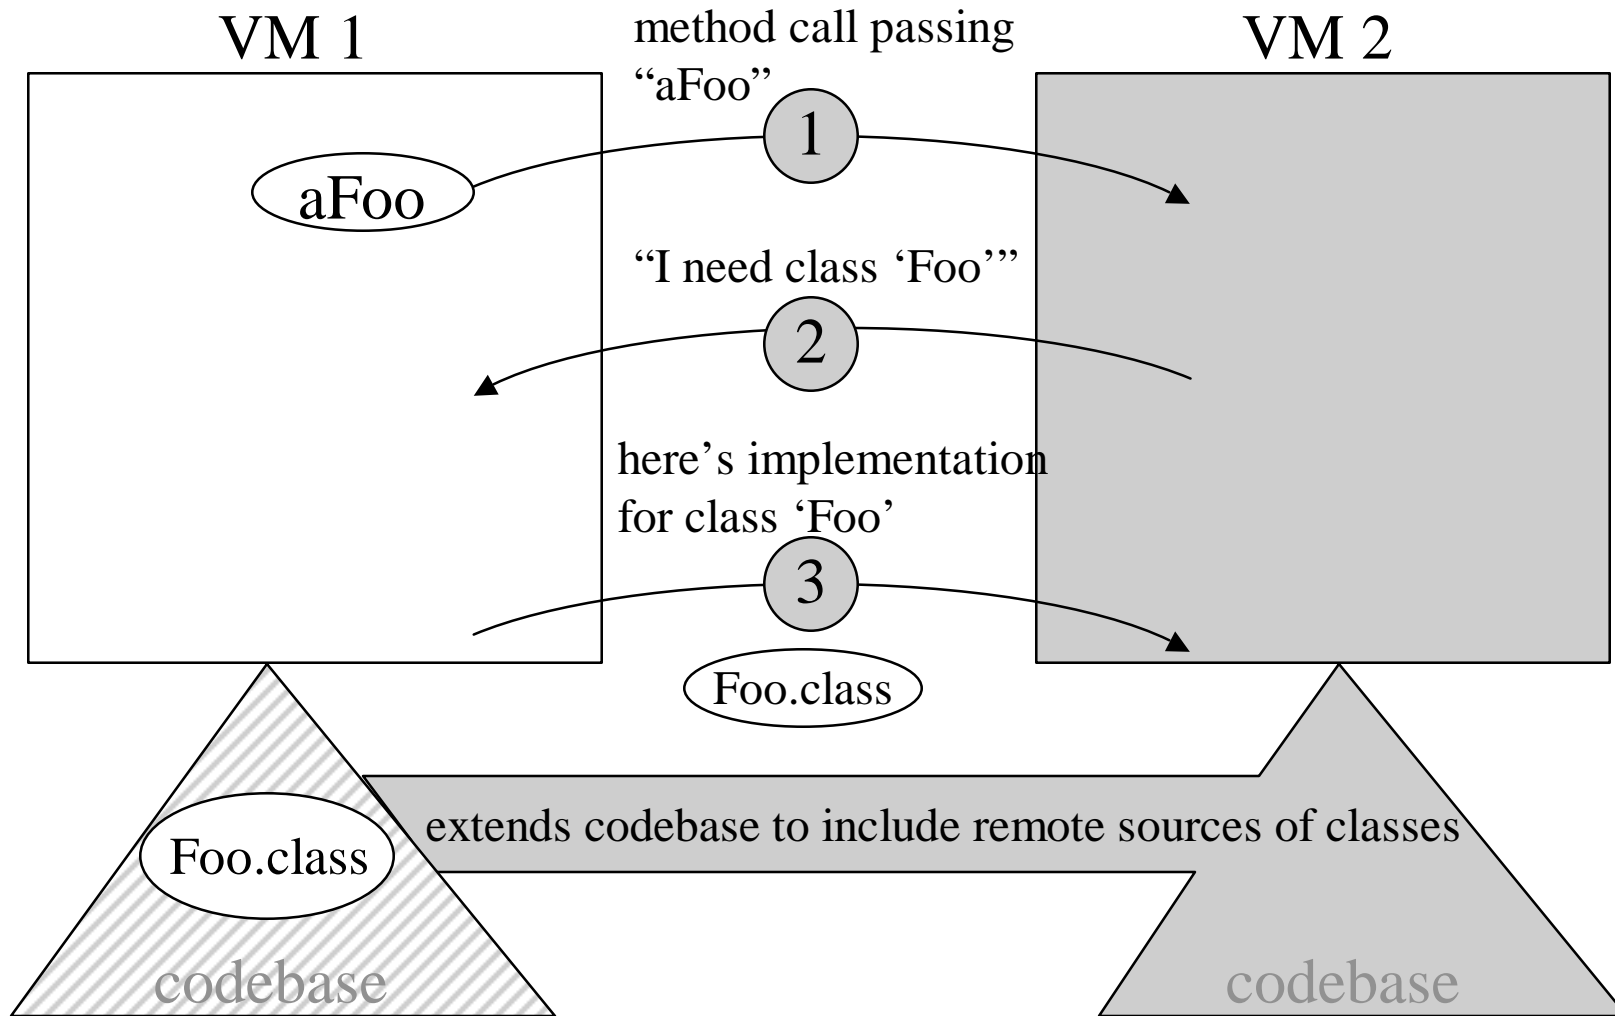
RMI Features/Limitations (Cont)

- Dynamic codebase
 - can send not only the state information (field values) but also the implementation (bytecode) of objects across the network
 - very powerful
 - send code to be executed by the server as part of performing work
 - example: a search object for searching a datastore
 - send an agent to a remote site to do work
 - send an object to a server to be stored persistently
 - etc.

RMI Features: Dynamic Codebase



RMI Features: Dynamic Codebase (cont)



RMI Features/Limitations (Cont)


- Performance
 - for designs which pass data only (no objects), and for which version mismatch detection is not needed, RMI performs more slowly than CORBA
 - exact numbers depend upon many things
 - platform
 - Java JIT/no-JIT
 - ORB vendor
 - problem domain
 - etc.
 - for designs which pass objects, RMI is currently the only ticket in town (CORBA can't pass objects)

RMI Features/Limitations (Cont)

- Performance (continued)
 - performance can be improved by custom serialization code
 - allows default implementation with targeted tuning for the 80/20 rule
 - BUT
 - using RMI features to solve problems produces a different design and can reduce or eliminate performance differences
 - passing objects, version mismatch, and distributed garbage collection are important issues when building distributed systems

RMI Features/Limitations (Cont)

- RMI is Maturing quickly



CORBA Features/Limitations

- Multi-language
 - C (requires layering object-oriented concepts on top)
 - C++
 - Smalltalk
 - Java - not yet standardized
 - COBOL
 - Ada

CORBA Features/Limitations (cont)

- Not free
 - costs can be substantial
 - developer licenses (-\$800+ per developer)
 - runtime licenses (cost varies with vendor)
 - standardized extensions (“Object services” in OMG-speak) add more costs

CORBA Features/Limitations (cont)

- Harder to use
 - must
 - write IDL to provide interfaces and define data structures for passing information
 - transform IDL to Java stub & skeleton source code
 - compile stub/skeleton code to bytecode
 - different vendor's stub/skeleton bytecodes (‘.class’ files) are not compatible
 - details of using CORBA are different for each vendor: changing ORB vendors requires changing source code (eek!)

CORBA Features/Limitations (cont)

- Harder to use (continued)
 - more source code required to perform same tasks (vs. RMI)
 - no standard way to bootstrap communication
 - some vendor-specific options exist, but nothing standard across different vendors
 - no support for detecting software version mismatches
 - wing it or hand build it
 - most often, this isn't handled and problems arise at runtime, sometimes resulting in silent errors

CORBA features/limitations (cont)

- Add-on to the language
 - does not dovetail with standard Java features
 - breaks the object-oriented model when doing distributed communication
 - no built-in support for flattening objects onto/off of the wire
 - must use pseudo-objects or do the flattening/inflating by hand
 - no distributed garbage collection
 - e.g. don't know when server objects & resources are no longer needed

CORBA features/limitations (cont)


- Add-on to the language (continued)
 - BUT: Allows implementing in multiple languages
 - good for accessing legacy systems
 - especially if Java is not supported on the legacy platform
 - good for extending existing CORBA-based systems with Java
 - good for leveraging expertise in other languages
 - as long as there is a CORBA binding for that language

CORBA features/limitations (cont)

- Structured communication
 - must define pseudo-objects using IDL and translated to data structures
 - can't send objects

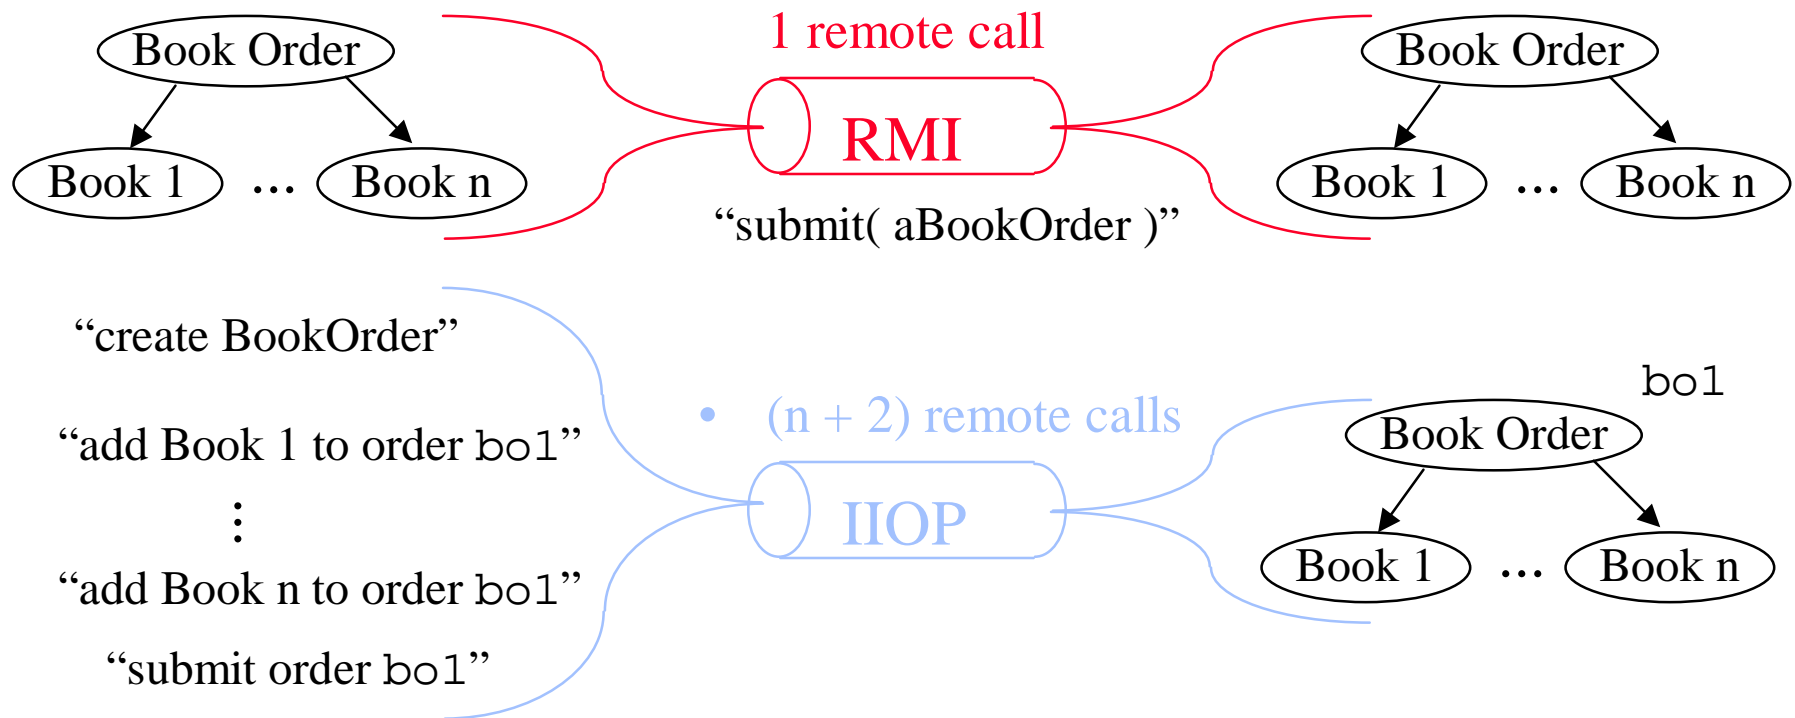
CORBA features/limitations (cont)

- Structured communication (continued)
 - there is substantial hand-coding to flatten/inflate a group of objects when sending their state information to a remote location



CORBA features/limitations (cont)

- Structured communication (continued)
 - lack of support for sending objects hampers design options and results in more remote calls across the wire



CORBA features/limitations (cont)

- Structured communication (continued)
 - assume an order for 3 books
 - we include marshalling/unmarshalling time
 - if 1 RMI call takes an average of 35 milliseconds (ms)
 - if 1 CORBA call takes
 - 21 ms => CORBA is • 105 ms & RMI is 35 ms
 - 10 ms => CORBA is • 50 ms & RMI is 35 ms
 - 6 ms => CORBA is • 35 ms & RMI is 35 ms

(numbers are illustrative only)

CORBA features/limitations (cont)

- Static codebase
 - can't send implementations across the wire, therefore
 - can't send previously unknown types of objects where they are needed to make computation more efficient
 - can't dynamically extend the behavior of running VM
 - using DII only allows discovering new remote types and making remote method calls upon them

CORBA features/limitations (cont)


- Performance
 - performance of CORBA IIOP is not as good as custom protocols provided by ORB vendors
 - this talk only discusses IIOP because inter-vendor support is important
 - traditional client/server speed is better in ORB implementations than in RMI/JDK 1.1.x.
 - how much better varies, depending on how the application is structured, what OS you are running, etc.

CORBA features/limitations (cont)

- Performance (continued)
 - Compared to RMI, CORBA limits possible designs and removes at least some of the performance gain when including time for
 - marshalling/unmarshalling of objects into/out of IDL structures
 - more calls across the wire, which are substantially more costly than local calls
 - code to perform version mismatch detection
 - distributed garbage collection issues

CORBA features/limitations (cont)

- CORBA is Maturing slowly



The Future of RMI (JDK 1.2)

- Support for mixing different reference types
 - secure, multi-cast, etc.
- Performance improvements
 - RMI-specific improvements
 - JDK improvements: HotSpot + VM/library tuning
- Better control over (de)serialization of objects and ability to unexport a “Remote” object in your VM
- Activation of server objects
 - launching a VM to run a server object

The Future of RMI (2.0?)

- RMI over IIOP
 - there is an initiative by the OMG and Javasoft to extend CORBA's IIOP to support features needed to run RMI over IIOP
 - no known availability date
 - requires OMG's "pass by-value" (see next slide)

The Future of CORBA

- (3.0?) OMG is looking into supporting pass by-value
 - first response to first RFP was completed end of January
 - no known availability date
- (???) Other proposed features which have varying degrees of vendor support
 - see <http://www.omg.org>

The Future of CORBA (Cont)

- There are hard problems to be solved in order to support pass by-value in a heterogeneous environment
 - different languages at each end (e.g. C++ vs. COBOL)
 - Might only allow passing state, not implementation
 - introduces new failure modes
 - appears to be current OMG direction (except for Java)
 - If only passing state, might fail if appropriate implementation not available at receiving end
 - introduces new failure modes
 - what about security?
 - no security “sandbox” except in Java

The Future of CORBA (Cont)

- Supporting pass by-value (continued)
 - almost all languages don't support dynamic code inclusion in a running program (e.g. C, COBOL, C++)
 - introduces new failure modes
 - version mismatch problem must be solved for a multi-language environment. If not implemented
 - systems development becomes much more error prone
 - can result in “silent” errors (bad computation)
 - would introduce new failure modes

The Future - CORBA continued

- Therefore: don't expect pass by-value to become standardized and supported on top of IIOP very soon

(my guess: not before 1Q'99, more likely 3Q'99 or later)

Useful URL's

- RMI pages at Javasoft
 - www.javasoft.com/products/jdk/rmi
- Object Management Group (OMG)
 - www.omg.org
- CORBA FAQ
 - www.cerfnet.com/~mpcline/corba-faq

Useful Newsgroups/email lists

- `rmi-users@java.sun.com`
- `comp.lang.java.*`
- `comp.object`
- `comp.object.corba`
- general CORBA mailing list?

Questions/Comments

- Now
- Later: gilstrap@inlink.com